

An introduction to hash tables

Outline

Discuss storing unrelated/unordered data

- IP addresses and domain names

Consider conversions between these two forms

Introduce the idea of hashing:

- Reducing $O(\ln(n))$ operations to $O(1)$

Consider some of the weaknesses

Supporting Example

Suppose we have a system which is associated with approximately 150 error conditions where

- Each of which is identified by an 8-bit number from 0 to 255, and
- When an identifier is received, a corresponding error-handling function must be called

We could create an array of 150 function pointers and to then call the appropriate function....

Supporting Example

```
#include <iostream>

void a() {
    std::cout
        << "Calling 'void a()'"
        << std::endl;
}

void b() {
    std::cout
        << "Calling 'void b()'"
        << std::endl;
}
```

```
int main() {
    void (*function_array[150])();
    unsigned char error_id[150];

    function_array[0] = a;
    error_id[0] = 3;
    function_array[1] = b;
    error_id[1] = 8;

    function_array[0]();
    function_array[1]();

    return 0;
}
```

Output:

```
% ./a.out
Calling 'void a()'
Calling 'void b()'
```

Supporting Example

Unfortunately, this is slow—we would have to do some form of binary search in order to determine which of the 150 slots corresponds to, for example, error-condition identifier `id = 198`

This would require approximately 6 comparisons per error condition

If there was a possibility of dynamically adding new error conditions or removing defunct conditions, this would substantially increase the effort required...

Supporting Example

A better solution:

- Create an array of size 256
- Assign those entries corresponding to valid error conditions

```
int main() {  
    void (*function_array[256])();  
    for ( int i = 0; i < 256; ++i ) {  
        function_array[i] = nullptr;  
    }  
  
    function_array[3] = a;  
    function_array[8] = b;  
  
    function_array[3]();  
    function_array[8]();  
  
    return 0;  
}
```

Question: }

- Is the increased speed worth the allocation of additional memory?

Keys

Our goal:

Store data so that all operations are $\Theta(1)$ time

Requirement:

The memory requirement should be $\Theta(n)$

In our supporting example, the corresponding function can be called in $\Theta(1)$ time and the array is less than twice the optimal size

Keys

In our example, we:

- Created an array of size 256
- Store each of 150 objects in one of the 256 entries
- The error code indicated which bin the corresponding function pointer was stored

In general, we would like to:

- Create an array of size M
- Store each of n objects in one of the M bins
- Have some means of determining the bin in which an object is stored

IP Addresses

Examples:

Suppose we want to associate IP addresses and any corresponding domain names

Recall that a 32-bit IP address are often written as four byte values from 0 to 255

- Consider 10000001 01100001 00001010 10110011₂
- This can be written as `http://129.97.10.179/`
- We use domain names because IP addresses are not human readable

IP Addresses

Similarly, the University of Waterloo has 2^{16} IP address control of names within its domain

- Any IP address starting with 129.97 belongs to UW
- This gives UW $256^2 = 65535$ IP addresses

The University of Waterloo currently uses ~60 % of the IP addresses

IP Addresses

Given an IP address, if we wanted to quickly find any associated domain name, we could create an array of size 65536 of strings:

```
int const MAX_IP_ADDRESSES = 65536  
string domain_name[MAX_IP_ADDRESSES];
```

For example, my computer is churchill.uwaterloo.ca and its IP address is `http://129.97.10.179/`

- The prefix 129.97 is common to all uWaterloo IP addresses
- As $179 + 256 \times 10 = 2739$, it follows that

```
domain_name[2739] = "churchill";
```

IP Addresses

Under IPv6, IP addresses are 128 bits

- It combines what is now implemented as subnets as well as allowing for many more IP addresses

Suppose uWaterloo is allocated 2^{64} IP addresses under this scheme

- We cannot allocate an array of size 2^{64}


IP Addresses

What if we want to associate domain names with IP addresses?

- Which entry in the array should be associated with the domain name `churchill.uwaterloo.ca` ?

Consider core routers on the Internet:

- They must associate each IP address with the next router that address should be passed to



Simpler problem

Let's try a simpler problem

- How do I store your examination grades so that I can access your grades in $\Theta(1)$ time?

Recall that each student is issued an 8-digit number

- How do I store your examination grades so that I can access your grades in $\Theta(1)$ time?
- Suppose Jane Doe has the number 20123456
- I can't create an array of size $10^8 \approx 1.5 \times 2^{26}$

Simpler problem

I could create an array of size 1000

- How could you convert an 8-digit number into a 3-digit number?
- First three digits might cause a problem: almost all students start with 201, 202, 203, 204, or 205
- The last three digits, however, are essentially random

Therefore, I could store Jane's examination grade

`grade[456] = 86;`

Simpler problem

Question:

- What is the likelihood that in a class of size 100 that no two students will have the same last three digits?
- Not very high

Simpler problem

Consequently, I have a function that maps a student onto a 3-digit number

- I can store something in that location
- Storing it, accessing it, and erasing it is $\Theta(1)$
- Problem: two or more students may map to the same number:
 - Wěi Wang has ID 20173456 and scored 85
 - Alma Ahmed has ID 2024456 and scored 87

⋮	⋮
454	
455	
456	86
457	
458	
459	
460	
461	
462	
463	79
464	
465	
⋮	⋮

The hashing problem

The process of mapping an object or a number onto an integer in a given range is called *hashing*

Problem: multiple objects may hash to the same value

- Such an event is termed a *collision*

Hash tables use a hash function together with a mechanism for dealing with collisions

IP Addresses

Going back to the problem with IP addresses

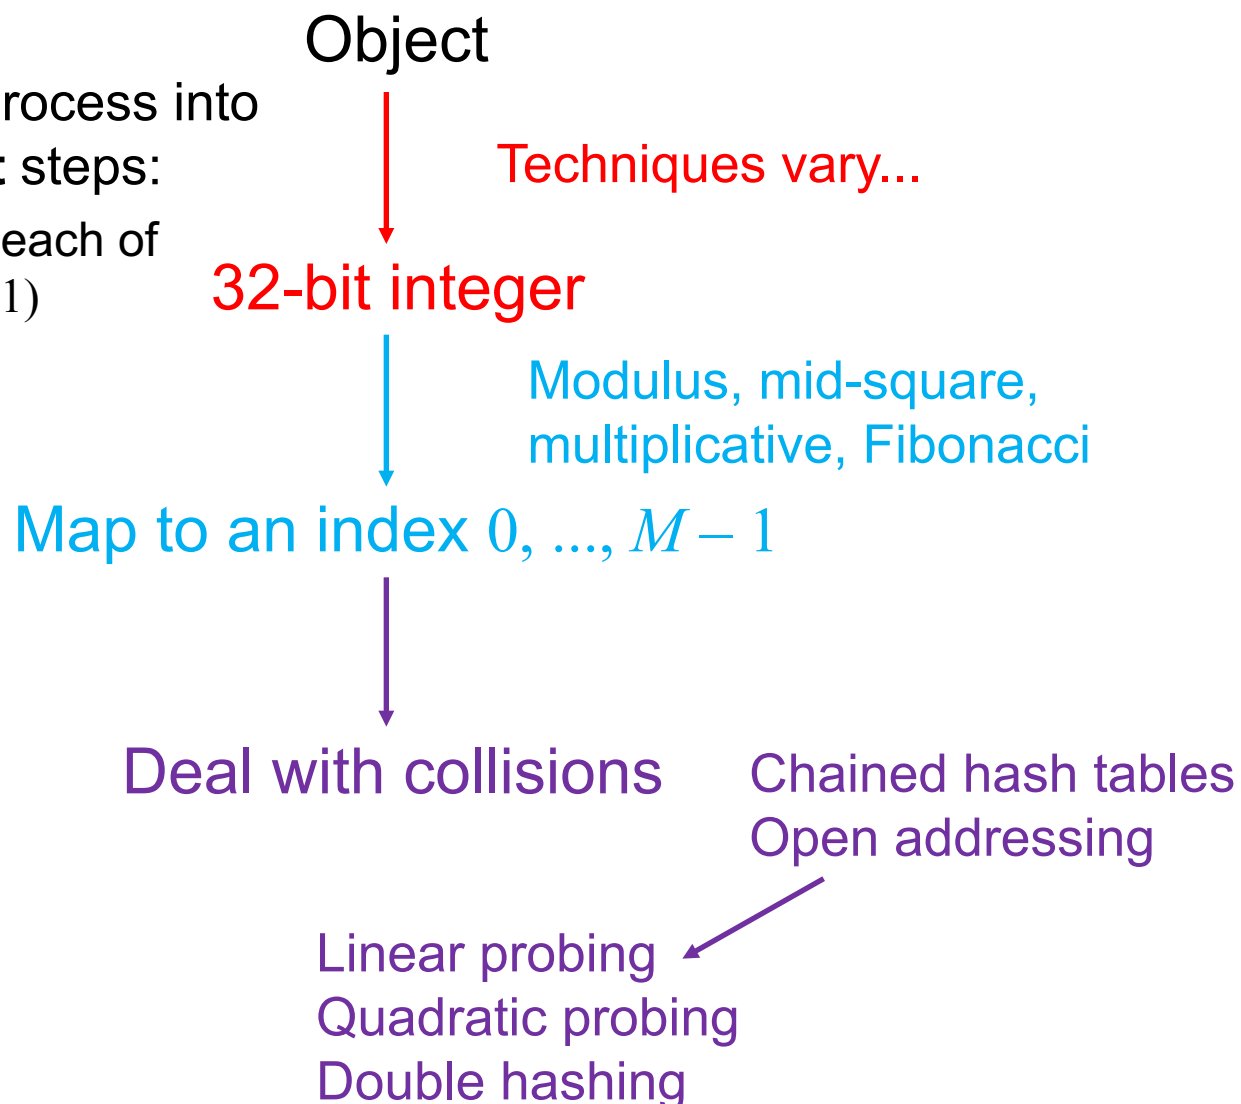
- We need a hash function to map an IP address to a smaller range
- We need a hash function to map a domain name to a smaller range

Mapping 129.97.10.179 onto a smaller range may seem easier, but a mechanism for mapping churchill.uwaterloo.ca onto a small range of integers may be more interesting

The hash process

We will break the process into three **independent** steps:

- We will try to get each of these down to $\Theta(1)$



Summary

Discuss storing unordered data

Discuss IP addresses and domain names

Consider conversions between these two forms

Introduce the idea of using a smaller array

- Converted “large” numbers into valid array indices
- Reduces $O(\ln(n))$ in arrays and AVL trees to $O(1)$

Discussed the issues with collisions

References

Wikipedia, http://en.wikipedia.org/wiki/Hash_table

- [1] Cormen, Leiserson, and Rivest, *Introduction to Algorithms*, McGraw Hill, 1990.
- [2] Weiss, *Data Structures and Algorithm Analysis in C++*, 3rd Ed., Addison Wesley.

These slides are provided for the ECE 250 *Algorithms and Data Structures* course. The material in it reflects Douglas W. Harder's best judgment in light of the information available to him at the time of preparation. Any reliance on these course slides by any party for any other purpose are the responsibility of such parties. Douglas W. Harder accepts no responsibility for damages, if any, suffered by any party as a result of decisions made or actions based on these course slides for any other purpose than that for which it was intended.

Hash functions

Outline

In this talk, we will discuss

- Finding 32-bit hash values using:
 - Predetermined hash values
 - Auto-incremented hash values
 - Address-based hash values
 - Arithmetic hash values
- Example: strings

Definitions

What is a hash of an object?

From Merriam-Webster:

a restatement of something that is already known

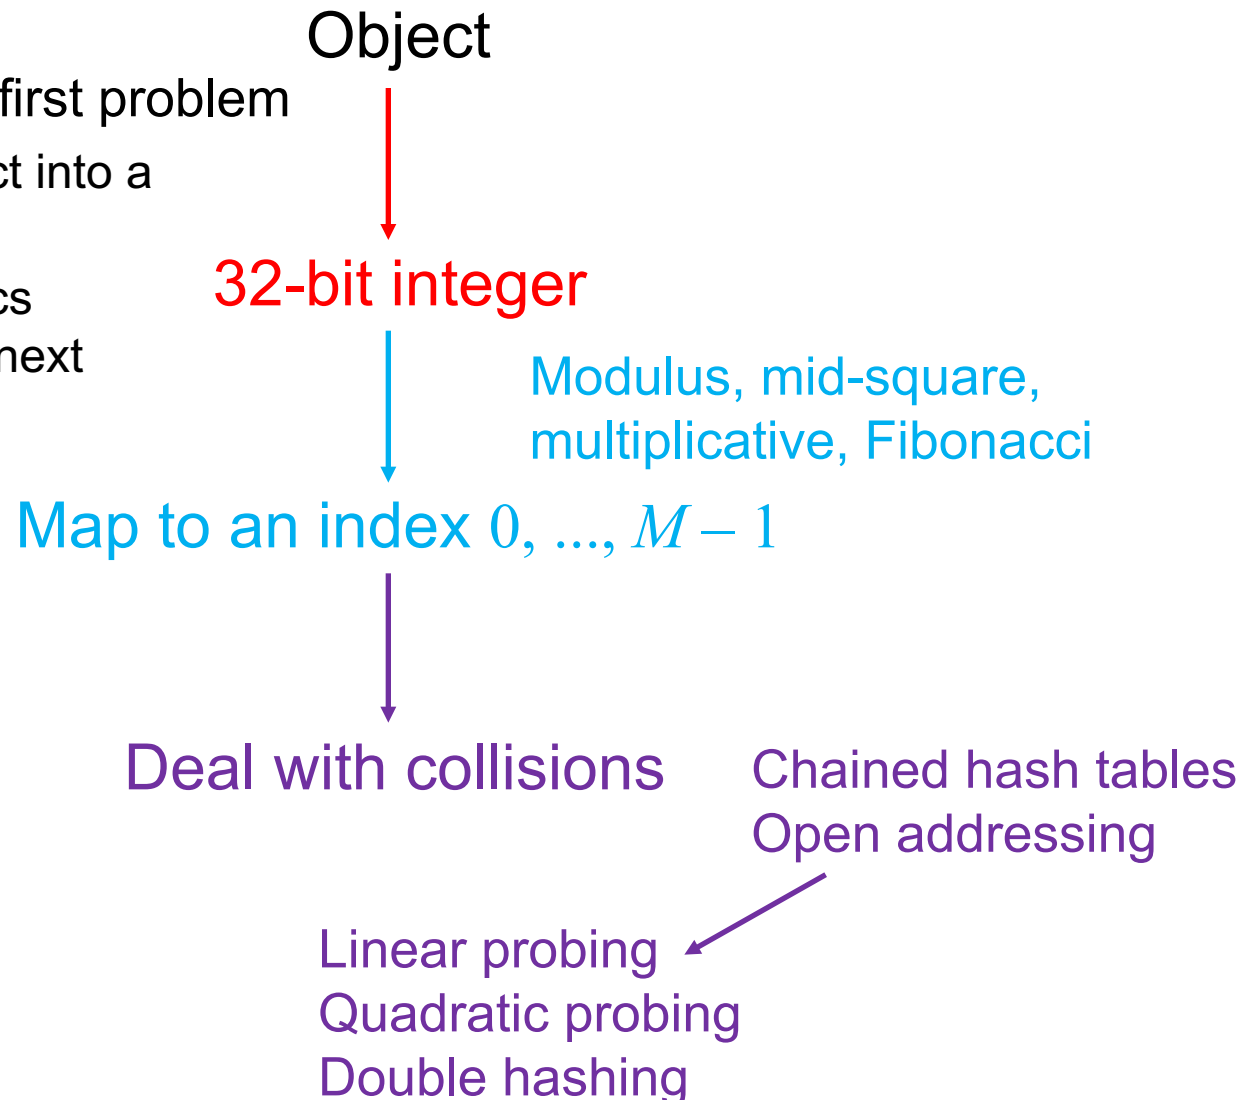
The ultimate goal is to map onto an integer range

$0, 1, 2, \dots, M - 1$

The hash process

We will look at the first problem

- Hashing an object into a 32-bit integer
- Subsequent topics will examine the next steps



Properties

Necessary properties of such a hash function h are:

- 1a. Should be fast: ideally $\Theta(1)$
- 1b. The hash value must be *deterministic*
 - It must always return the same 32-bit integer each time
- 1c. Equal objects hash to equal values
 - $x = y \Rightarrow h(x) = h(y)$
- 1d. If two objects are randomly chosen, there should be only a one-in- 2^{32} chance that they have the same hash value

Types of hash functions

We will look at two classes of hash functions

- Predetermined hash functions (explicit)
- Arithmetic hash functions (implicit)

Predetermined hash functions

The easiest solution is to give each object a unique number

```
class Class_name {
private:
    unsigned int hash_value; // int:           $-2^{31}, \dots, 2^{31} - 1$ 
                             // unsigned int:  $0, \dots, 2^{32} - 1$ 
public:
    Class_name();
    unsigned int hash() const;
};

Class_name::Class_name() {
    hash_value = ???;
}

unsigned int Class_name::hash() const {
    return hash_value;
}
```

Predetermined hash functions

For example, an auto-incremented static member variable

```
class Class_name {  
    private:  
        unsigned int hash_value;  
        static unsigned int hash_count;  
    public:  
        Class_name();  
        unsigned int hash() const;  
};  
  
unsigned int Class_name::hash_count = 0;
```

```
Class_name::Class_name() {  
    hash_value = hash_count;  
    ++hash_count;  
}  
  
unsigned int Class_name::hash() const {  
    return hash_value;  
}
```

Predetermined hash functions

Examples: All UW co-op student have two hash values:

- UW Student ID Number
- Social Insurance Number

Any 9-digit-decimal integer yields a 32-bit integer

$$\lg(10^9) = 29.897$$

Predetermined hash functions

If we only need the hash value while the object exists in memory, use the address:

```
unsigned int Class_name::hash() const {  
    return reinterpret_cast<unsigned int>( this );  
}
```


Predetermined hash functions

Predetermined hash values give each object a unique hash value

This is not always appropriate:

- Objects which are conceptually equal:

```
Rational x(1, 2);
```

```
Rational y(3, 6);
```

- Strings with the same characters:

```
string str1 = "Hello world!";
```

```
string str2 = "Hello world!";
```

These hash values must depend on the member variables

- Usually this uses arithmetic functions

Arithmetic Hash Values

An arithmetic hash value is a deterministic function that is calculated from the relevant member variables of an object

We will look at arithmetic hash functions for:

- Rational numbers, and
- Strings

Rational number class

What if we just add the numerator and denominator?

```
class Rational {  
    private:  
        int numer, denom;  
    public:  
        Rational( int, int );  
};  
  
unsigned int Rational::hash() const {  
    return static_cast<unsigned int>( numer ) +  
        static_cast<unsigned int>( denom );  
}
```

Rational number class

We could improve on this: multiply the denominator by a large prime:

```
class Rational {  
    private:  
        int numer, denom;  
    public:  
        Rational( int, int );  
};  
  
unsigned int Rational::hash() const {  
    return static_cast<unsigned int>( numer ) +  
        429496751*static_cast<unsigned int>( denom );  
}
```

Rational number class


For example, the output of

```
int main() {
    cout << Rational( 0, 1 ).hash() << endl;
    cout << Rational( 1, 2 ).hash() << endl;
    cout << Rational( 2, 3 ).hash() << endl;
    cout << Rational( 99, 100 ).hash() << endl;

    return 0;
}
```

is

```
429496751
858993503
1288490255
2239
```



<http://xkcd.com/571/>

Recall that arithmetic operations wrap on overflow

Rational number class

This hash function does not generate unique values

- The following pairs have the same hash values:

0/1	1327433019/800977868
-----	----------------------

1/2	534326814/1480277007
-----	----------------------

2/3	820039962/1486995867
-----	----------------------

- Finding rational numbers with matching hash values is very difficult:
- Finding these required the generation of 1 500 000 000 random rational numbers
- It is fast: $\Theta(1)$
- It does produce an even distribution

Rational number class

Problem:

- The rational numbers $1/2$ and $2/4$ have different values
- The output of

```
int main() {  
    cout << Rational( 1, 2 ).hash();  
    cout << Rational( 2, 4 ).hash();  
    return 0;  
}
```

is

```
858993503  
1717987006
```

Rational number class

Solution: divide through by the greatest common divisor

```
Rational::Rational( int a, int b ):numer(a), denom(b) {  
    int divisor = gcd( numer, denom );  
    numer /= divisor;  
    denom /= divisor;  
}  
  
int gcd( int a, int b) {  
    while( true ) {  
        if ( a == 0 ) {  
            return (b >= 0) ? b : -b;  
        }  
  
        b %= a;  
  
        if ( b == 0 ) {  
            return (a >= 0) ? a : -a;  
        }  
        a %= b;  
    }  
}
```


Rational number class

Problem:

- The rational numbers $\frac{1}{2}$ and $-\frac{1}{2}$ (negative) have different values
- The output of

```
int main() {  
    cout << Rational( 1, 2 ).hash();  
    cout << Rational( -1, -2 ).hash();  
    return 0;  
}
```

is

858993503

3435973793

Rational number class

Solution: define a normal form

- Require that the denominator is positive

```
Rational::Rational( int a, int b ):numer(a), denom(b) {  
    int divisor = gcd( numer, denom );  
    divisor = (denom >= 0) ? divisor : -divisor;  
    numer /= divisor;  
    denom /= divisor;  
}
```

String class

Two strings are equal if all the characters are equal and in the identical order

A string is simply an array of bytes:

- Each byte stores a value from 0 to 255

Any hash function must be a function of these bytes

String class

We could, for example, just add the characters:

```
unsigned int hash( const string &str ) {  
    unsigned int hash_value = 0;  
  
    for ( int k = 0; k < str.length(); ++k ) {  
        hash_value += str[k];  
    }  
  
    return hash_value;  
}
```

Summary

We have seen how a number of objects can be mapped onto a 32-bit integer

We considered

- Predetermined hash functions
 - Auto-incremented variables
 - Addresses
- Hash functions calculated using arithmetic

Next: map a 32-bit integer onto a smaller range $0, 1, \dots, M-1$

References

Wikipedia, http://en.wikipedia.org/wiki/Hash_function

- [1] Cormen, Leiserson, and Rivest, *Introduction to Algorithms*, McGraw Hill, 1990.
- [2] Weiss, *Data Structures and Algorithm Analysis in C++*, 3rd Ed., Addison Wesley.

These slides are provided for the ECE 250 *Algorithms and Data Structures* course. The material in it reflects Douglas W. Harder's best judgment in light of the information available to him at the time of preparation. Any reliance on these course slides by any party for any other purpose are the responsibility of such parties. Douglas W. Harder accepts no responsibility for damages, if any, suffered by any party as a result of decisions made or actions based on these course slides for any other purpose than that for which it was intended.

Mapping down to $0, \dots, M - 1$

Outline

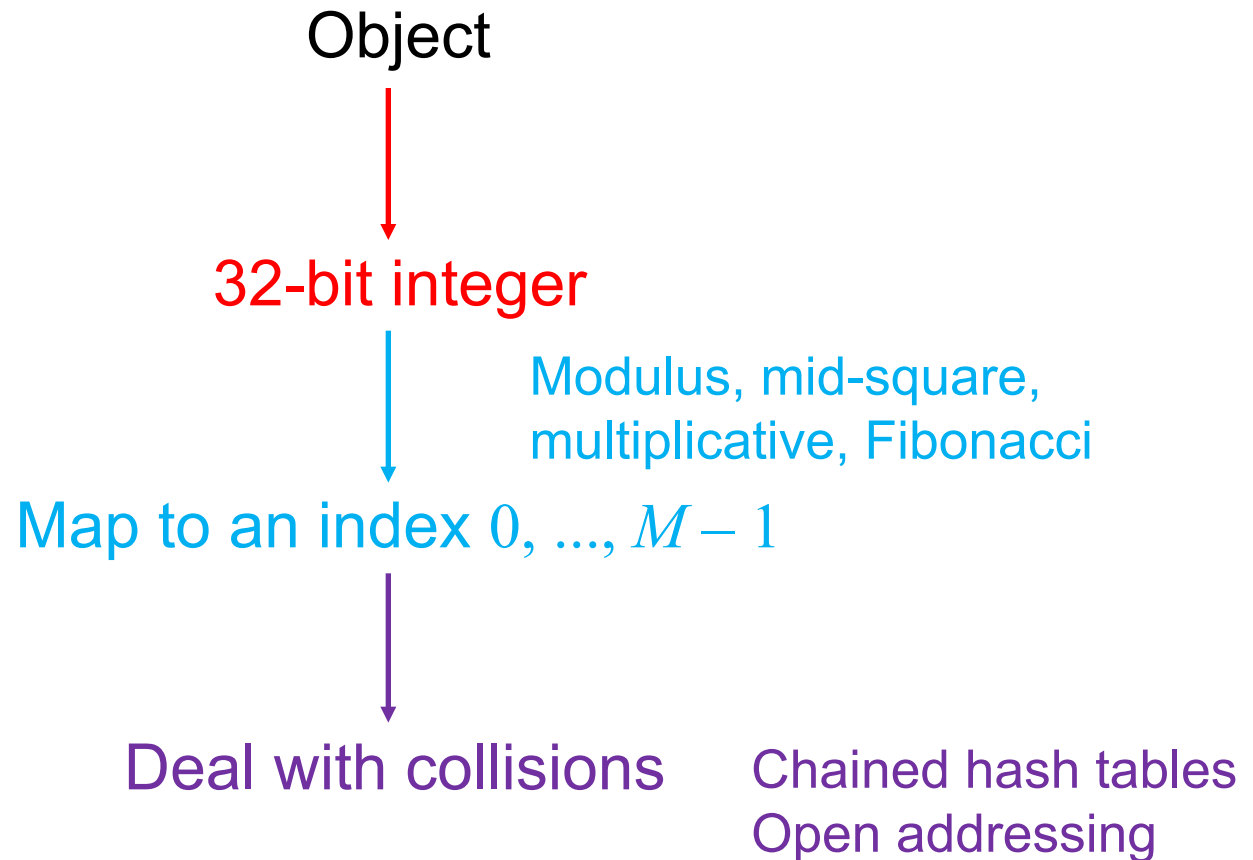
Previously, we considered means for calculating 32-bit hash values

- Explicitly defined hash values
- Implicitly calculated hash values

Practically, we will require a hash value on the range $0, \dots, M - 1$:

- The modulus operator %
- Review of bitwise operations

The hash process



Properties

Necessary properties of this mapping function h_M are:

- 2a. Must be fast: $\Theta(1)$
- 2b. The hash value must be *deterministic*
 - Given n and M , $h_M(n)$ must always return the same value
- 2c. If two objects are randomly chosen, there should be only a one-in- M chance that they have the same value from 0 to $M - 1$

Modulus operator

Easiest method: return the value modulus M

```
unsigned int hash_M( unsigned int n, unsigned int M ) {  
    return n % M;  
}
```

Unfortunately, calculating the modulus (or remainder) is expensive

- We can simplify this if $M = 2^m$
- We can use logic operations
 - We will review bitwise operations: left and right shift and bit-wise and

The bitwise operators: & << >>

Suppose I want to calculate

$$7985325 \% 100$$

The modulo is a power of ten: $100 = 10^2$

- In this case, take the last **two** decimal digits: 25

Similarly, $7985325 \% 10^3 = 325$

- We set the appropriate digits to 0:

000025 and **0000**325

The bitwise operators: & << >>

The same works in base 2:

$$100011100101_2 \% 10000_2$$

The modulo is a power of 2: $10000_2 = 2^4$

- In this case, take the last **four** bits: 0101

Similarly, $100011100101_2 \% 1000000_2 == 100101$,

- We set the appropriate digits to 0:

$$000000000101 \text{ and } 000000100101$$

The bitwise operators: & << >>

To zero all but the last n bits, select the last n bits using *bitwise and*:

$$1000\ 1110\ \textcolor{red}{0101}_2 \ \& \ 0000\ 0000\ \textcolor{red}{1111}_2 \rightarrow 0000\ 0000\ \textcolor{red}{0101}_2$$

$$1000\ 11\textcolor{blue}{10}\ \textcolor{blue}{0101}_2 \ \& \ 0000\ 00\textcolor{blue}{11}\ \textcolor{blue}{1111}_2 \rightarrow 0000\ 00\textcolor{blue}{10}\ \textcolor{blue}{0101}_2$$

The bitwise operators: & << >>

Similarly, multiplying or dividing by powers of 10 is easy:

$$7985325 * 100$$

The multiplier is a power of ten: $100 = 10^2$

- In this case, add **two** zeros: 798532500

Similarly, $7985325 / 10^3 = 7985$

- Just add the appropriate number of zeros or remove the appropriate number of digits

The bitwise operators: & << >>

The same works in base 2:

$$100011100101_2 * 10000_2$$

The modulo is a power of 2: $10000_2 = 2^4$

- In this case, add **four** zeros: 1000111001010000

Similarly, $100011100101_2 / 1000000_2 == 100011$

The bitwise operators: & << >>

This can be done mechanically by shifting the bits appropriately:

$$1000\ 1110\ 0101_2 \ll 4 == 1000\ 1110\ 0101\ 0000_2$$

$$1000\ 1110\ 0101_2 \gg 6 == 10\ 0011_2$$

Powers of 2 are now easy to calculate:

$$1_2 \ll 4 == 1\ 0000_2 \quad // \quad 2^4 = 16$$

$$1_2 \ll 6 == 100\ 0000_2 \quad // \quad 2^6 = 64$$

Modulo a power of two

The implementation using the modulus/remainder operator:

```
unsigned int hash_M( unsigned int n, unsigned int m ) {  
    return n & ((1 << m) - 1);  
}
```

Modulo a power of two

Problem:

- Suppose that the hash function h is always even
- An even number modulo a power of two is still even

Example: memory allocations are multiples of word size

- On a 64-bit computer, addresses returned by new will be multiples of 8
- Thus, if $x \neq y$, $h(x) \neq h(y)$
- However, the probability that $h_M(x) = h_M(y)$ is one in $M/8$
 - This is not one in M

Modulo a power of two

For some objects, it is worse:

- Instance of `Single_node<int>` on `ecelinux` always have 10000 as the last five bits
 - This increases the probability of a collision to one in $M/32$

Fortunately, the multiplicative method resolves this issue

Chained hash tables

Outline

We have:

- Discussed techniques for hashing
- Discussed mapping down to a given range $0, \dots, M - 1$

Now we must deal with collisions

- Numerous techniques exist
- Containers in general
 - Specifically linked lists

Background

First, a review:

- We want to store objects in an array of size M
- We want to quickly calculate the bin where we want to store the object
 - We came up with hash functions—hopefully $\Theta(1)$
 - Perfect hash functions (no collisions) are difficult to design
- We will look at some schemes for dealing with collisions

Implementation

Consider associating each bin with a linked list:

```
template <class Type>
class Chained_hash_table {
    private:
        int table_capacity;
        int table_size;
        Single_list<Type> *table;


        unsigned int hash( Type const & );

    public:
        Chained_hash_table( int = 16 );
        int count( Type const & ) const;
        void insert( Type const & );
        int erase( Type const & );
        // ...
};
```


Implementation

The constructor creates an array of n linked lists

```
template <class Type>
Chained_hash_table::Chained_hash_table( int n ):
    table_capacity( std::max( n, 1 ) ),
    table_size( 0 ),
    table( new Single_list<Type>[table_capacity] ) {
    // empty constructor
}
```




Implementation

The function hash will determine the bin of an object:

```
template <class Type>
int Chained_hash_table::hash( Type const &obj ) {
    return hash_M( obj.hash(), capacity() );
}
```

Recall:

- `obj.hash()` returns a 32-bit non-negative integer
- `unsigned int hash_M(obj, M)` returns a value in $0, \dots, M-1$




Implementation

Other functions mapped to corresponding linked list functions:


```
template <class Type>
void Chained_hash_table::insert( Type const &obj ) {
    unsigned int bin = hash( obj );

    if ( table[bin].count( obj ) == 0 ) {
        table[bin].push_front( obj );
        ++table_size;
    }
}
```



Implementation


```
template <class Type>
int Chained_hash_table::count( Type const &obj ) const {
    return table[hash( obj )].count( obj );
}
```



Implementation

```
template <class Type>
int Chained_hash_table::erase( Type const &obj ) {
    unsigned int bin = hash( obj );


    if ( table[bin].erase( obj ) ) {
        --table_size;
        return 1;
    } else {
        return 0;
    }
}
```



Example

As an example, let's store hostnames and allow a fast look-up of the corresponding IP address


- We will choose the bin based on the host name
- Associated with the name will be the IP address
- *E.g.*, ("optimal", 129.97.94.57)



Example

We will store strings and the hash value of a string will be the last 3 bits of the first character in the host name


- The hash of "optimal" is based on "o"



Example

The following is a list of the binary representation of each letter:

- "a" is 1 and it cycles from there...




a	01100 001	n	01101 110
b	01100 010	o	01101 111
c	01100 011	p	01110 000
d	01100 100	q	01110 001
e	01100 101	r	01110 010
f	01100 110	s	01110 011
g	01100 111	t	01110 100
h	01101 000	u	01110 101
i	01101 001	v	01110 110
j	01101 010	w	01110 111
k	01101 011	x	01111 000
l	01101 100	y	01111 001
m	01101 101	z	01111 010

Example


Our hash function is

```
unsigned int hash( string const &str ) {  
    // the empty string "" is hashed to 0  
    if str.length() == 0 ) {  
        return 0;  
    }  
  
    return str[0] & 7;  
}
```




Example

Starting with an array of 8 empty linked lists



Example


The pair ("optimal", 129.97.94.57) is entered into bin
01101**111** = 7



Example

Similarly, as "c" hashes to 3


- The pair ("cheetah", 129.97.94.45) is entered into bin 3



Example

The "w" in Wellington also hashes to 7

- ("wellington", 129.97.94.42) is entered into bin 7



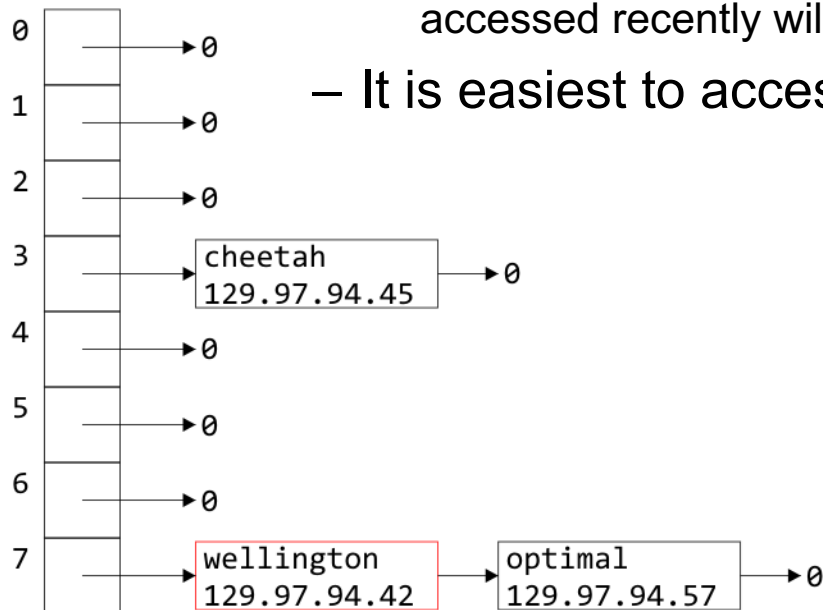
Example

Why did I use `push_front` from the linked list?

- Do I have a choice?
- A good heuristic is

“unless you know otherwise, data which has been accessed recently will be accessed again in the near future”


- It is easiest to access data at the front of a linked list



Heuristics include rules of thumb, educated guesses, and intuition


Example

Similarly we can insert the host names "agustin" and "lowpower"



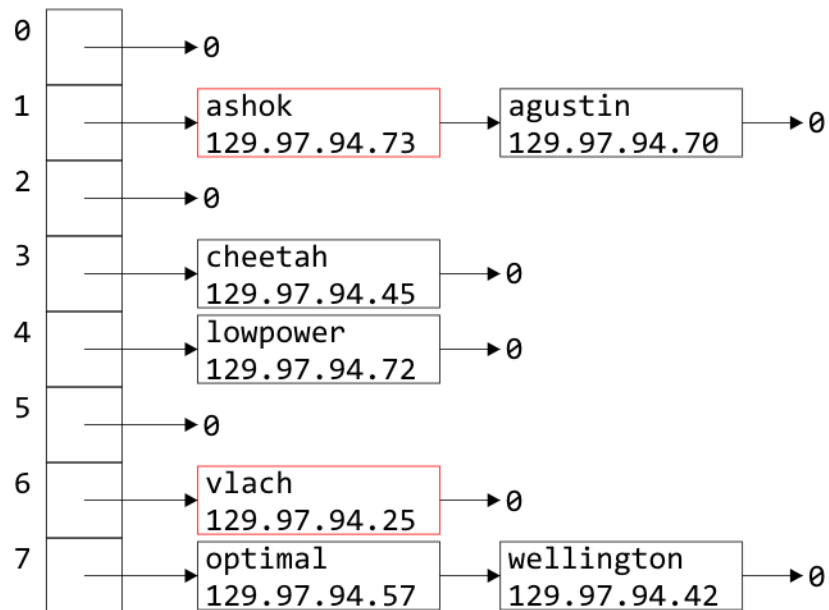
Example

If we now wanted the IP address for "optimal", we would simply hash "optimal" to 7, walk through the linked list, and access 129.97.94.57 when we access the node containing the relevant string



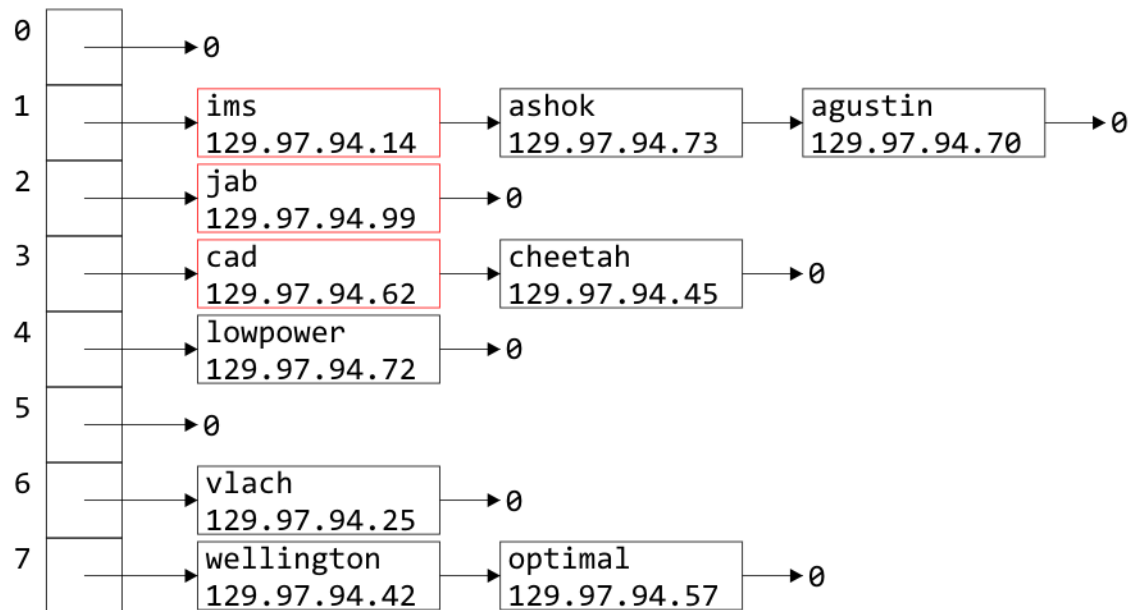
Example

Similarly, "ashok" and "vlach" are entered into bin 7



Example

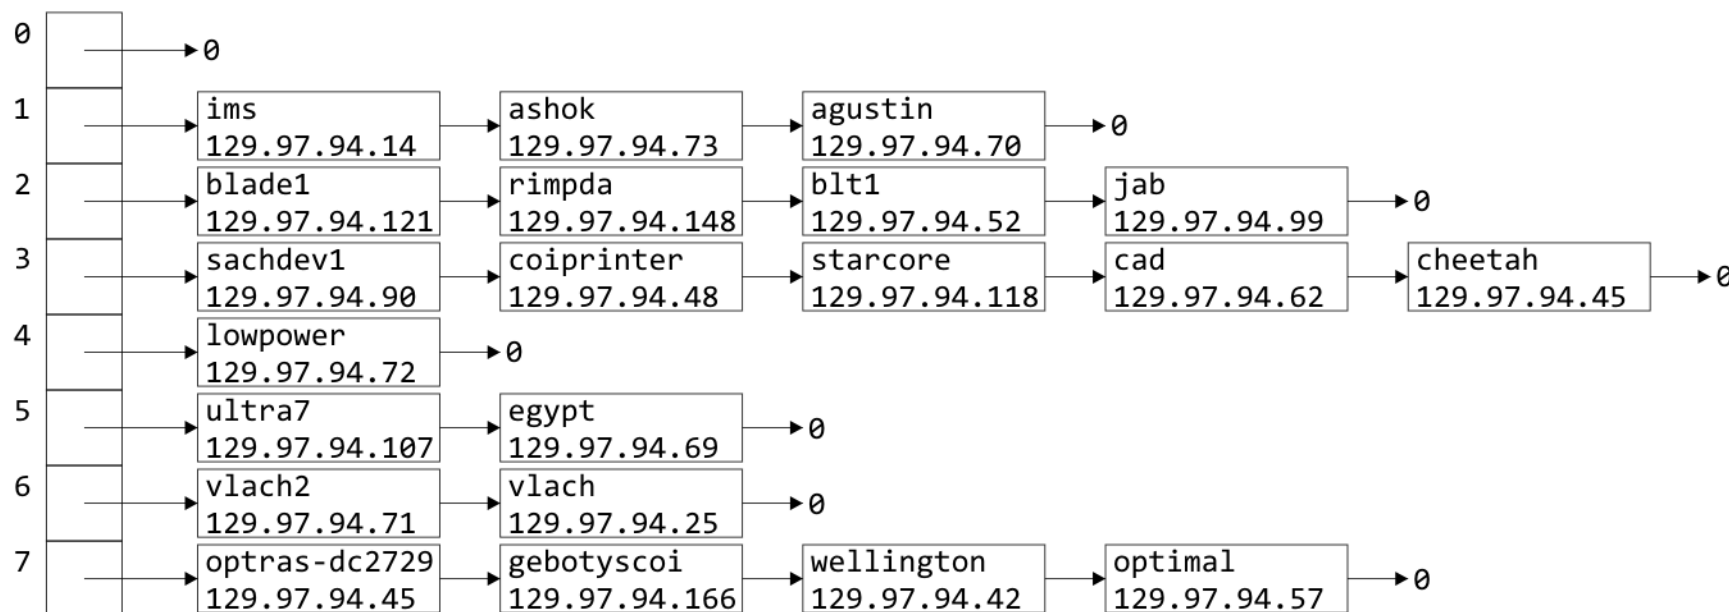
Inserting "ims", "jab", and "cad" doesn't even out the bins



Example

Indeed, after 21 insertions, the linked lists are becoming rather long

- We were looking for $\Theta(1)$ access time, but accessing something in a linked list with k objects is $\mathcal{O}(k)$



Load Factor

To describe the length of the linked lists, we define the *load factor* of the hash table:

$$= \frac{n}{M}$$

This is the average number of objects per bin

- This assumes an even distribution

Right now, the load factor is $\lambda = 21/8 = 2.625$

- The average bin has 2.625 objects

Load Factor

If the load factor becomes too large, access times will start to increase: $O(\lambda)$

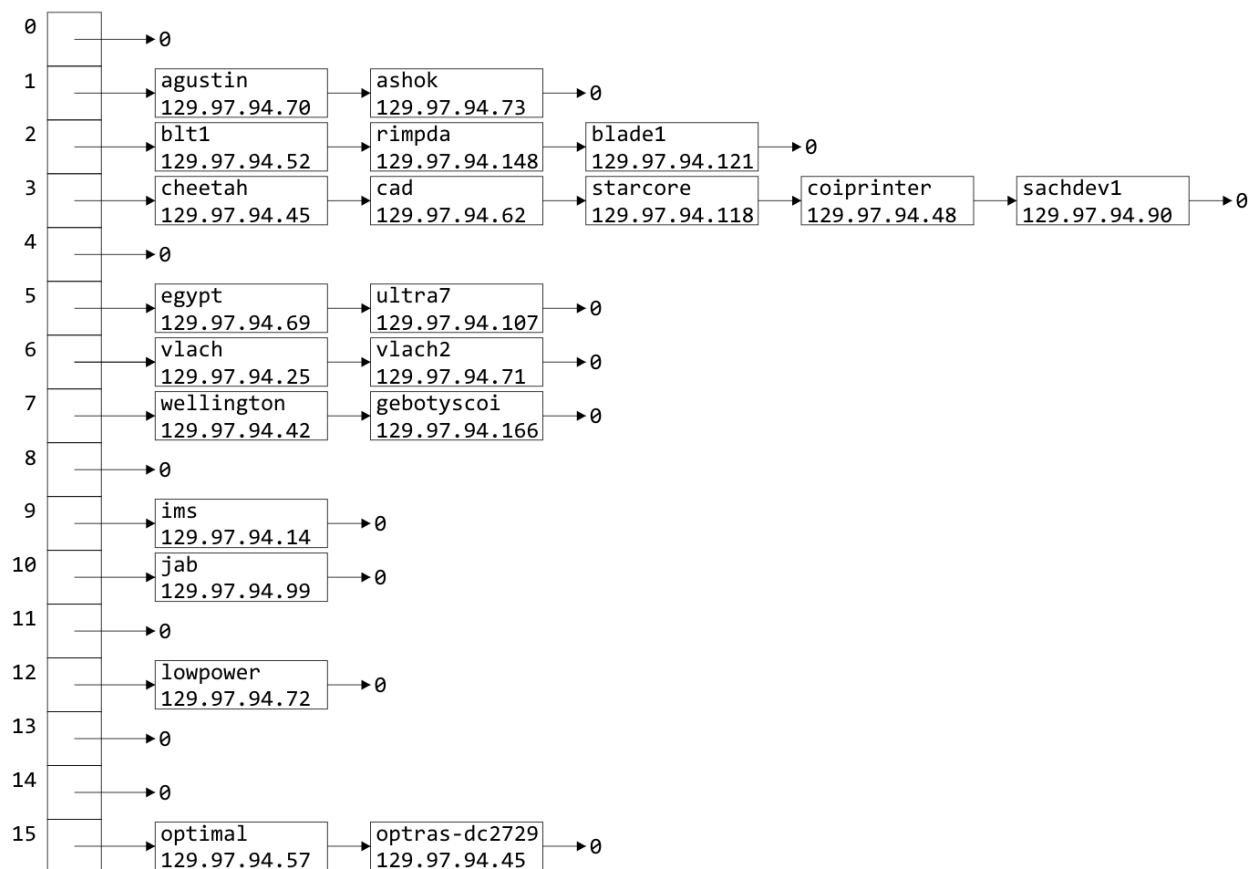
The most obvious solution is to double the size of the hash table

- Unfortunately, the hash function must now change
- In our example, the doubling the hash table size requires us to take, for example, the last four bits

Doubling Size

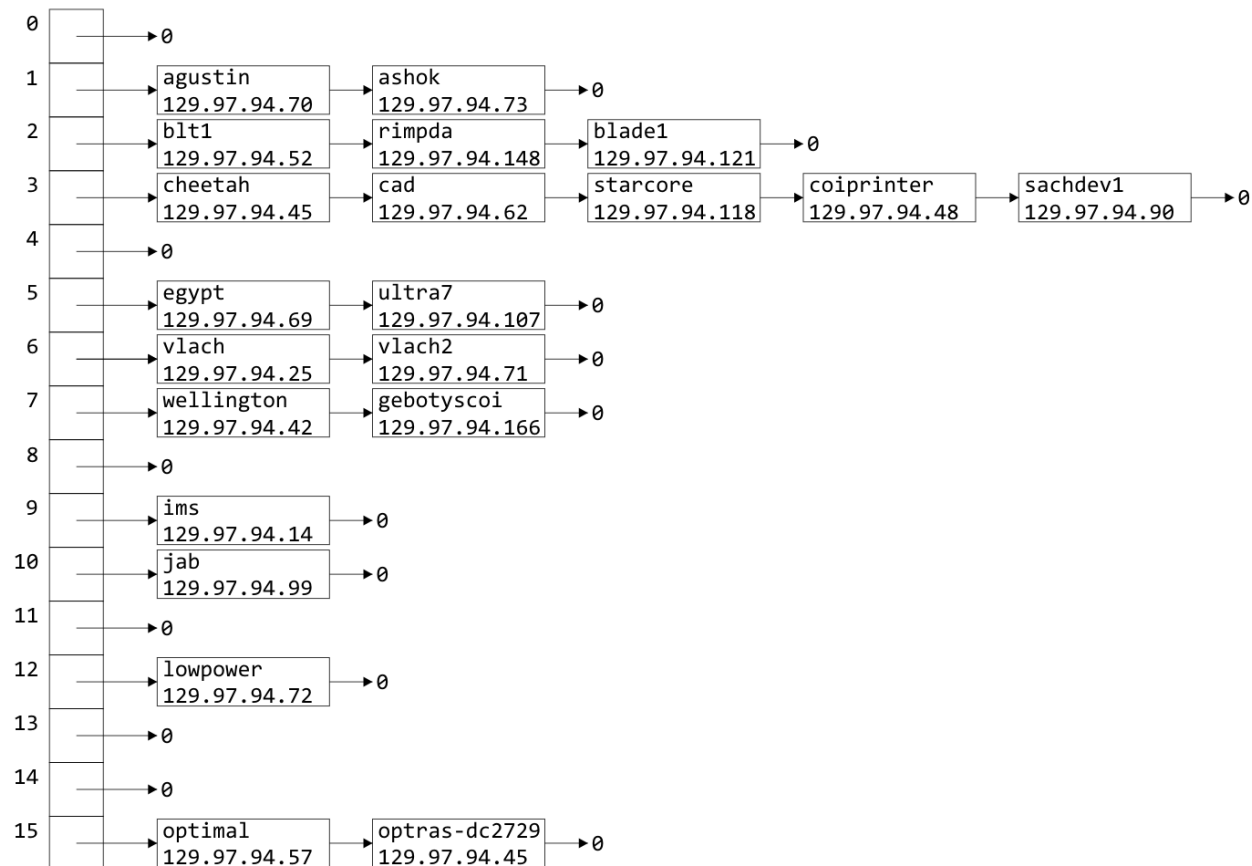
The load factor is now $\lambda = 1.3125$

- Unfortunately, the distribution hasn't improved much



Doubling Size

There is significant *clustering* in bins 2 and 3 due to the choice of host names



Choosing a Good Hash Function

We choose a very poor hash function:

- We looked at the first letter of the host name

Unfortunately, all these are also actual host names:

ultra7 ultra8 ultra9 ultra10 ultra11

ultra12 ultra13 ultra14 ultra15 ultra16 ultra17

blade1 blade2 blade3 blade4 blade5

This will cause clustering in bins 2 and 5

- Any hash function based on anything other than every letter will cause clustering

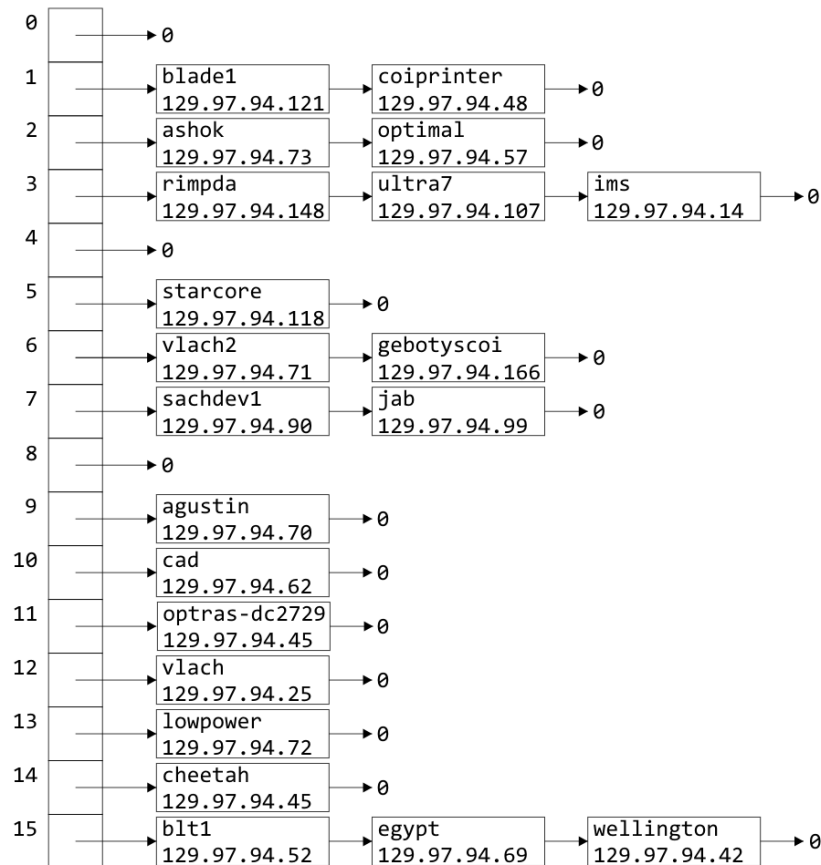
Choosing a Good Hash Function

Let's go back to the hash function defined previously:

```
unsigned int hash( string const &str ) {  
    unsigned int hash_value = 0;  
  
    for ( int k = 0; k < str.length(); ++k ) {  
        hash_value = 12347*hash_value + str[k];  
    }  
  
    return hash_value;  
}
```

Choosing a Good Hash Function


This hash function yields a much nicer distribution:



Choosing a Good Hash Function



When we insert the names of all 55 names, we would have a load factor $\lambda = 3.4375$

- Clearly there are not exactly 3.4375 objects per bin
- How can we tell if this is a good hash function?
- Can we expect exactly 3.4375 objects per bin?
 - Clearly no...
- The answer is with statistics...



Choosing a Good Hash Function

We would expect the number of bins which hold k objects to approximately follow a Poisson distribution



Problems with Linked Lists

One significant issue with chained hash tables using linked lists

- It requires extra memory
- It uses dynamic memory allocation

Total memory:

16 bytes

- A pointer to an array, initial and current number of bins, and the size
- + 12M bytes (8M if we remove count from `Single_list`)
- + 8n bytes if each object is 4 bytes

Problems with linked lists

For faster access, we could replace each linked list with an AVL tree (assuming we can order the objects)

- The access time drops to $O(\ln(\lambda))$
- The memory requirements are increased by $\Theta(n)$, as each node will require two pointers

We could look at other techniques:

- Scatter tables: use other bins and link the bins
- Use an alternate memory allocation model

Black Board Example

Use the hash function

```
unsigned int hash( unsigned int n ) { return n % 10; }
```

to enter the following 15 numbers into a hash table with 10 bins:

534, 415, 465, 459, 869, 442, 840, 180, 450, 265, 23, 946, 657, 3, 29

Summary

The easiest way to deal with collisions is to associate each bin with a container

We looked at bins of linked lists

- The example used host names and IP addresses
- We defined the load factor $\lambda = n/M$
- Discussed doubling the number of bins
- Our goals are to choose a good hash function and to keep the load factor low
- We discussed alternatives

Next we will see a different technique using only one array of bins: open addressing

References

Wikipedia, http://en.wikipedia.org/wiki/Hash_function

- [1] Cormen, Leiserson, and Rivest, *Introduction to Algorithms*, McGraw Hill, 1990.
- [2] Weiss, *Data Structures and Algorithm Analysis in C++*, 3rd Ed., Addison Wesley.

These slides are provided for the ECE 250 *Algorithms and Data Structures* course. The material in it reflects Douglas W. Harder's best judgment in light of the information available to him at the time of preparation. Any reliance on these course slides by any party for any other purpose are the responsibility of such parties. Douglas W. Harder accepts no responsibility for damages, if any, suffered by any party as a result of decisions made or actions based on these course slides for any other purpose than that for which it was intended.