

Graph theory and the Graph ADT

Outline

A graph is an abstract data type for storing adjacency relations

- We start with definitions:
 - Vertices, edges, degree and sub-graphs
- We will describe paths in graphs
 - Simple paths and cycles
- Definition of connectedness
- Weighted graphs
- We will then reinterpret these in terms of directed graphs
- Directed acyclic graphs

Outline

We will define an Undirected Graph ADT as a collection of *vertices*

$$V = \{v_1, v_2, \dots, v_n\}$$

- The number of vertices is denoted by

$$|V| = n$$

- Associated with this is a collection E of unordered pairs $\{v_i, v_j\}$ termed *edges* which connect the vertices

There are a number of data structures that can be used to implement abstract undirected graphs


- Adjacency matrices
- Adjacency lists

Graphs

Consider this collection of vertices

$$V = \{v_1, v_2, \dots, v_9\}$$

where $|V| = n$




Undirected graphs

Associated with these vertices are $|E| = 5$ edges

$$E = \{\{v_1, v_2\}, \{v_3, v_5\}, \{v_4, v_8\}, \{v_4, v_9\}, \{v_6, v_9\}\}$$

- The pair $\{v_j, v_k\}$ indicates that both vertex v_j is adjacent to vertex v_k and vertex v_k is adjacent to vertex v_j



Undirected graphs

We will assume in this course that a vertex is never adjacent to itself

- For example, $\{v_1, v_1\}$ will not define an edge

The maximum number of edges in an undirected graph is

$$|E| \leq \binom{|V|}{2} = \frac{|V|(|V|-1)}{2} = O(|V|^2)$$


An undirected graph

Example: given the $|V| = 7$ vertices

$$V = \{A, B, C, D, E, F, G\}$$

and the $|E| = 9$ edges

$$E = \{\{A, B\}, \{A, D\}, \{A, E\}, \{B, C\}, \{B, D\}, \{B, E\}, \{C, E\}, \{C, F\}, \{D, E\}\}$$



Degree


The degree of a vertex is defined as the number of adjacent vertices

$$\text{degree}(A) = \text{degree}(D) = \text{degree}(C) = 3$$

$$\text{degree}(B) = \text{degree}(E) = 4$$

$$\text{degree}(F) = 1$$



$$\text{degree}(G) = 0$$



Those vertices adjacent to a given vertex are its *neighbors*



Sub-graphs

A *sub-graph* of a graph is a subset of the vertices and a subset of the edges that connected the subset of vertices in the original graph



Vertex-induced sub-graphs

A *vertex-induced sub-graph* is a subset of the vertices where the edges are all edges in the original graph that originally



Paths

A path in an undirected graph is an ordered sequence of vertices

$(v_0, v_1, v_2, \dots, v_k)$


where $\{v_{j-1}, v_j\}$ is an edge for $j = 1, \dots, k$

- Termed *a path from v_0 to v_k*
- The length of this path is k

Paths

A path of length 4:


(A, B, E, C, F)



Paths

A path of length 5:


(A, B, E, C, B, D)



Paths

A *trivial* path of length 0:

(A)



Simple paths

A *simple path* has no repetitions other than perhaps the first and last vertices

A *simple cycle* is a simple path of at least two vertices with the first and last vertices equal

- Note: these definitions are not universal




<http://xkcd.com/140/>


Connectedness

Two vertices v_i, v_j are said to be *connected* if there exists a path from v_i to v_j

A graph is connected if there exists a path between any two vertices



A connected graph




An unconnected graph

Weighted graphs

A weight may be associated with each edge in a graph

- This could represent distance, energy consumption, cost, etc.
- Such a graph is called a *weighted graph*


Pictorially, we will represent weights by numbers next to the edges



Weighted graphs

The *length* of a path within a weighted graph is the sum of all of the edges which make up the path


- The length of the path (A, D, G) in the following graph is $5.1 + 3.7 = 8.8$



Weighted graphs

Different paths may have different weights


- Another path is (A, C, F, G) with length $1.2 + 1.4 + 4.5 = 7.1$



Weighted graphs

Problem: find the shortest path between two vertices




- Here, the shortest path from A to H is (A, C, F, D, E, G) with length 5.7



Trees

A graph is a tree if it is connected and there is a unique path between any two vertices

- Three trees on the same eight vertices



Consequences:

- The number of edges is $|E| = |V| - 1$
- The graph is *acyclic*, that is, it does not contain any cycles
- Adding one more edge must create a cycle
- Removing any one edge creates two disjoint non-empty sub-graphs

Trees





Any tree can be converted into a rooted tree by:

- Choosing any vertex to be the root
- Defining its neighboring vertices as its children

and then recursively defining:

- All neighboring vertices other than that one designated its parent are now defined to be that vertex's children

Given this tree, here are three rooted trees associated with it



Forests


A forest is any graph that has no cycles

Consequences:

- The number of edges is $|E| < |V|$
- The number of trees is $|V| - |E|$
- Removing any one edge adds one more tree to the forest

Here is a forest with 22 vertices and 18 edges

- There are four trees



Directed graphs

In a *directed graph*, the edges on a graph are be associated with a direction

- Edges are ordered pairs (v_j, v_k) denoting a connection from v_j to v_k
- The edge (v_j, v_k) is different from the edge (v_k, v_j)

Streets are directed graphs:


- In most cases, you can go two ways unless it is a one-way street

Directed graphs

Given our graph of nine vertices $V = \{v_1, v_2, \dots, v_9\}$

- These six pairs (v_j, v_k) are *directed edges*

$$E = \{(v_1, v_2), (v_3, v_5), (v_5, v_3), (v_6, v_9), (v_8, v_4), (v_9, v_4)\}$$



Directed graphs

The maximum number of directed edges in a directed graph is

$$|E| \leq 2 \binom{|V|}{2} = 2 \frac{|V|(|V|-1)}{2} = |V|(|V|-1) = O(|V|^2)$$

In and out degrees


The degree of a vertex must be modified to consider both cases:

- The *out-degree* of a vertex is the number of vertices which are adjacent to the given vertex
- The *in-degree* of a vertex is the number of vertices which this vertex is adjacent to

In this graph:

$$\text{in_degree}(v_1) = 0 \quad \text{out_degree}(v_1) = 2$$

$$\text{in_degree}(v_5) = 2 \quad \text{out_degree}(v_5) = 3$$




Sources and sinks

Some definitions:

- Vertices with an in-degree of zero are described as *sources*
- Vertices with an out-degree of zero are described as *sinks*

In this graph:

- Sources: v_1, v_6, v_7
- Sinks: v_2, v_9



Paths

A path in a directed graph is an ordered sequence of vertices

$$(v_0, v_1, v_2, \dots, v_k)$$


where (v_{j-1}, v_j) is an edge for $j = 1, \dots, k$

A path of length 5 in this graph is

$$(v_1, v_4, v_5, v_3, v_5, v_2)$$

A simple cycle of length 3 is

$$(v_8, v_4, v_5, v_8)$$




Connectedness

Two vertices v_j, v_k are said to be *connected* if there exists a path from v_j to v_k

- A graph is *strongly connected* if there exists a directed path between any two vertices
- A graph is *weakly connected* there exists a path between any two vertices that ignores the direction

In this graph:


- The sub-graph $\{v_3, v_4, v_5, v_8\}$ is strongly connected
- The sub-graph $\{v_1, v_2, v_3, v_4, v_5, v_8\}$ is weakly connected



Weighted directed graphs

In a weighted directed graphs, each edge is associated with a value

Unlike weighted undirected graphs, if both (v_j, v_k) and (v_k, v_j) are edges, it is not required that they have the same weight




Directed acyclic graphs


A *directed acyclic graph* is a directed graph which has no cycles

- These are commonly referred to as DAGs
- They are graphical representations of partial orders on a finite number of elements

These two are DAGs:



This directed graph is not acyclic:



Directed acyclic graphs

Applications of directed acyclic graphs include:


- The parse tree constructed by a compiler
- A reference graph that can be garbage collected using simple reference counting
- Dependency graphs such as those used in instruction scheduling and **makefiles**
- Dependency graphs between classes formed by inheritance relationships in object-oriented programming languages
- Information categorization systems, such as folders in a computer
- Directed acyclic word graph data structure to memory-efficiently store a set of strings (words)

Reference: http://en.wikipedia.org/wiki/Directed_acyclic_graph

Representations

How do we store the adjacency relations?

- Binary-relation list
- Adjacency matrix
- Adjacency list




Binary-relation list

The most inefficient is a relation list:

- A container storing the edges

$\{(1, 2), (1, 4), (3, 5), (4, 2), (4, 5), (5, 2), (5, 3), (5, 8), (6, 9), (7, 9), (8, 4)\}$

- Requires $\Theta(|E|)$ memory
- Determining if v_j is adjacent to v_k is $O(|E|)$
- Finding all neighbors of v_j is $\Theta(|E|)$




Adjacency matrix

Requiring more memory but also faster, an adjacency matrix

- The matrix entry (j, k) is set to true if there is an edge (v_j, v_k)

	1	2	3	4	5	6	7	8	9
1		T		T					
2									
3					T				
4		T			T				
5		T	T					T	
6									T
7								T	
8									
9									

- Requires $\Theta(|V|^2)$ memory
- Determining if v_j is adjacent to v_k is $O(1)$
- Finding all neighbors of v_j is $\Theta(|V|)$



Adjacency list


Most efficient for algorithms is an adjacency list

- Each vertex is associated with a list of its neighbors

```

1  • → 2 → 4
2  •
3  • → 5
4  • → 2 → 5
5  • → 2 → 3 → 8
6  • → 9
7  • → 9
8  • → 4
  
```

- Requires $\Theta(|V| + |E|)$ memory
- On average:
 - Determining if v_j is adjacent to v_k is
 - Finding all neighbors of v_j is



The Graph ADT

The Graph ADT describes a container storing an adjacency relation

- Queries include:
 - The number of vertices
 - The number of edges
 - List the vertices adjacent to a given vertex
 - Are two vertices adjacent?
 - Are two vertices connected?
- Modifications include:
 - Inserting or removing an edge
 - Inserting or removing a vertex (and all edges containing that vertex)

The run-time of these operations will depend on the representation

Summary

In this topic, we have covered:

- Basic graph definitions
 - Vertex, edge, degree, adjacency
- Paths, simple paths, and cycles
- Connectedness
- Weighted graphs
- Directed graphs
- Directed acyclic graphs

We will continue by looking at a number of problems related to graphs

References

Wikipedia, http://en.wikipedia.org/wiki/Adjacency_matrix
http://en.wikipedia.org/wiki/Adjacency_list

- [1] Donald E. Knuth, *The Art of Computer Programming, Volume 1: Fundamental Algorithms*, 3rd Ed., Addison Wesley, 1997, §2.2.1, p.238.
- [2] Cormen, Leiserson, and Rivest, *Introduction to Algorithms*, McGraw Hill, 1990, §11.1, p.200.
- [3] Weiss, *Data Structures and Algorithm Analysis in C++*, 3rd Ed., Addison Wesley, §3.6, p.94.
- [4] David H. Laidlaw, Course Notes, <http://cs.brown.edu/courses/cs016/lectures/13%20Graphs.pdf>

These slides are provided for the ECE 250 *Algorithms and Data Structures* course. The material in it reflects Douglas W. Harder's best judgment in light of the information available to him at the time of preparation. Any reliance on these course slides by any party for any other purpose are the responsibility of such parties. Douglas W. Harder accepts no responsibility for damages, if any, suffered by any party as a result of decisions made or actions based on these course slides for any other purpose than that for which it was intended.