

Deep Learning Part I

Instructor: Prof. Yi Fang

yfang@nyu.edu

Python tutorial: <http://learnpython.org/>

TensorFlow tutorial: <https://www.tensorflow.org/tutorials/>

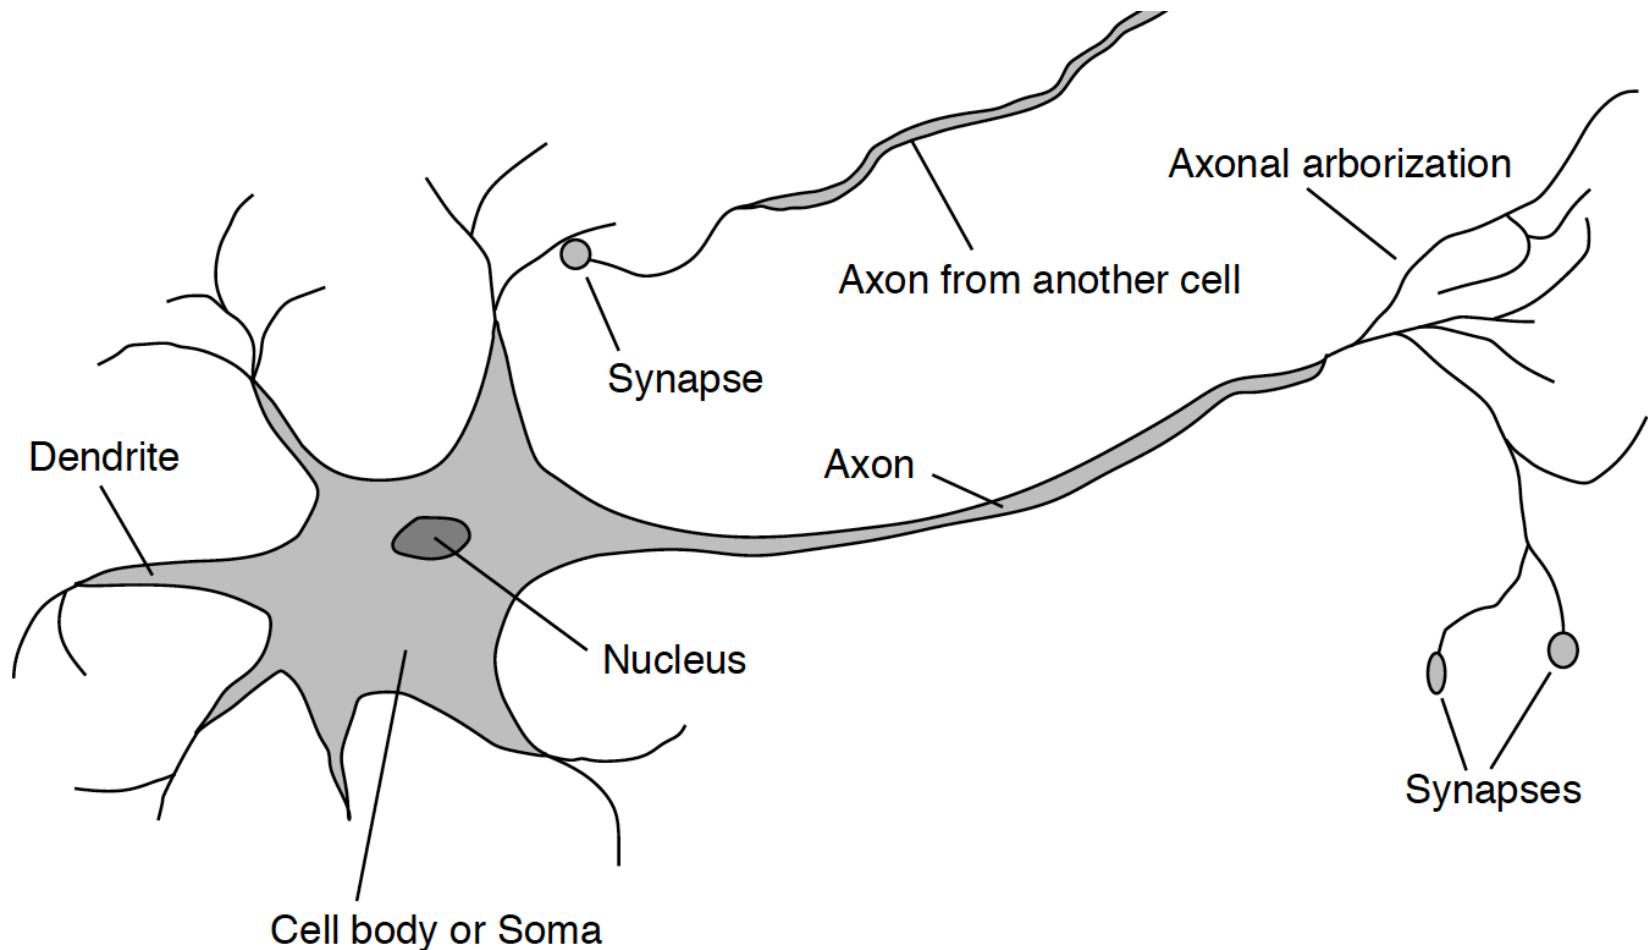
PyTorch tutorial: <https://pytorch.org/tutorials/>

Overview

- Human Brain
- Neural Network
- Perceptrons
- Multilayer Perceptrons
- Convolution Neural Network
 - Overview architecture
 - Convolution operation
 - Convolution layer
 - Pooling layer
 - Fully connection layer

Human Brain

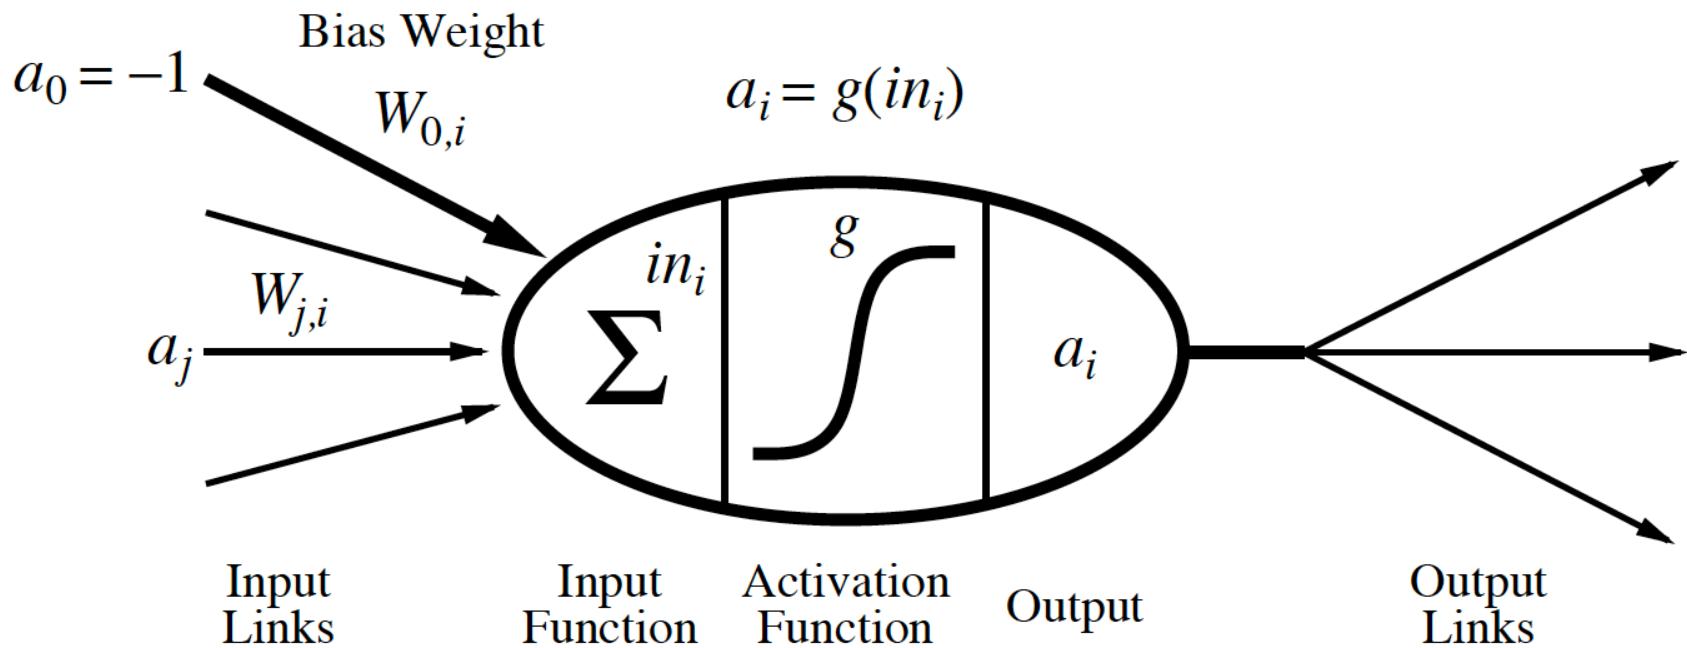
- 10^{11} neurons of > 20 types, 10^{14} synapses, 1ms-10ms cycle time
- Signals are noisy “spike trains” of electrical potential



Neurons

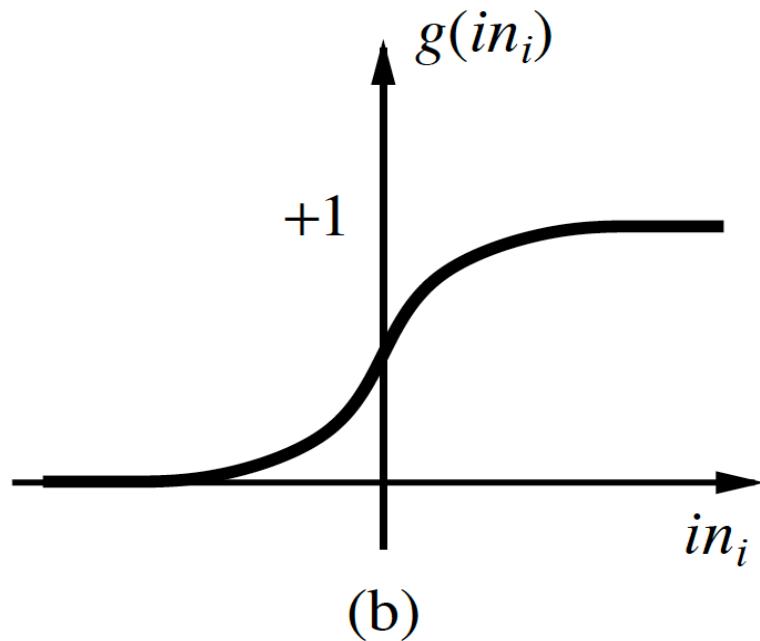
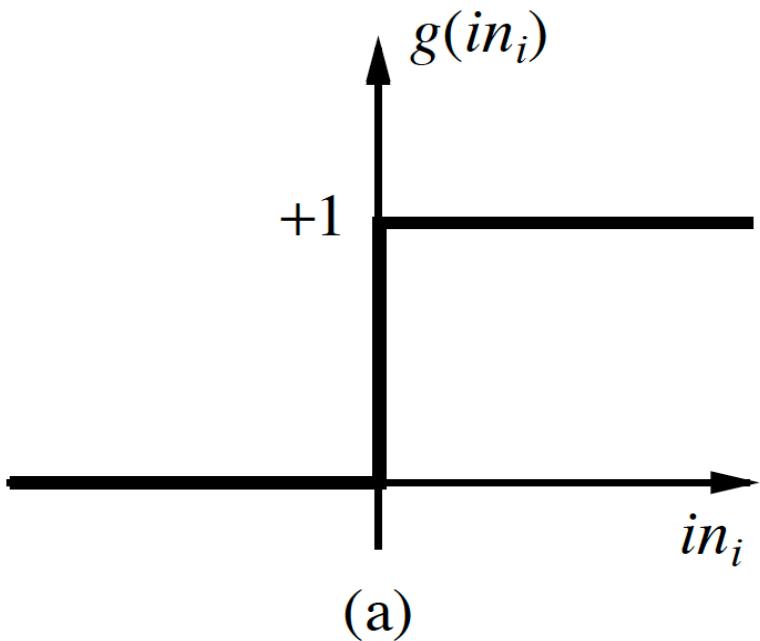
- A gross oversimplification of real neurons

$$a_i \leftarrow g(in_i) = g\left(\sum_j W_{j,i} a_j\right)$$



Human Brain

- Activation function



(a) is a step function or threshold function

(b) is a sigmoid function $1/(1 + e^{-x})$

Changing the bias weight $W_{0,i}$ moves the threshold location

Network structure

- Feed-forward Networks
 - Single-layer perceptrons
 - Multi-layer perceptrons
- Recurrent Network
 - Hopfield networks
 - Boltzmann machines

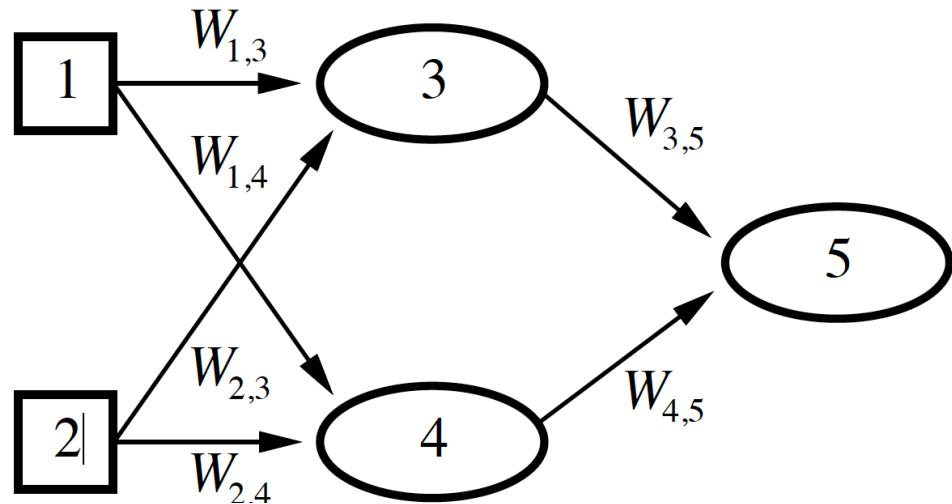
Feed-forward networks implement functions which have no internal state

Recurrent neural nets have directed cycles with delays and have internal state (like ip-ops)

Feed-forward Networks

- Feed-forward network = a parameterized family of nonlinear functions
- Adjusting weights changes the function

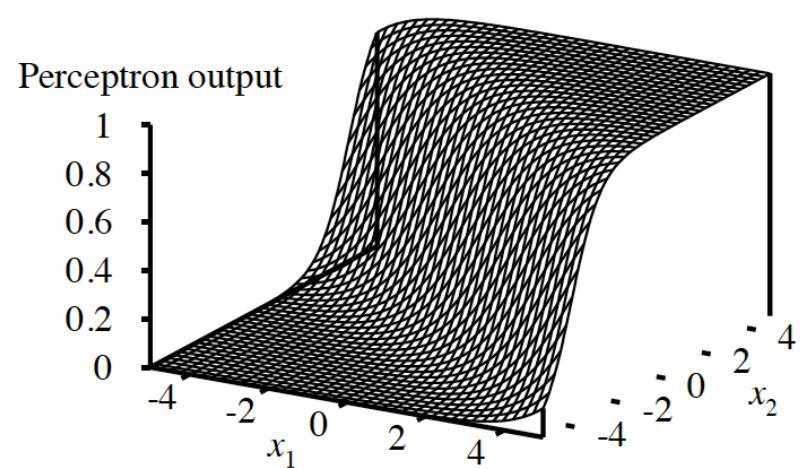
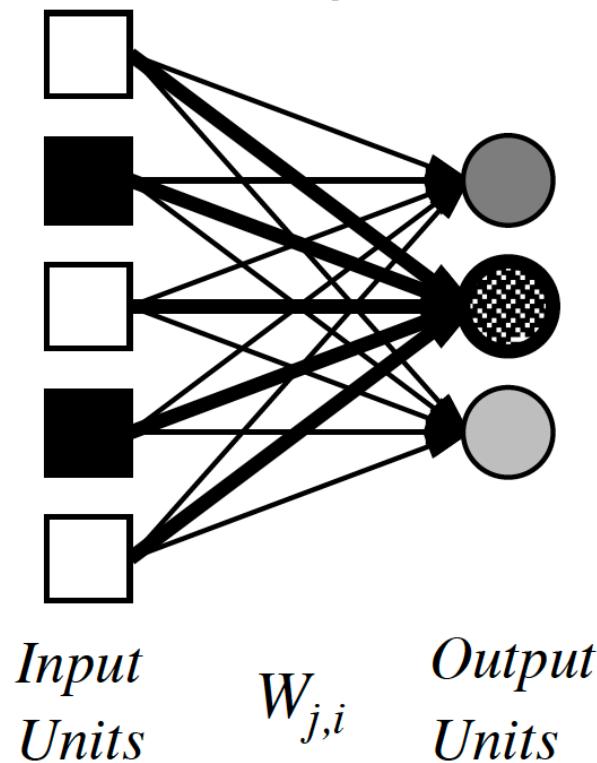
An easy example:



$$\begin{aligned}a_5 &= g(W_{3,5} \cdot a_3 + W_{4,5} \cdot a_4) \\&= g(W_{3,5} \cdot g(W_{1,3} \cdot a_1 + W_{2,3} \cdot a_2) + W_{4,5} \cdot g(W_{1,4} \cdot a_1 + W_{2,4} \cdot a_2))\end{aligned}$$

Single-layer perceptrons

- Output units all operate separately and no shared weights
- Adjusting weights moves the location, orientation, and steepness of cliff



Perceptron learning

Learn by adjusting weights to reduce **error** on training set

The **squared error** for an example with input \mathbf{x} and true output y is

$$E = \frac{1}{2} Err^2 \equiv \frac{1}{2} (y - h_{\mathbf{W}}(\mathbf{x}))^2 ,$$

Perform optimization search by gradient descent:

$$\begin{aligned}\frac{\partial E}{\partial W_j} &= Err \times \frac{\partial Err}{\partial W_j} = Err \times \frac{\partial}{\partial W_j} (y - g(\sum_{j=0}^n W_j x_j)) \\ &= -Err \times g'(in) \times x_j\end{aligned}$$

Simple weight update rule:

$$W_j \leftarrow W_j + \alpha \times Err \times g'(in) \times x_j$$

Multilayer perceptrons

- Layers are usually fully connected
- Numbers of hidden units typically chosen by hand

Output units

a_i

$W_{j,i}$

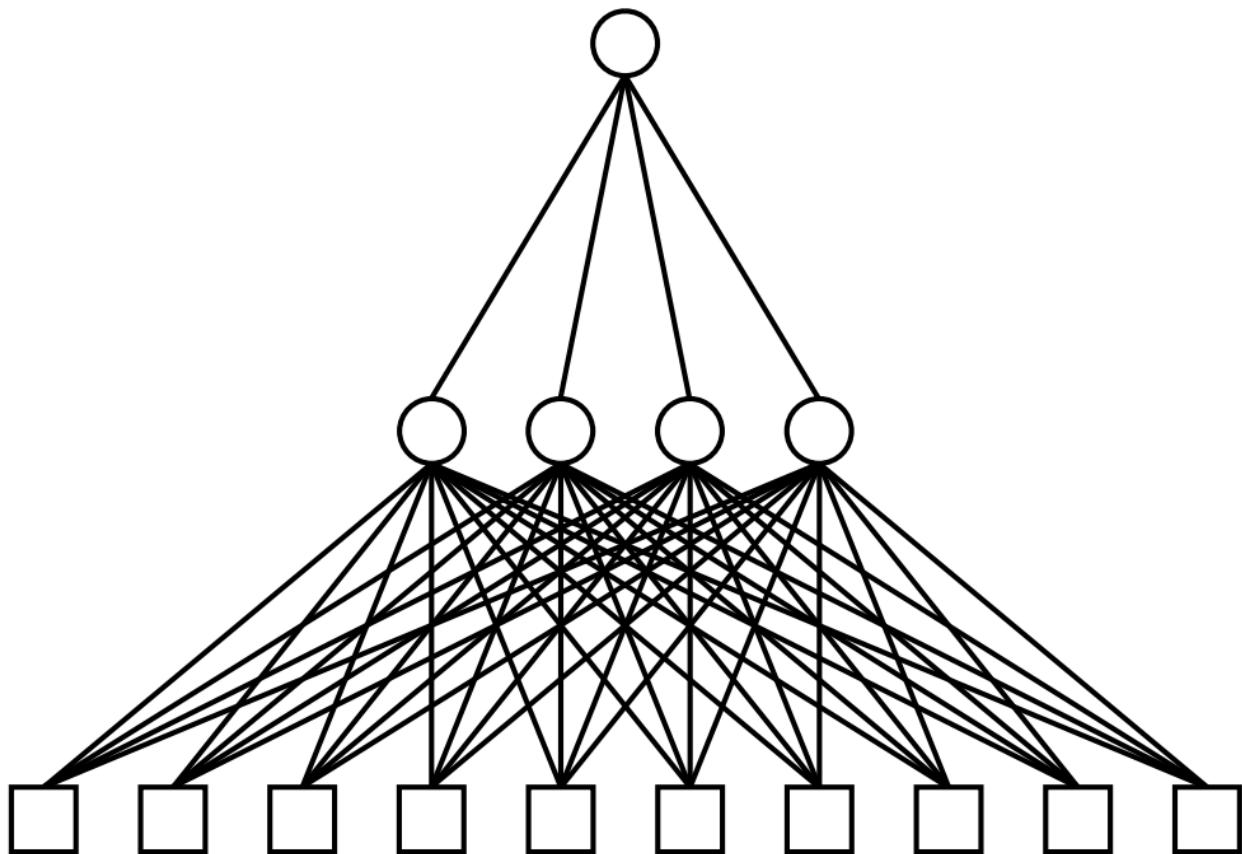
Hidden units

a_j

$W_{k,j}$

Input units

a_k



Back-propagation learning

Output layer: same as for single-layer perceptron,

$$W_{j,i} \leftarrow W_{j,i} + \alpha \times a_j \times \Delta_i$$

where $\Delta_i = Err_i \times g'(in_i)$

Hidden layer: **back-propagate** the error from the output layer:

$$\Delta_j = g'(in_j) \sum_i W_{j,i} \Delta_i .$$

Update rule for weights in hidden layer:

$$W_{k,j} \leftarrow W_{k,j} + \alpha \times a_k \times \Delta_j .$$

Back-propagation example

- Loss is defined as:

$$E = \frac{1}{2} \sum_i (y_i - a_i)^2 ,$$

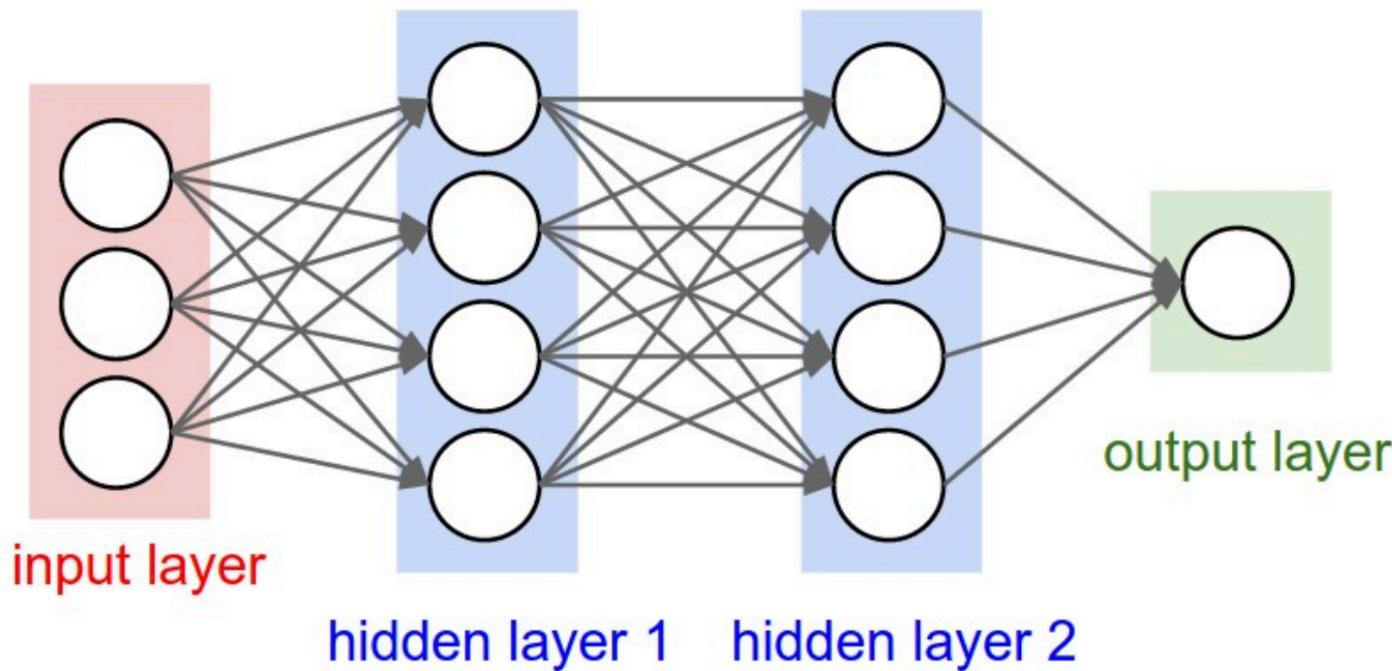
- where the sum is over the nodes in the output layer

$$\begin{aligned}\frac{\partial E}{\partial W_{j,i}} &= -(y_i - a_i) \frac{\partial a_i}{\partial W_{j,i}} = -(y_i - a_i) \frac{\partial g(in_i)}{\partial W_{j,i}} \\ &= -(y_i - a_i) g'(in_i) \frac{\partial in_i}{\partial W_{j,i}} = -(y_i - a_i) g'(in_i) \frac{\partial}{\partial W_{j,i}} \left(\sum_j W_{j,i} a_j \right) \\ &= -(y_i - a_i) g'(in_i) a_j = -a_j \Delta_i\end{aligned}$$

Convolution Neural Network (CNN)

Why CNN?

- A regular 3-layer Neural Network.

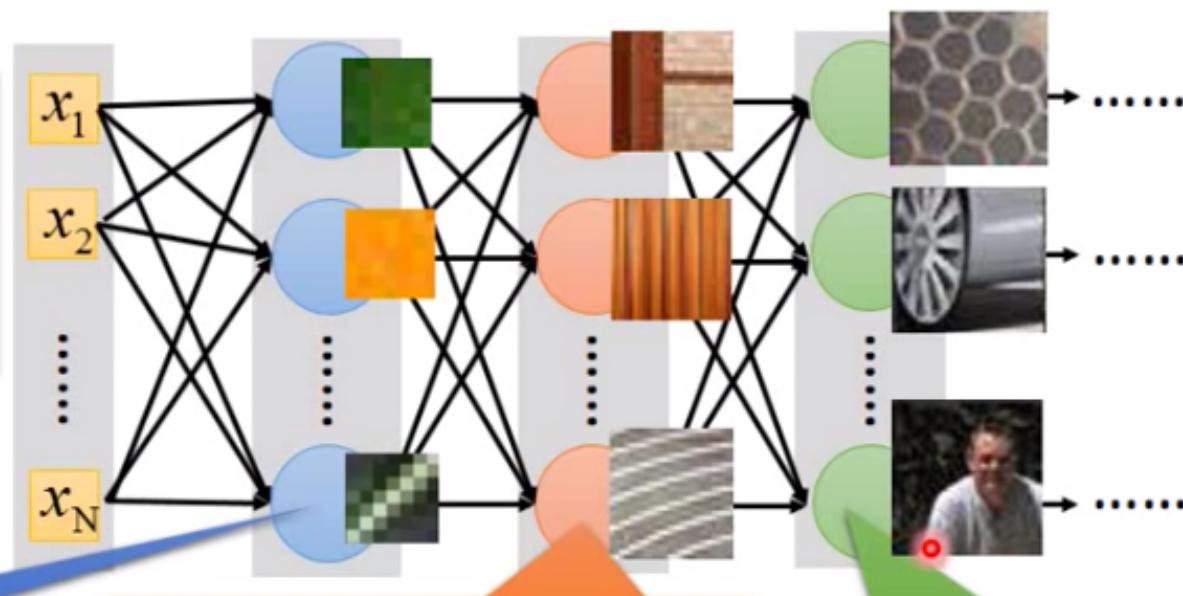


Why CNN for Image?

[Zeiler, M. D., ECCV 2014]



Represented
as pixels



The most basic
classifiers

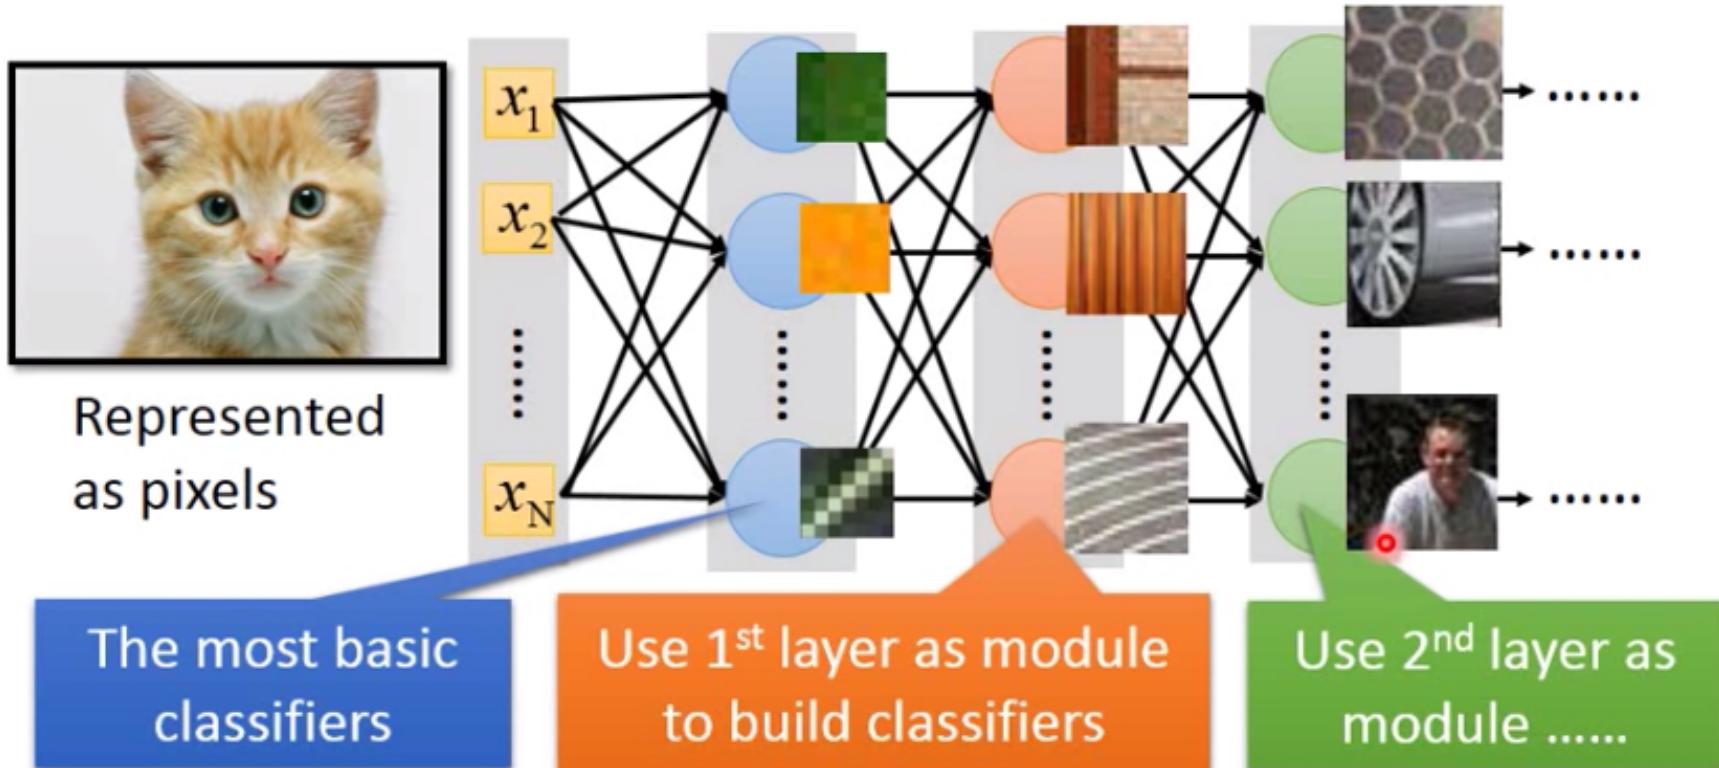
Use 1st layer as module
to build classifiers

Use 2nd layer as
module

Redundant Parameters?

Why CNN for Image?

[Zeiler, M. D., ECCV 2014]



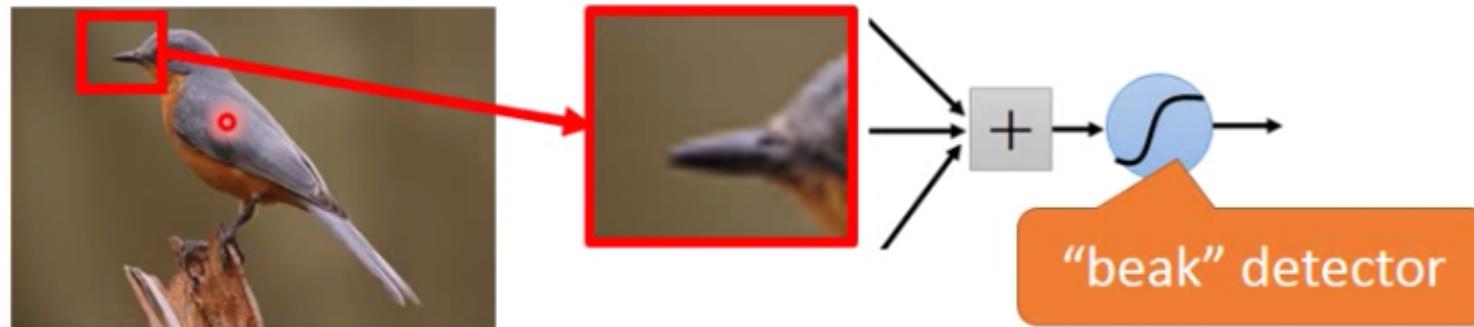
Can the Network be simplified?

- Focus on local region
- Insensitive to location
- Sub-sample operator

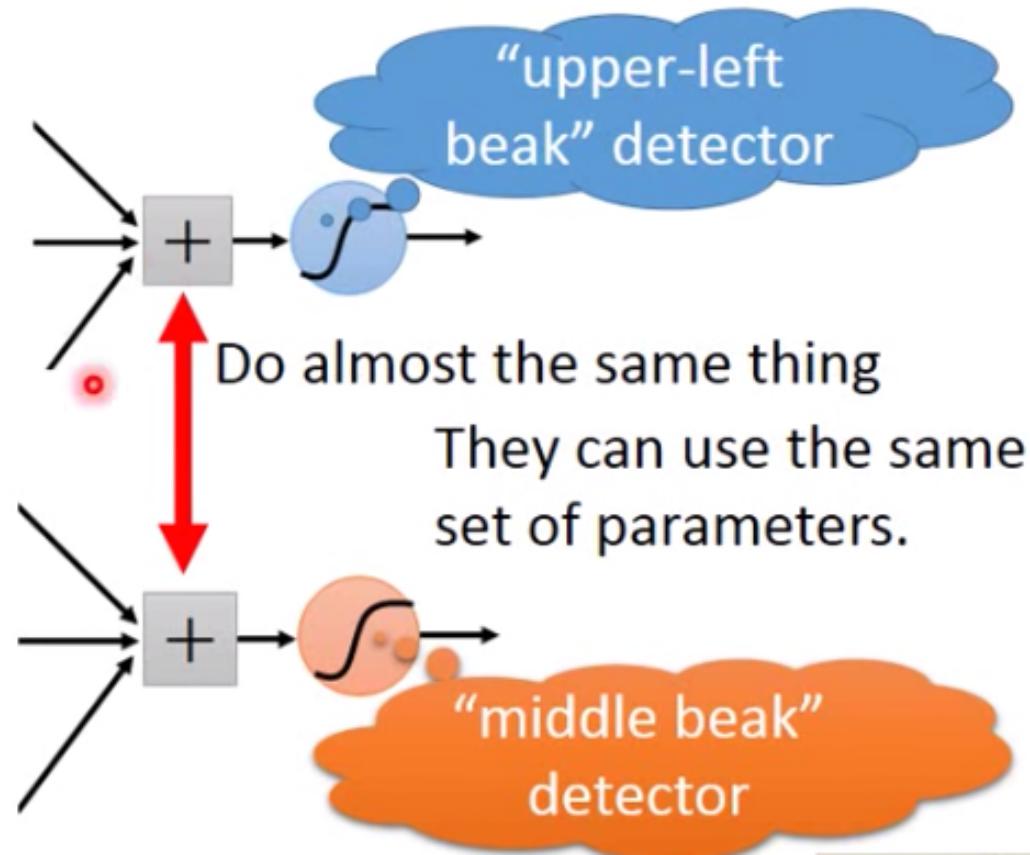
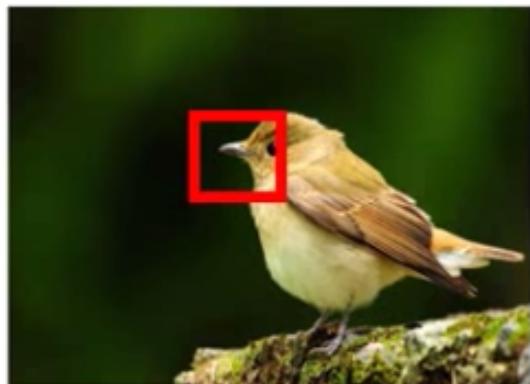
Focus on local region

A neuron does not have to see the whole image to discover the pattern.

Connecting to small region with less parameters



Inensitive to location



Sub-sampling

- Subsampling the pixels will not change the object



We can subsample the pixels to make image smaller

→ Less parameters for the network to process the image

The whole CNN



Property 1

- Some patterns are much smaller than the whole image

Property 2

- The same patterns appear in different regions.

Property 3

- Subsampling the pixels will not change the object

Convolution

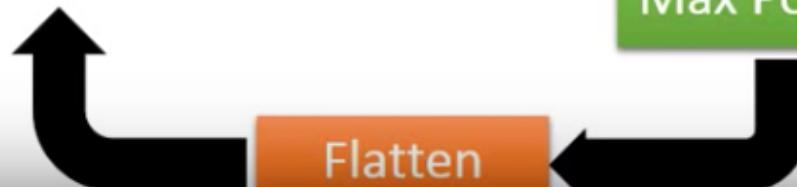
Max Pooling

Convolution

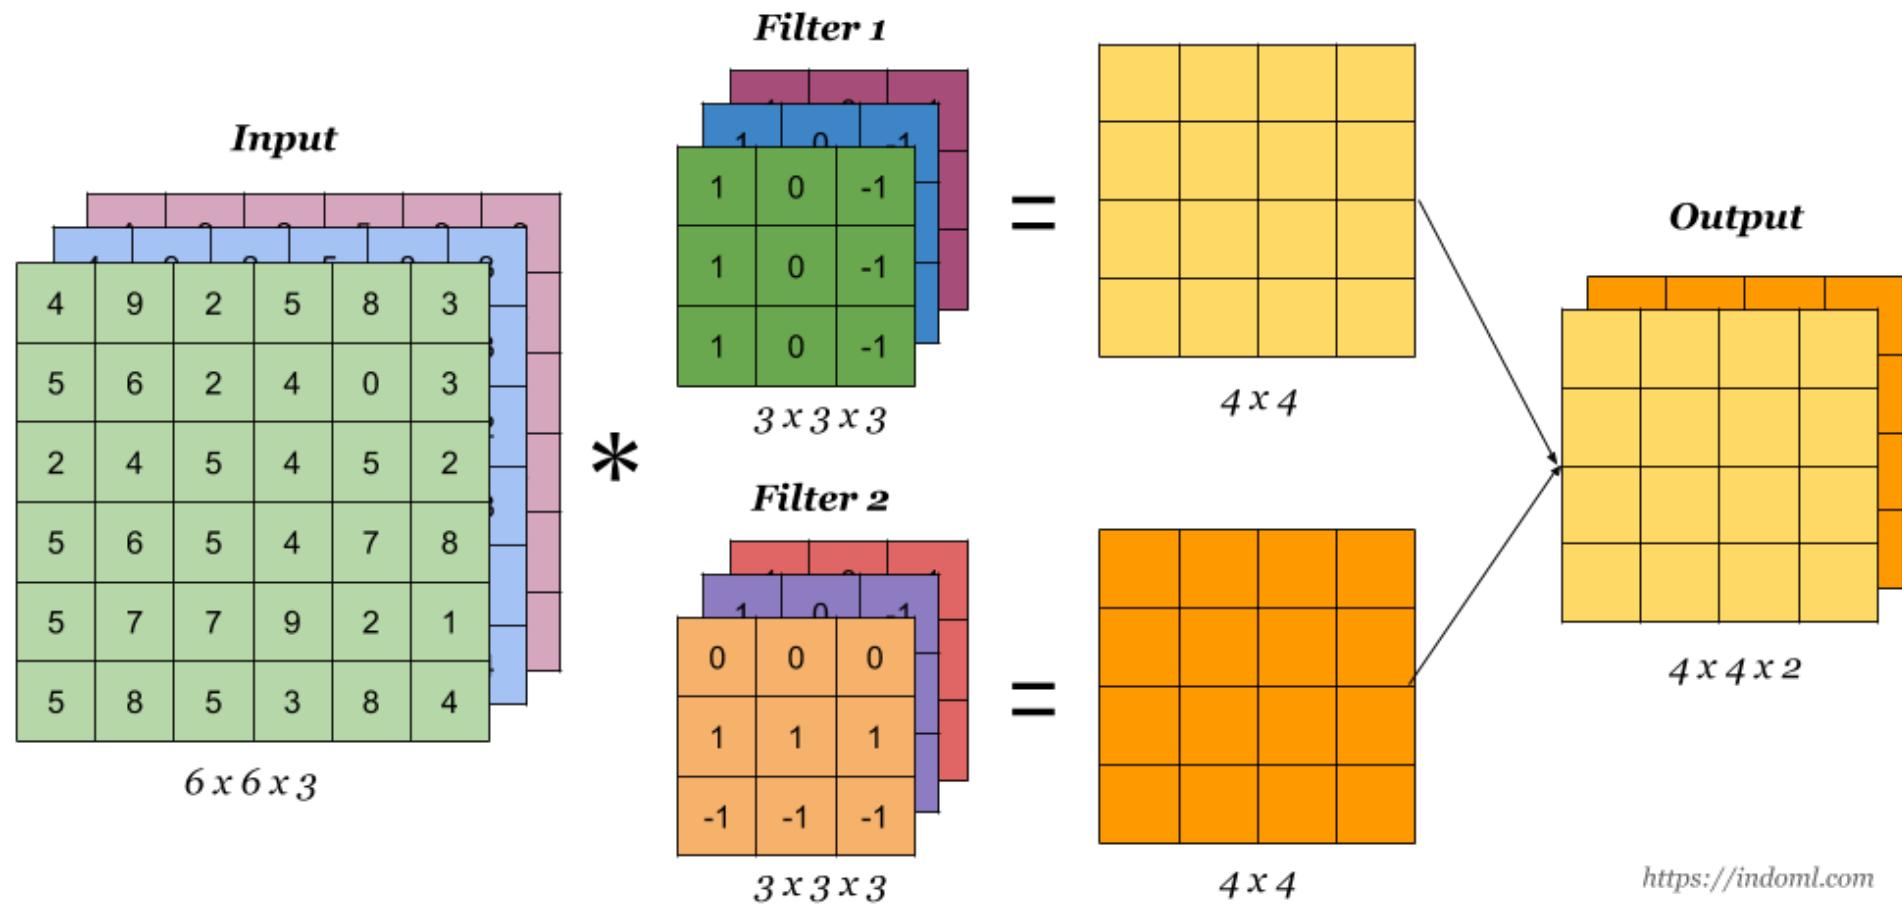
Max Pooling

Flatten

Can repeat many times



Convolution Neural Network



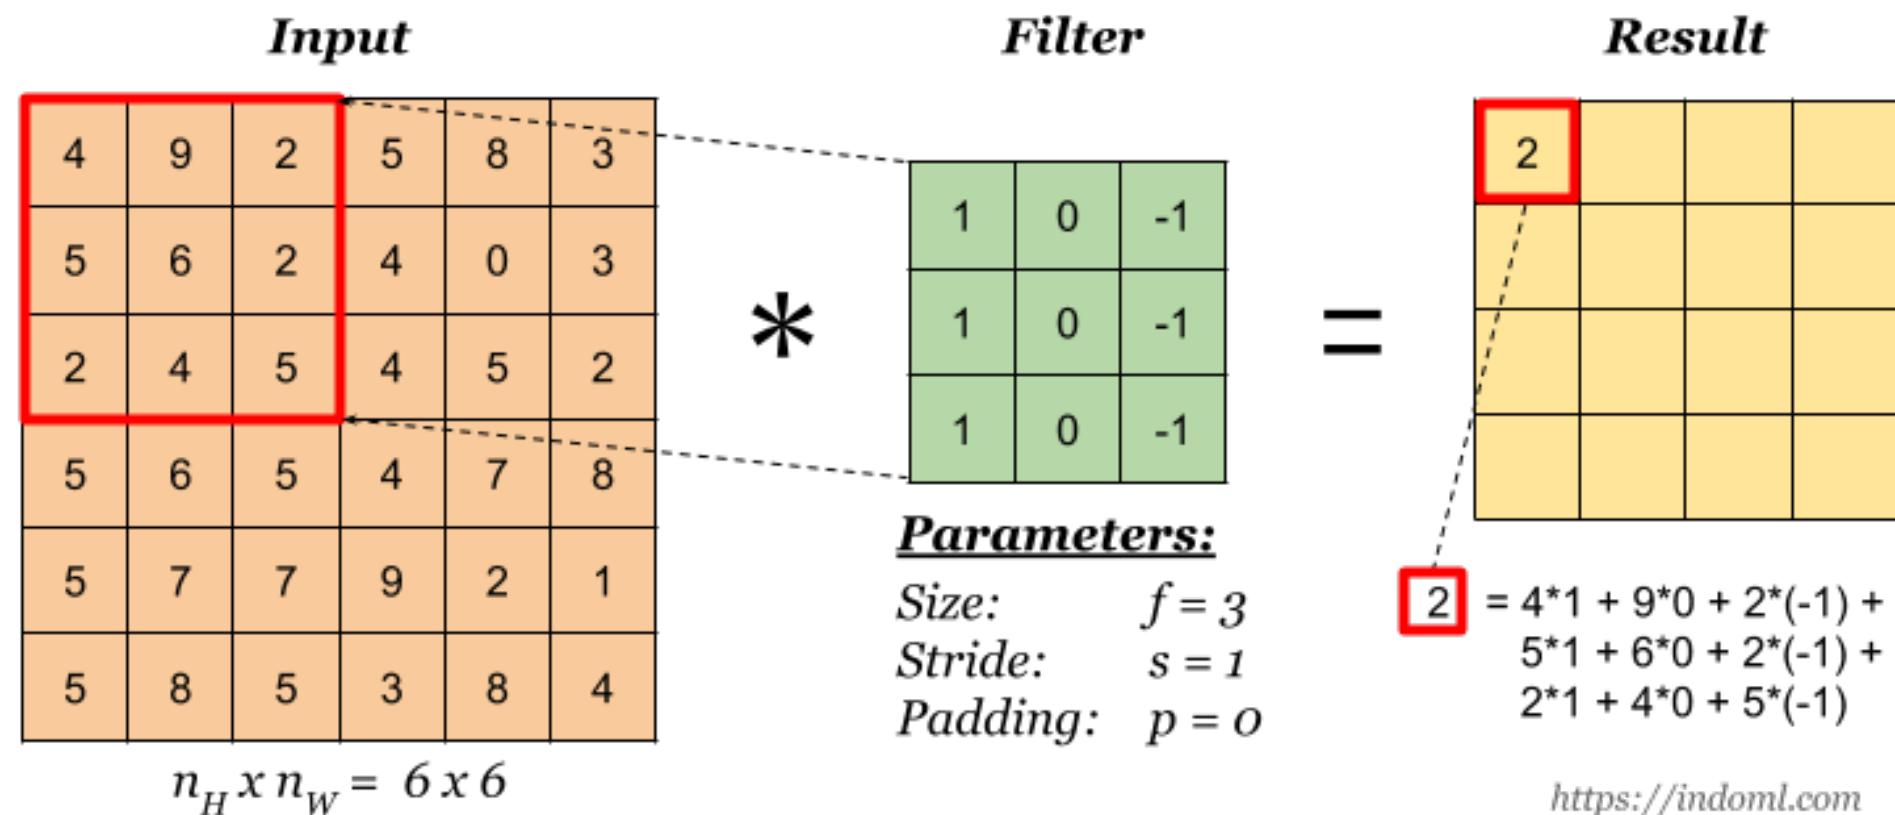
<https://indoml.com>

CNN Architectures

- Basics of CNN
- Classic Networks

CNN Basics

Basic Convolution Operation



$$n_H \times n_W = 6 \times 6$$

Input

4	9	2	5	8	3
6	2	4	0	3	
2	4	5	4	5	2
5	6	5	4	7	8
5	7	7	9	2	1
5	8	5	3	8	4

$$n_H \times n_W = 6 \times 6$$

Filter

1	0	-1
1	0	-1
1	0	-1

*

Result

2	6		

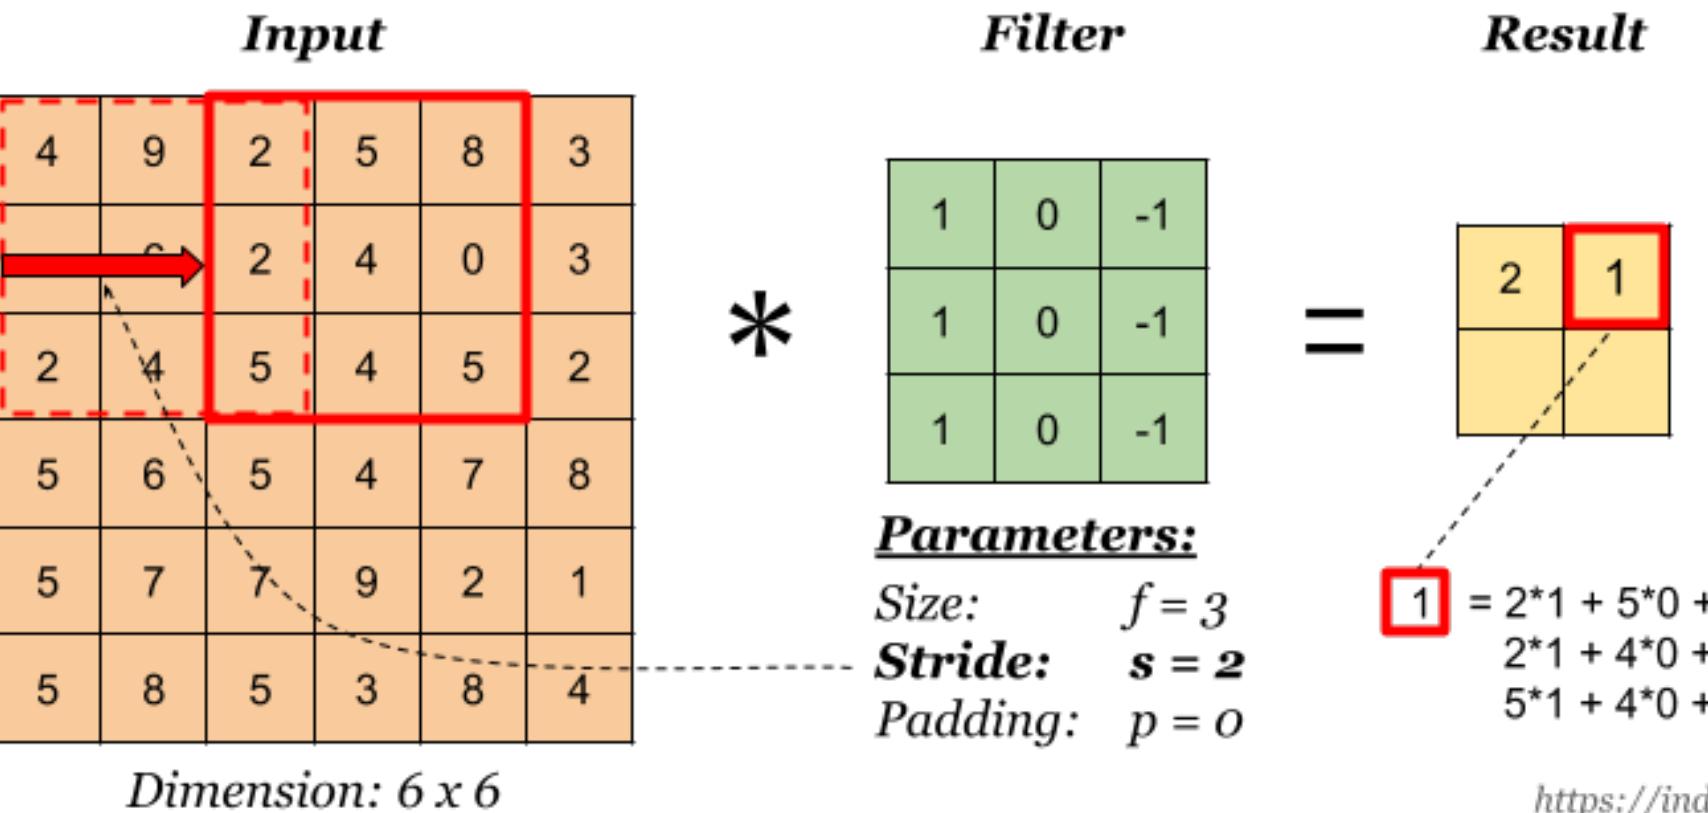
=

Parameters:Size: $f = 3$ Stride: $s = 1$ Padding: $p = 0$

$\boxed{6} = 9*1 + 2*0 + 5*(-1) +$
 $6*1 + 2*0 + 4*(-1) +$
 $4*1 + 5*0 + 4*(-1)$

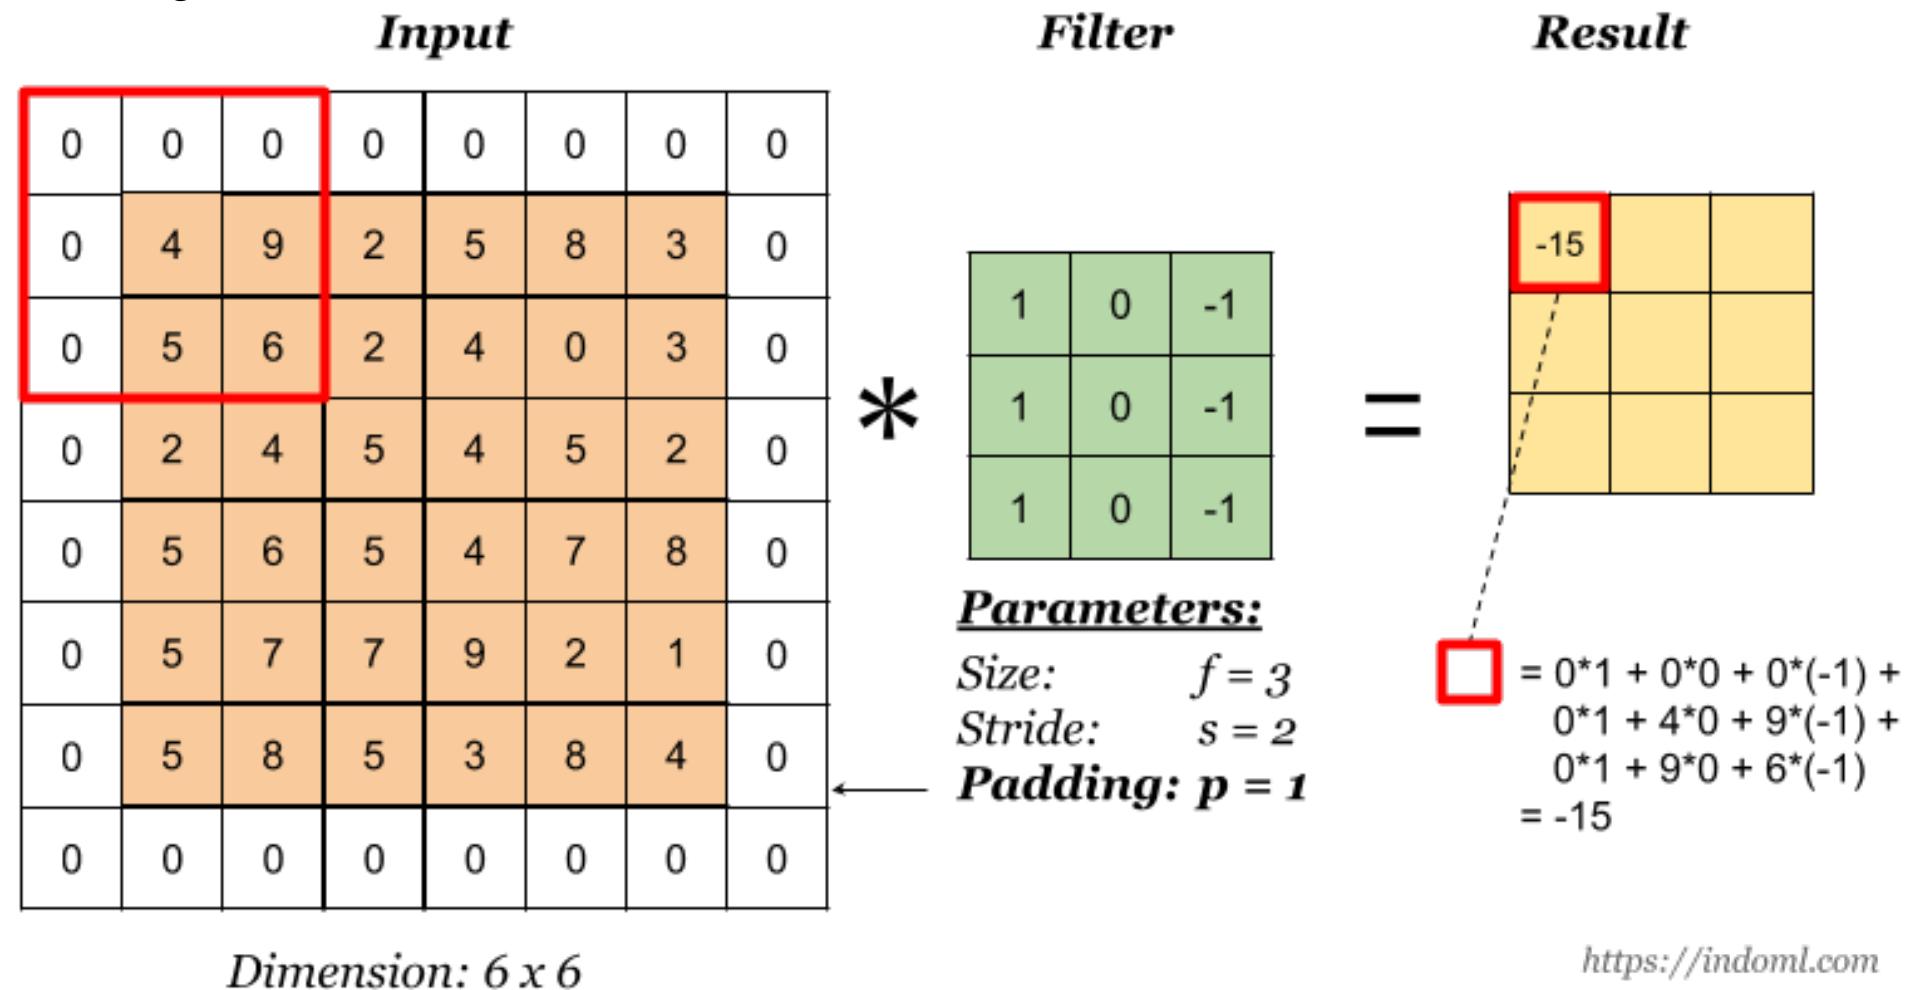
<https://indoml.com>

Stride:

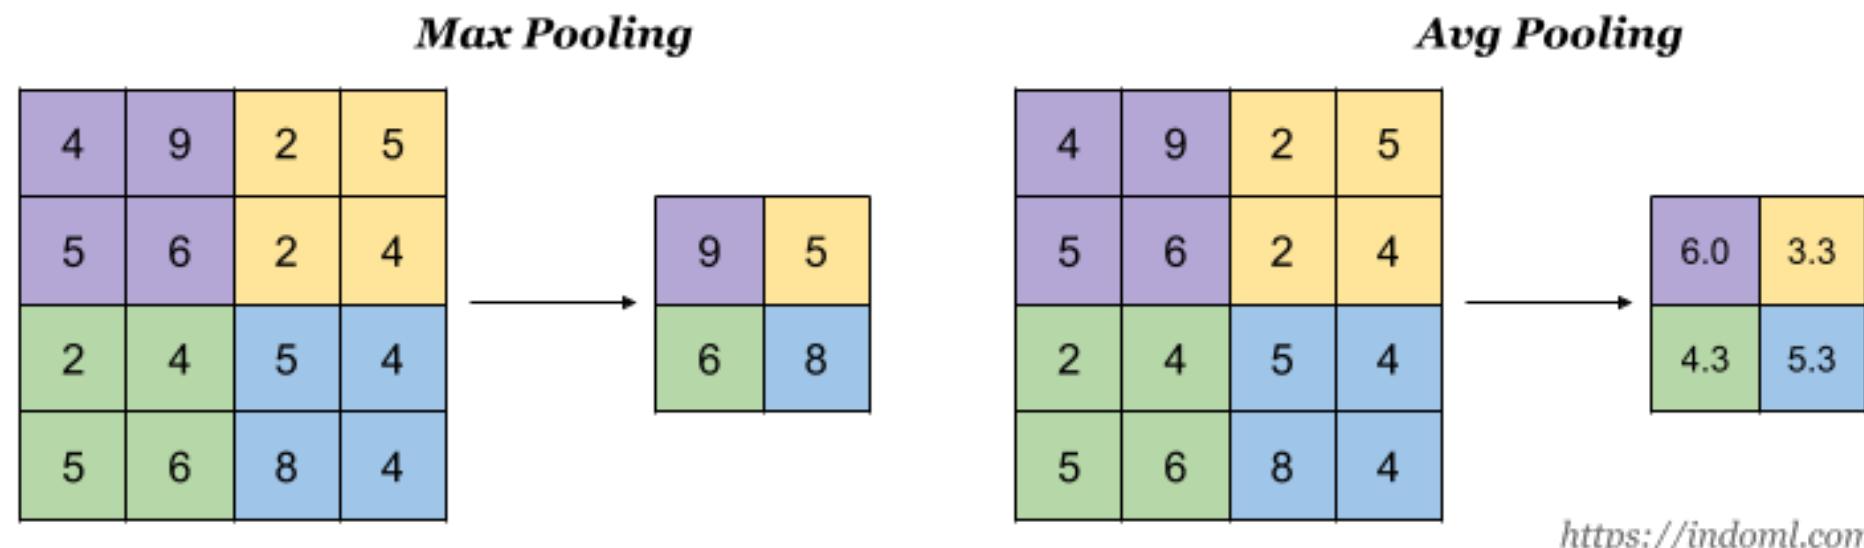


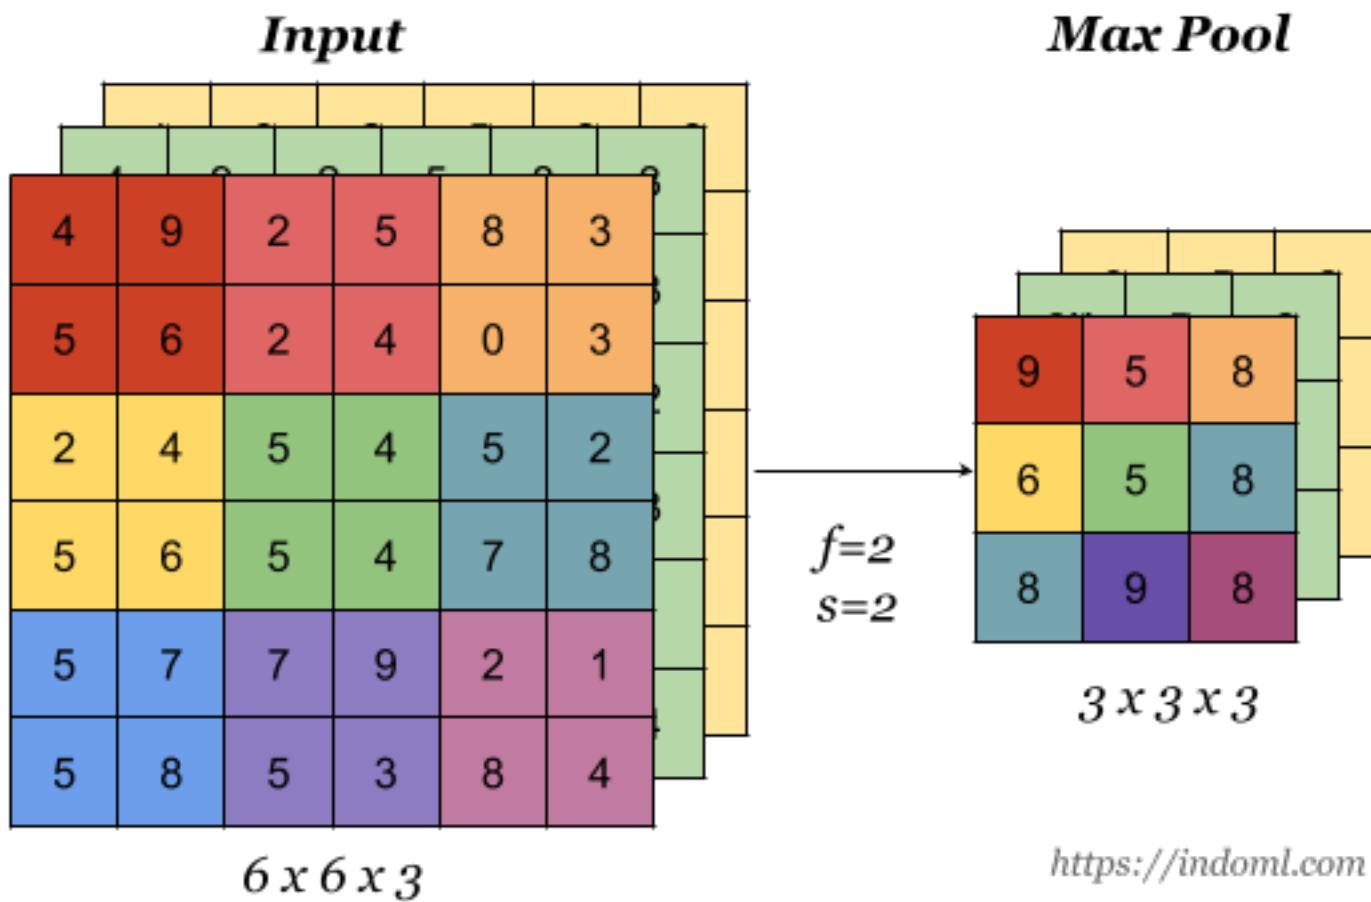
<https://indoml.com>

Padding:



Pooling layer:





Softmax layer:

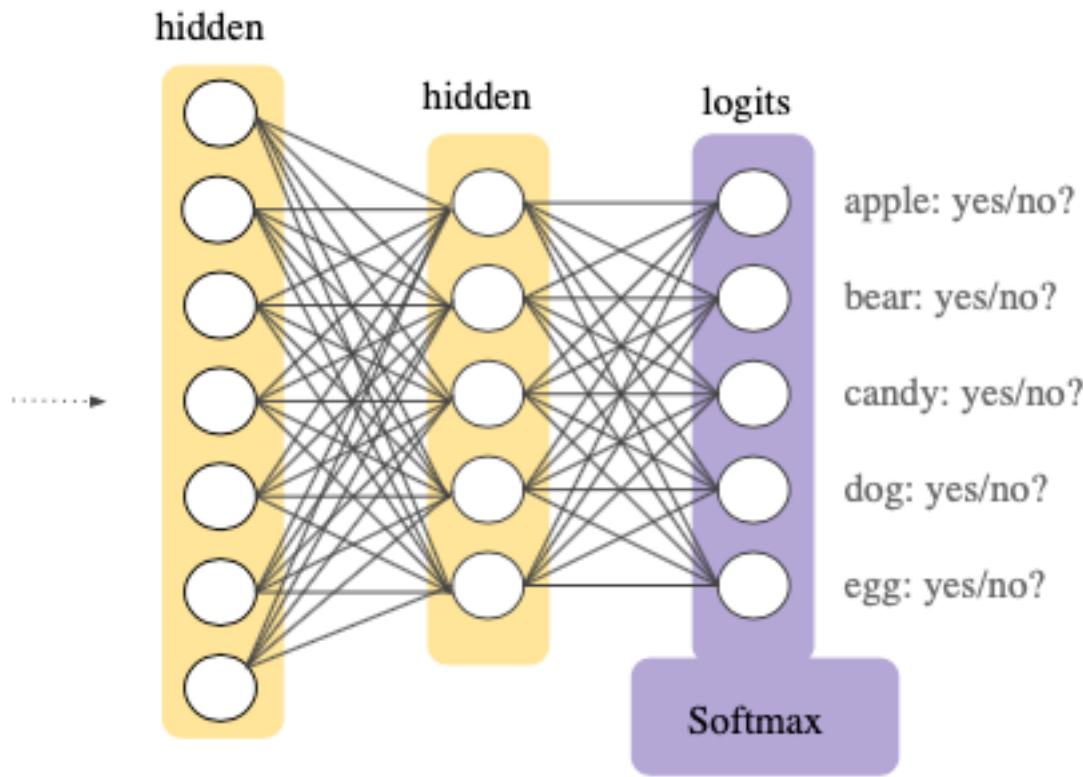
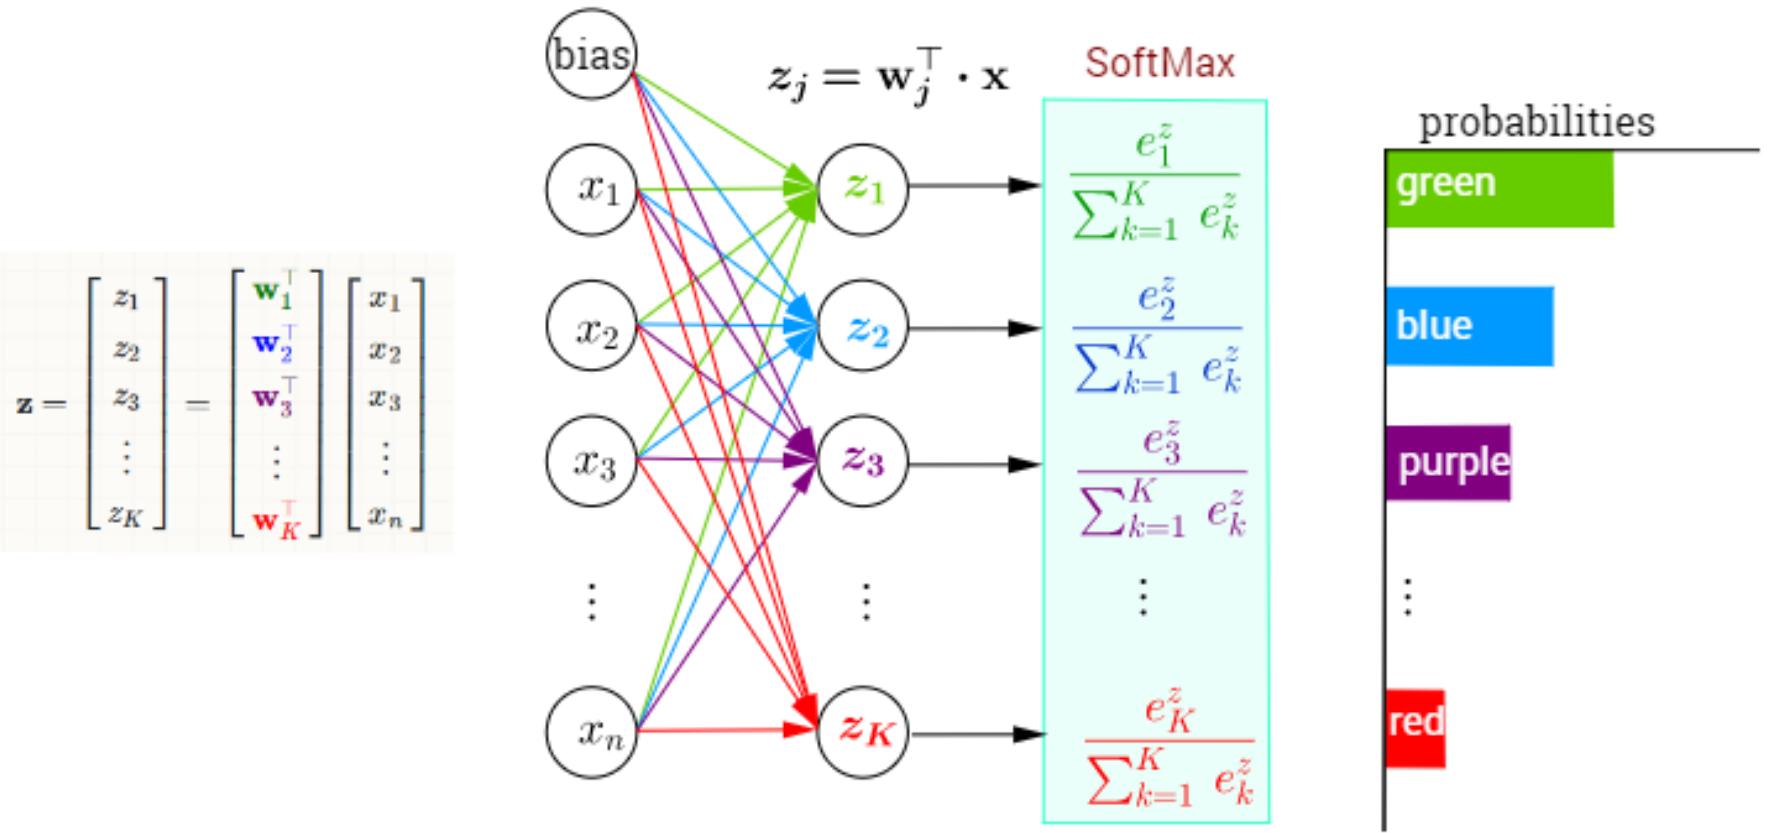
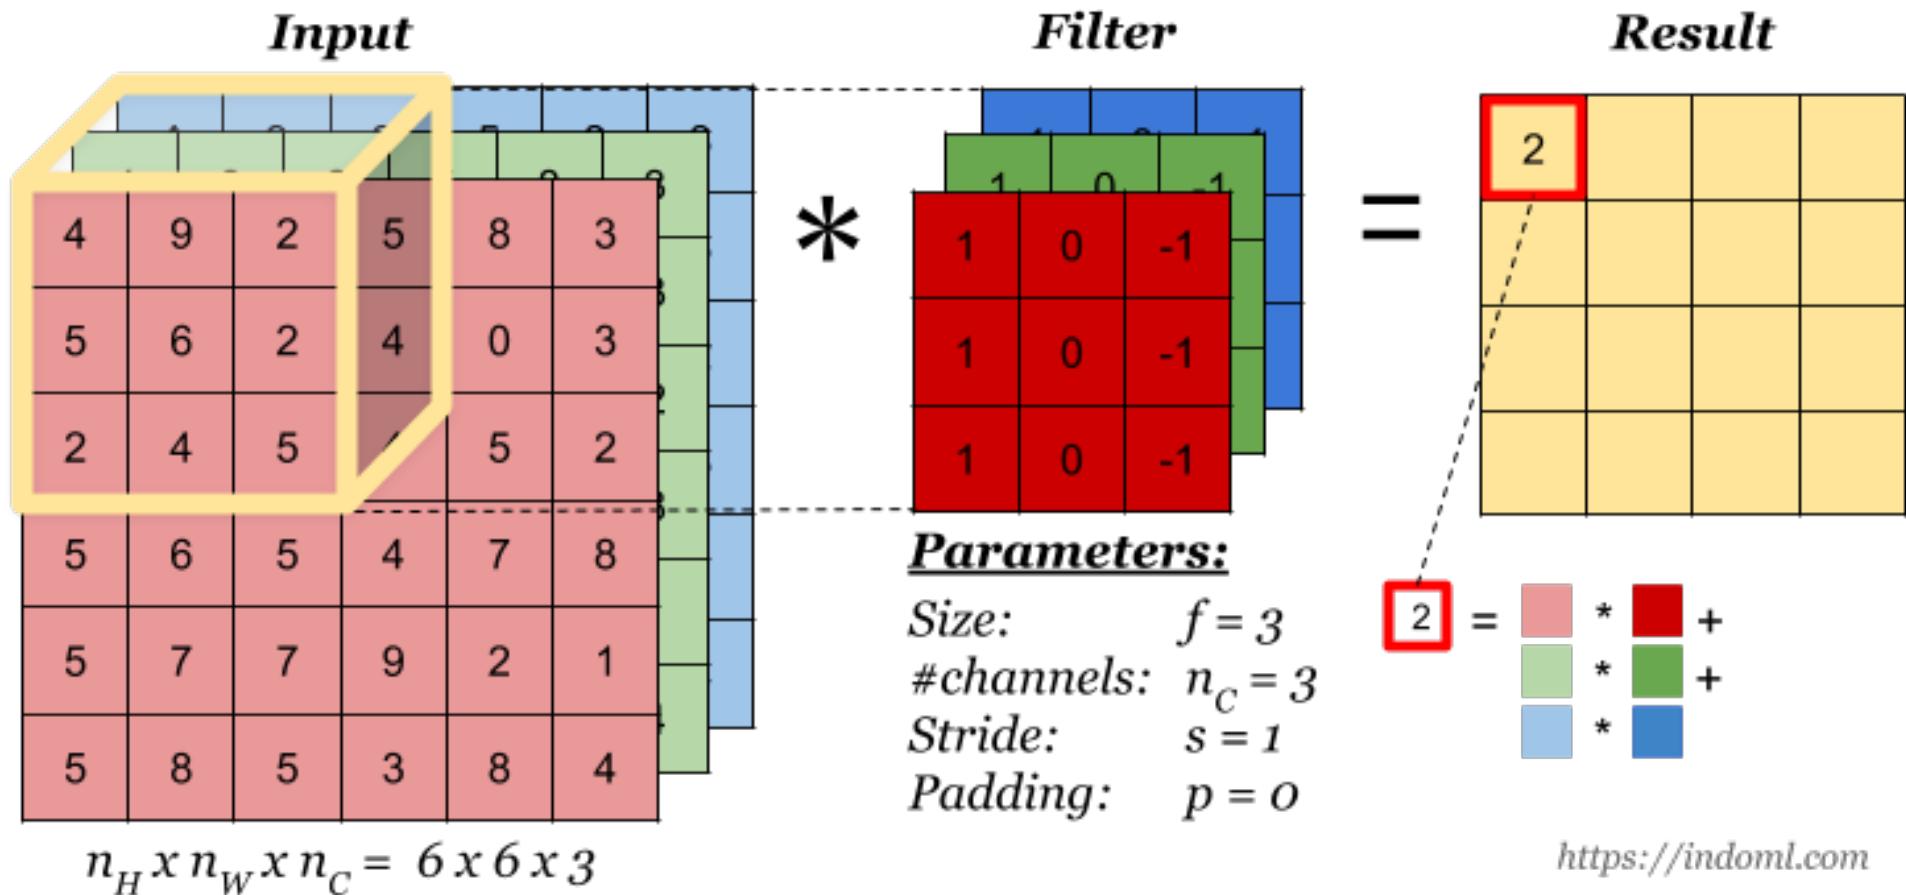


Image source: <https://developers.google.com/machine-learning/crash-course/multi-class-neural-networks/softmax>



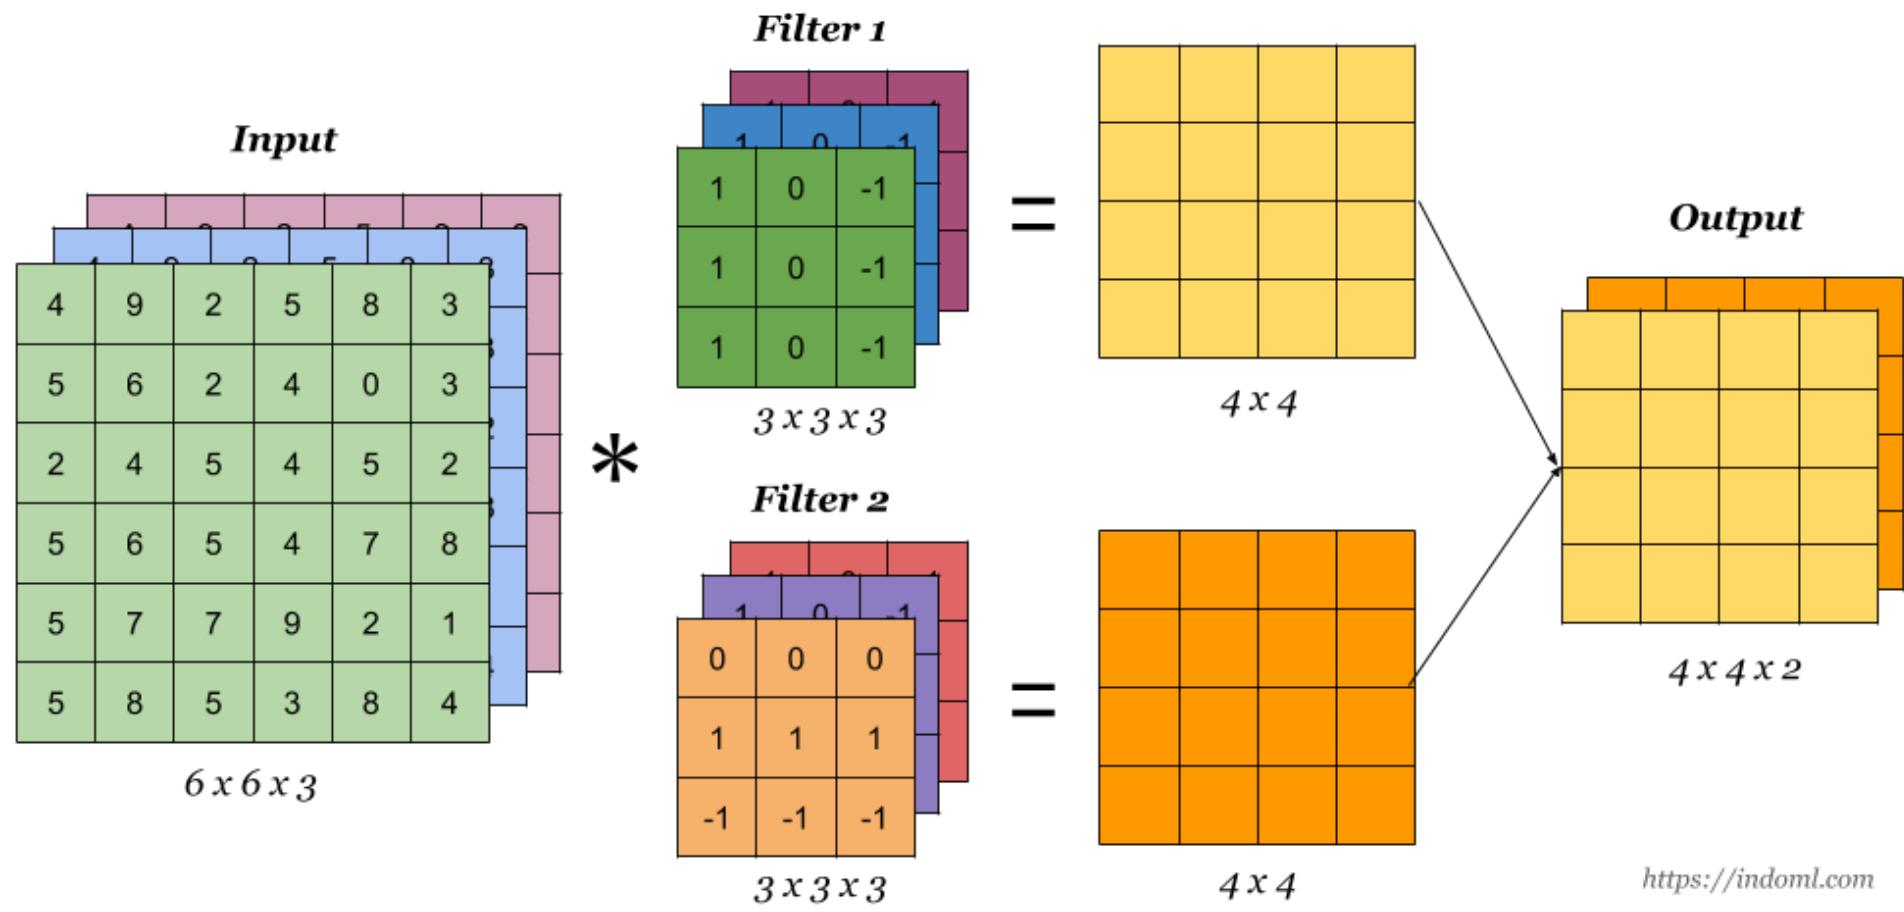
Source: <https://stats.stackexchange.com/questions/273465/neural-network-softmax-activation>

CNN on Volume (Tensors)



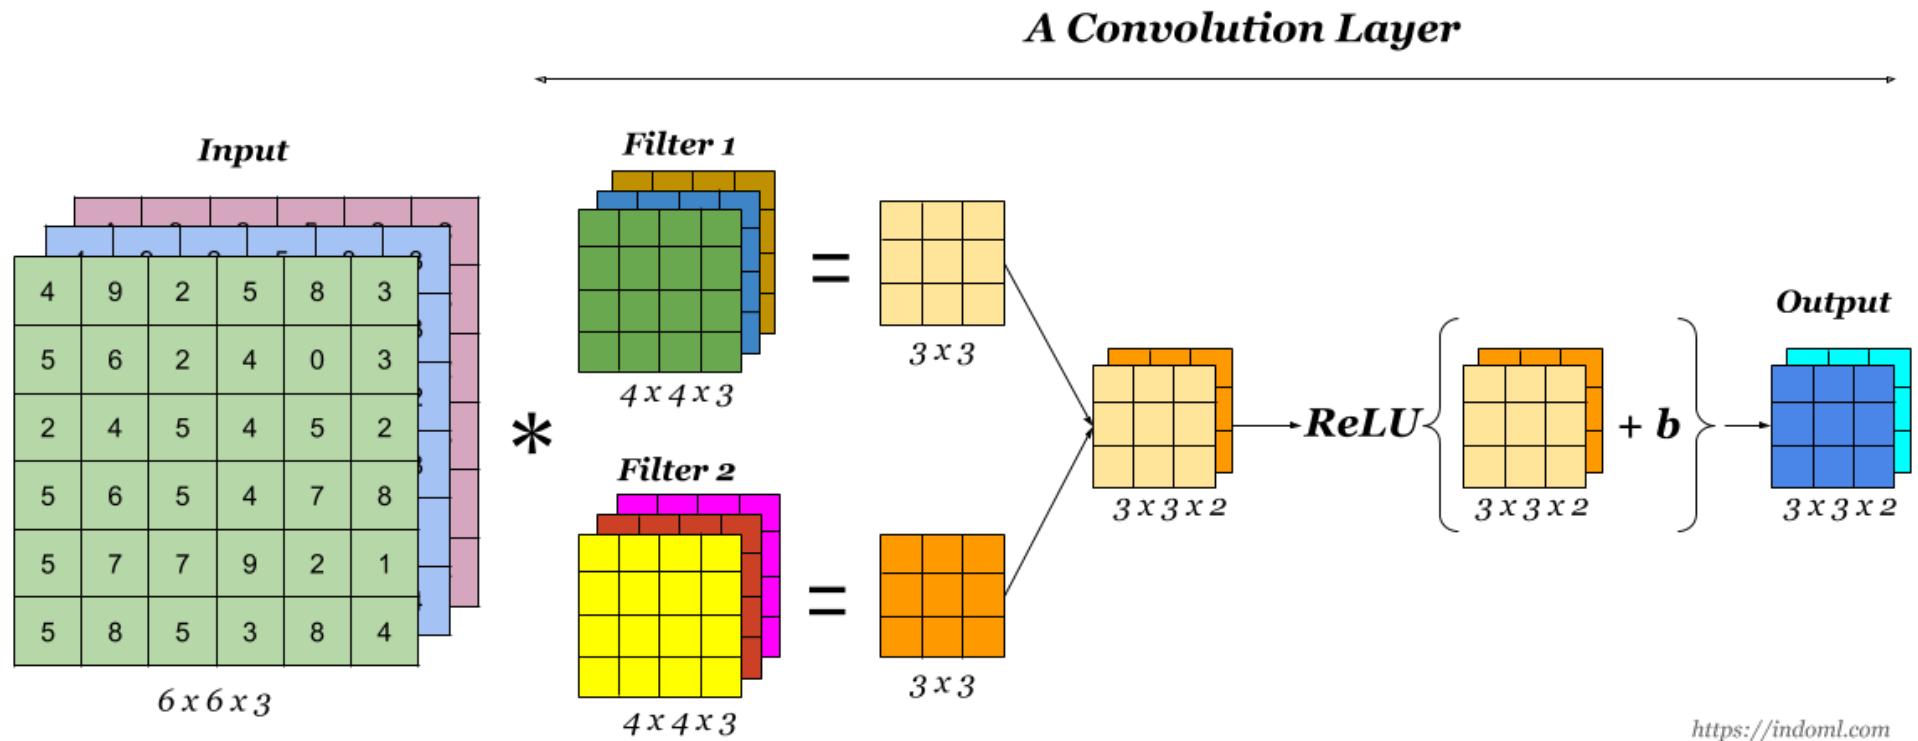
<https://indoml.com>

Multiple Filters

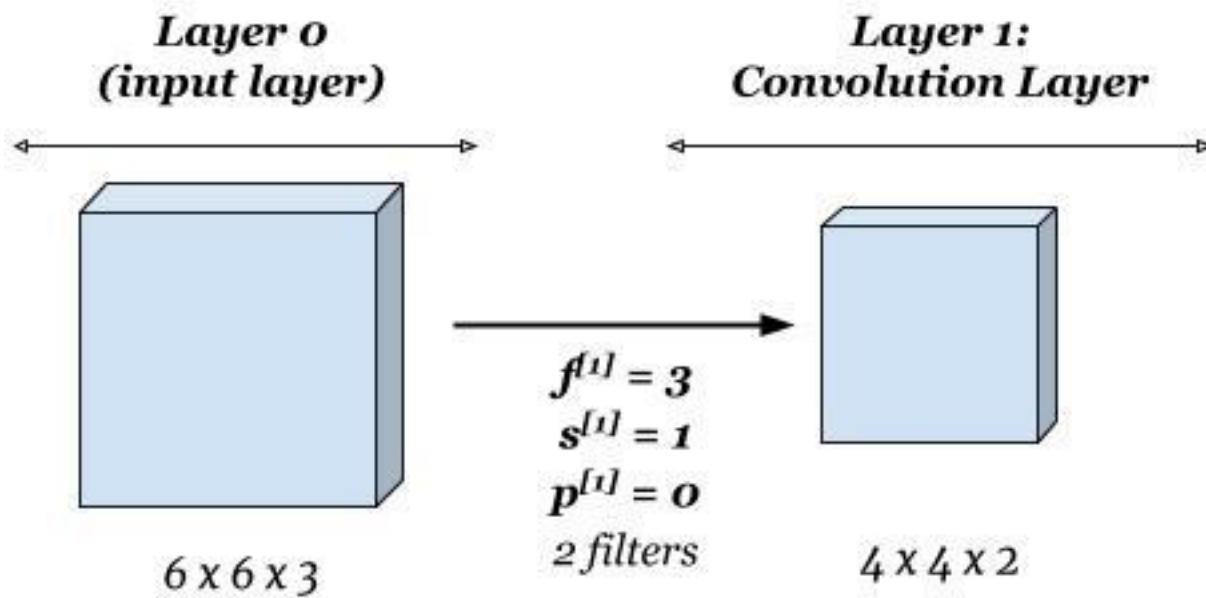


<https://indoml.com>

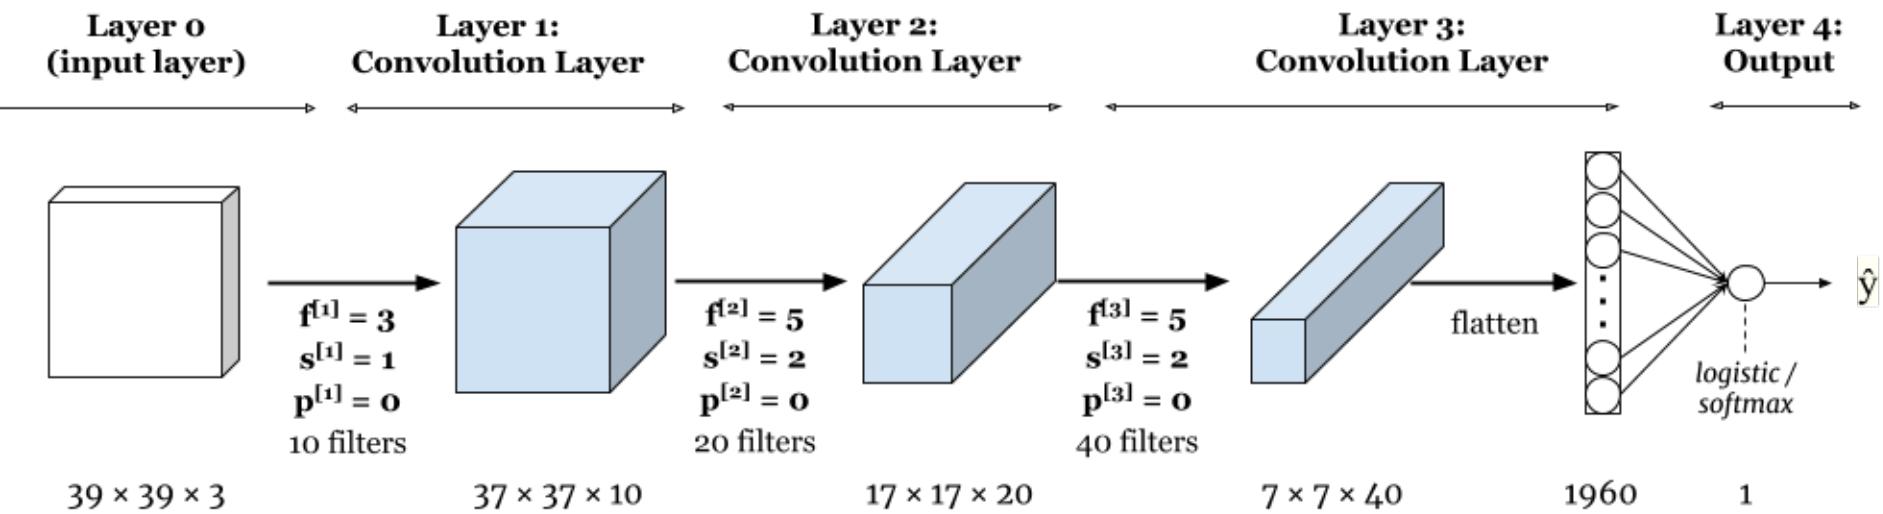
One Layer Representation



Shorthand Representation

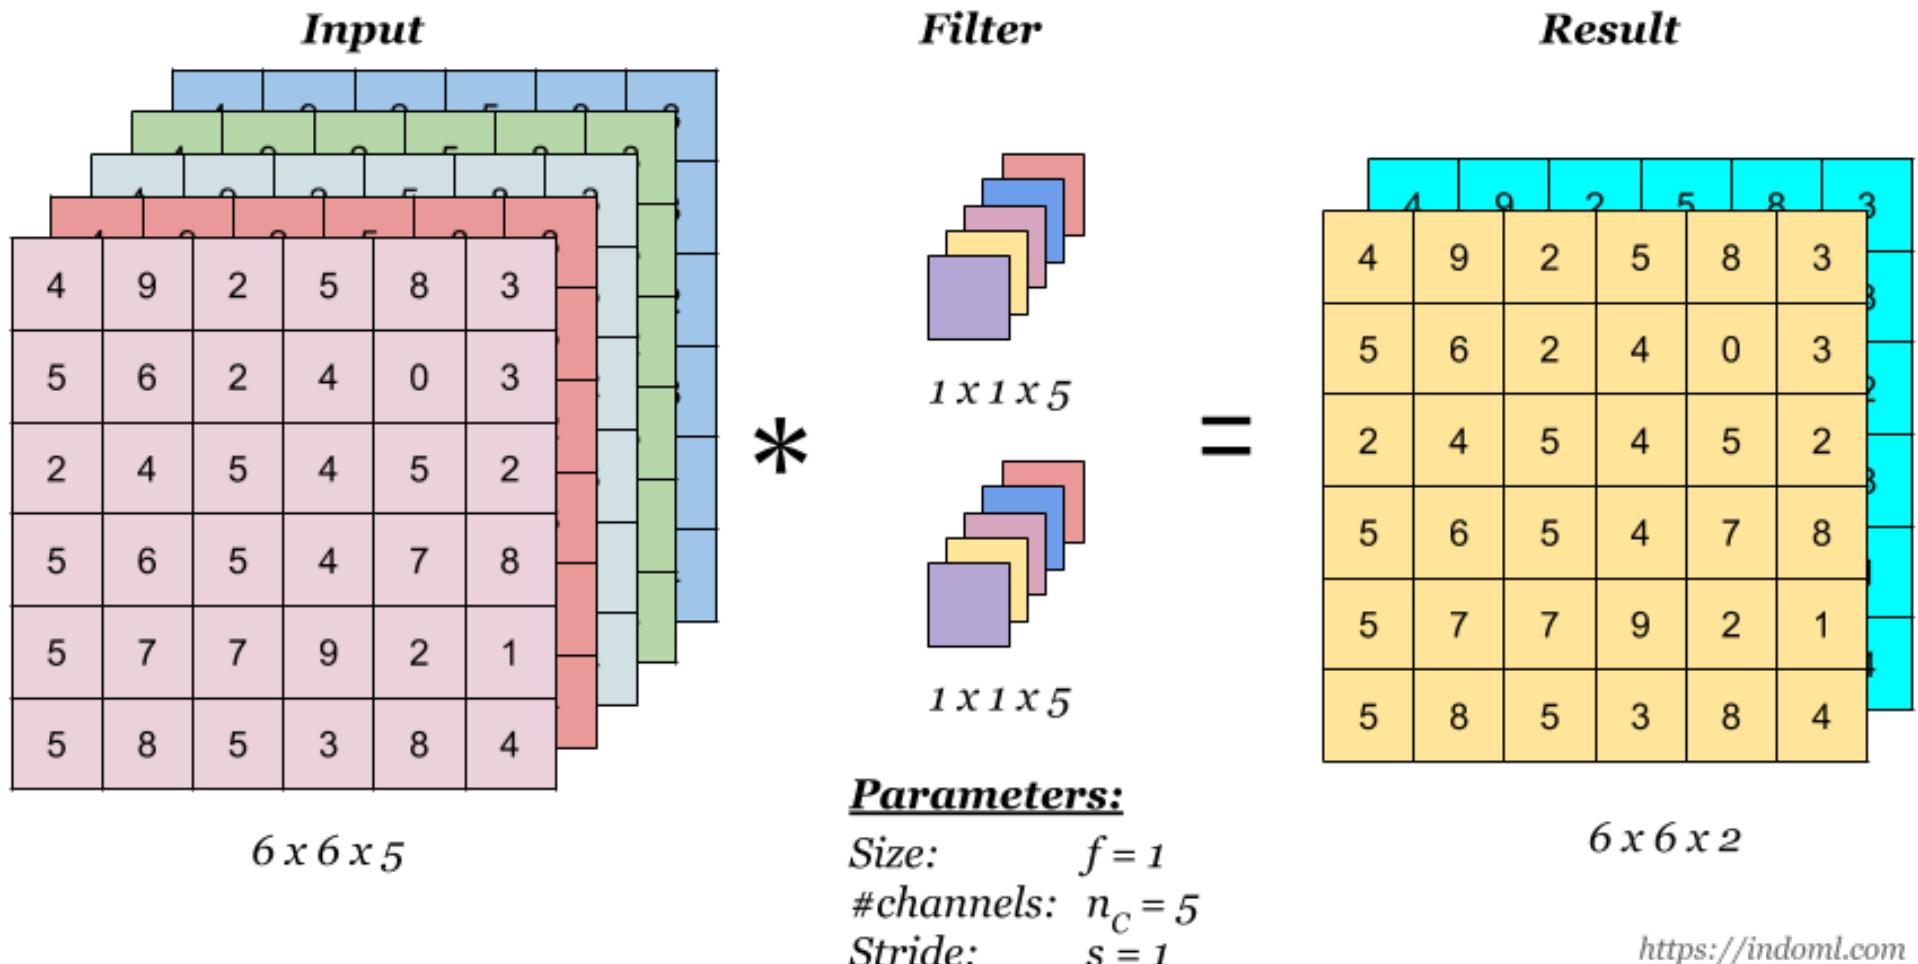


Sample Network Structure

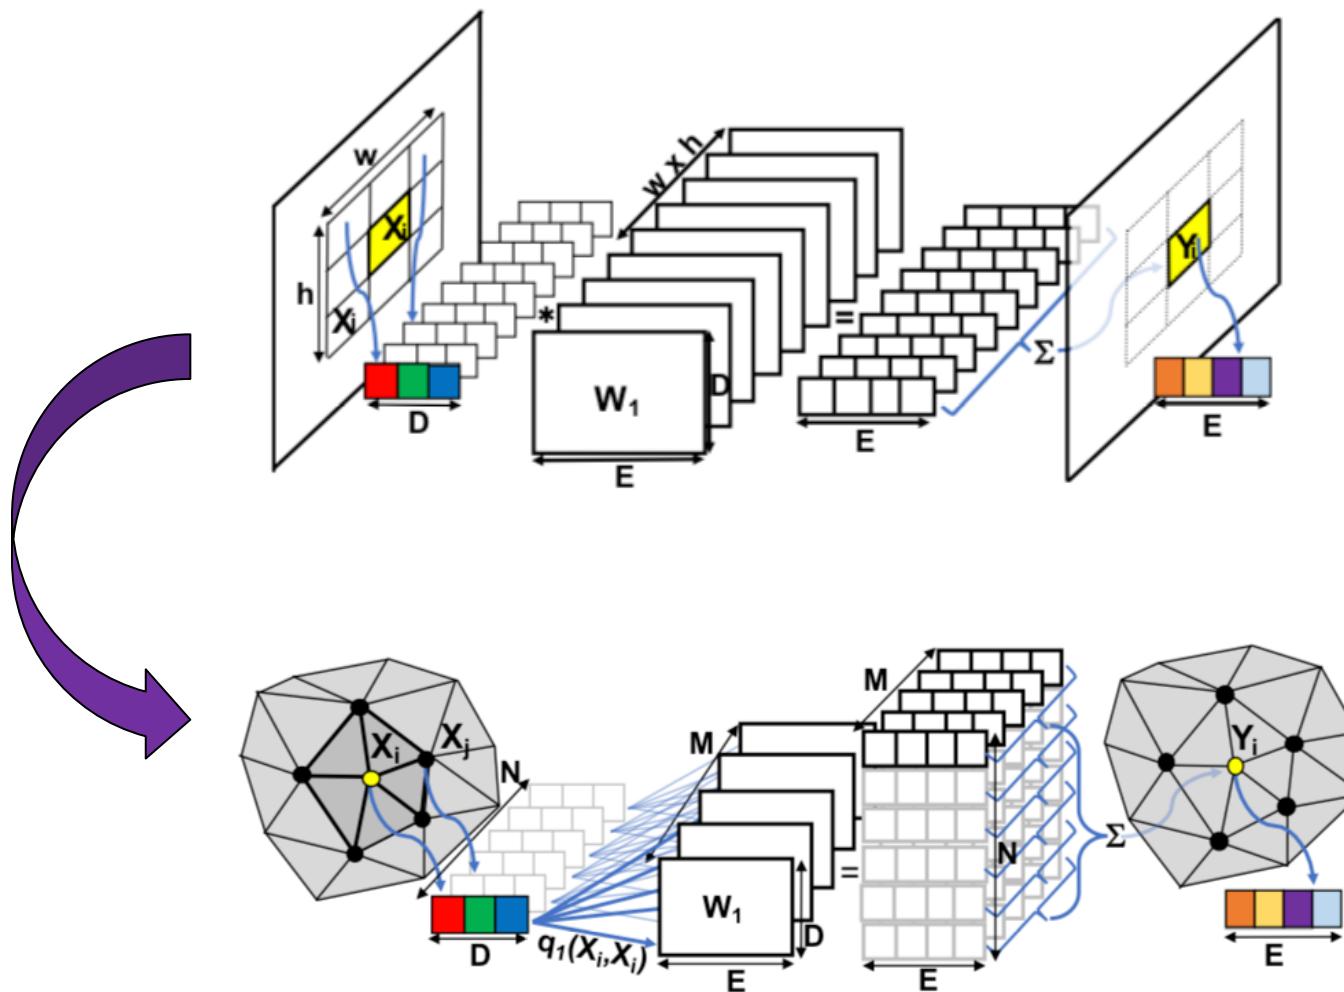


<https://indoml.com>

1*1 Convolutions



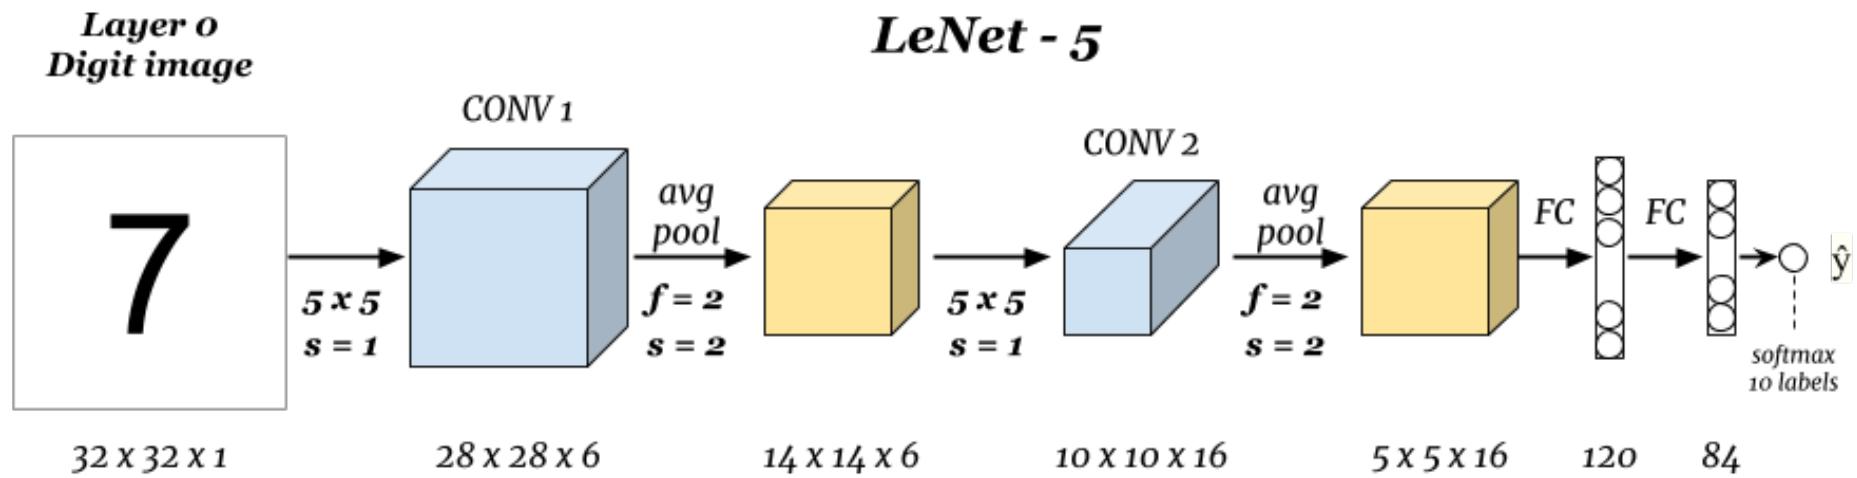
Generalize the CNN



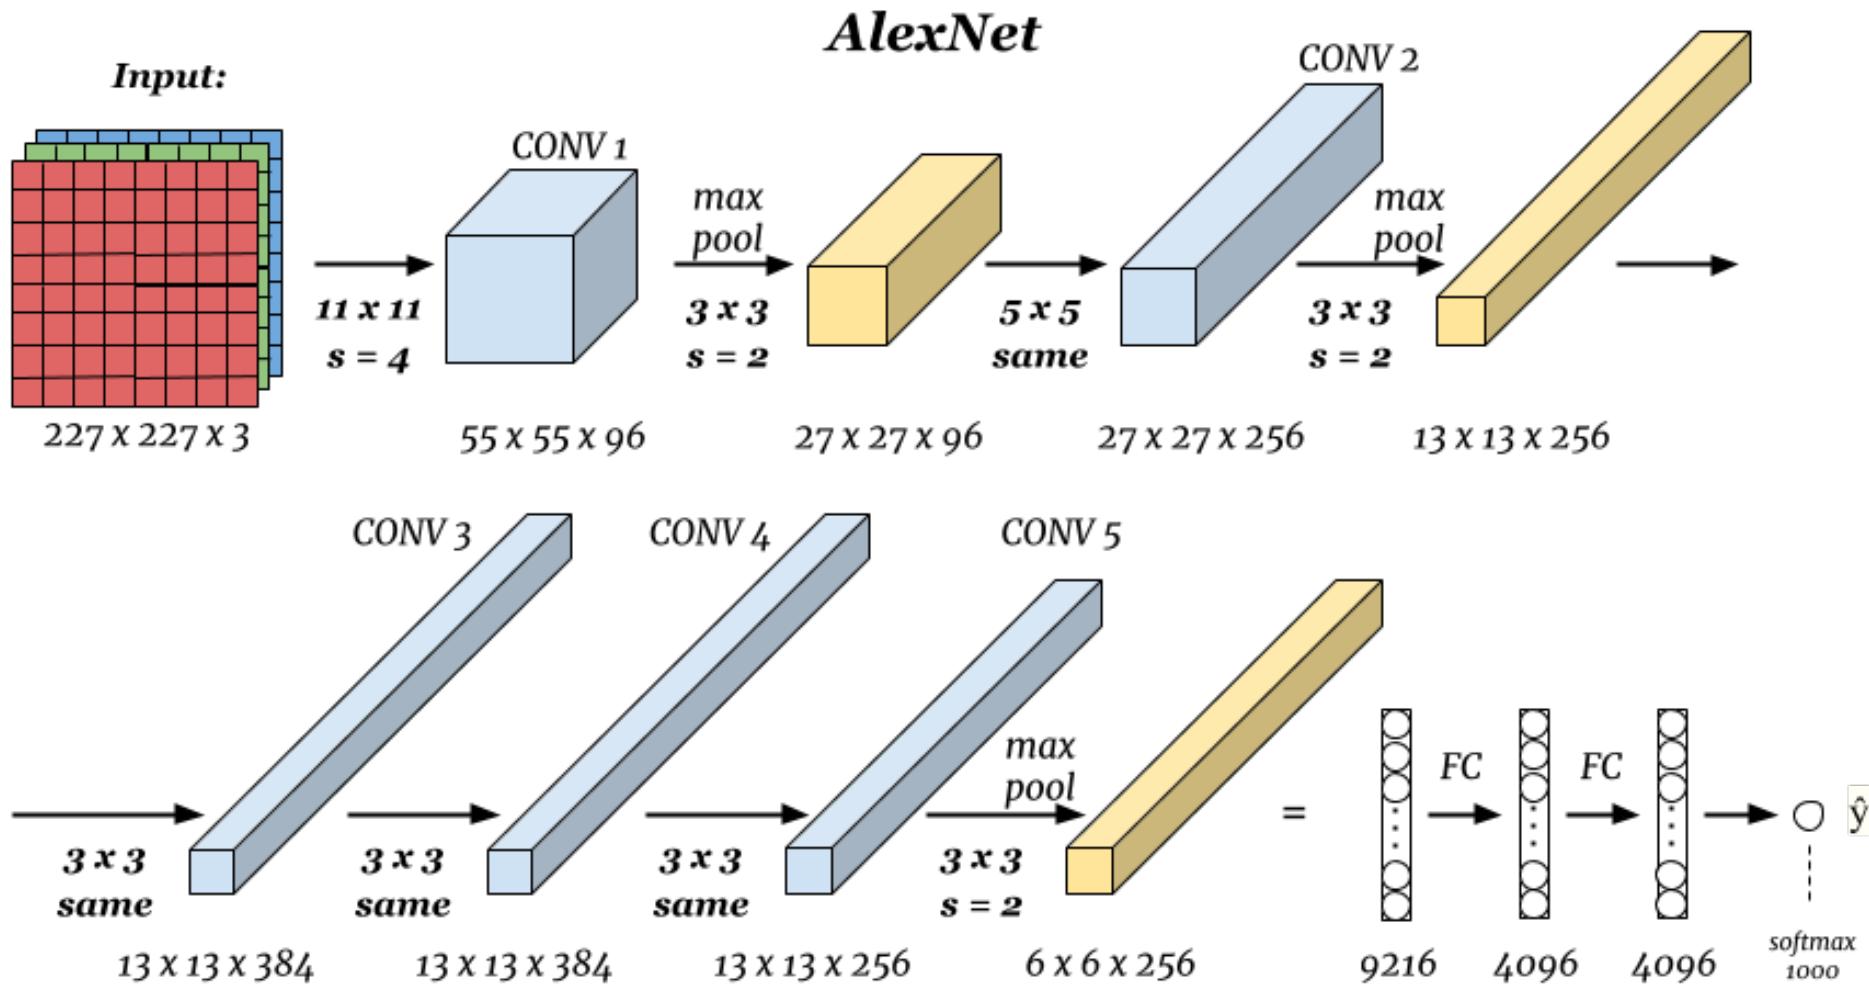
Classic Networks

- LeNet
- AlexNet
- VGG Net
- ResNet
- GooLeNet

LeNet

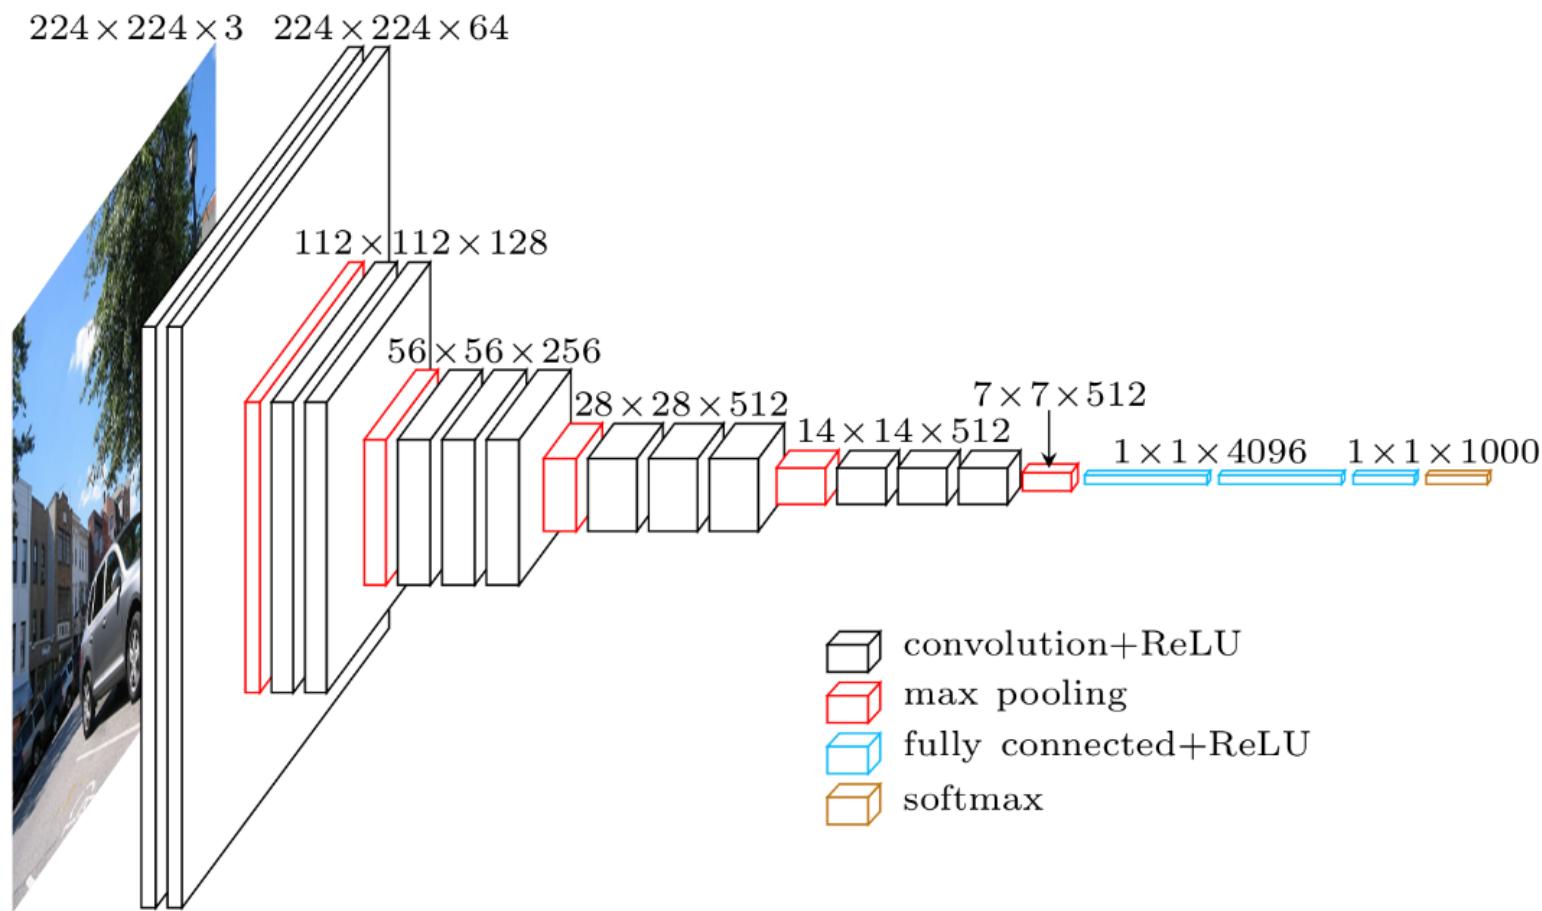


AlexNet

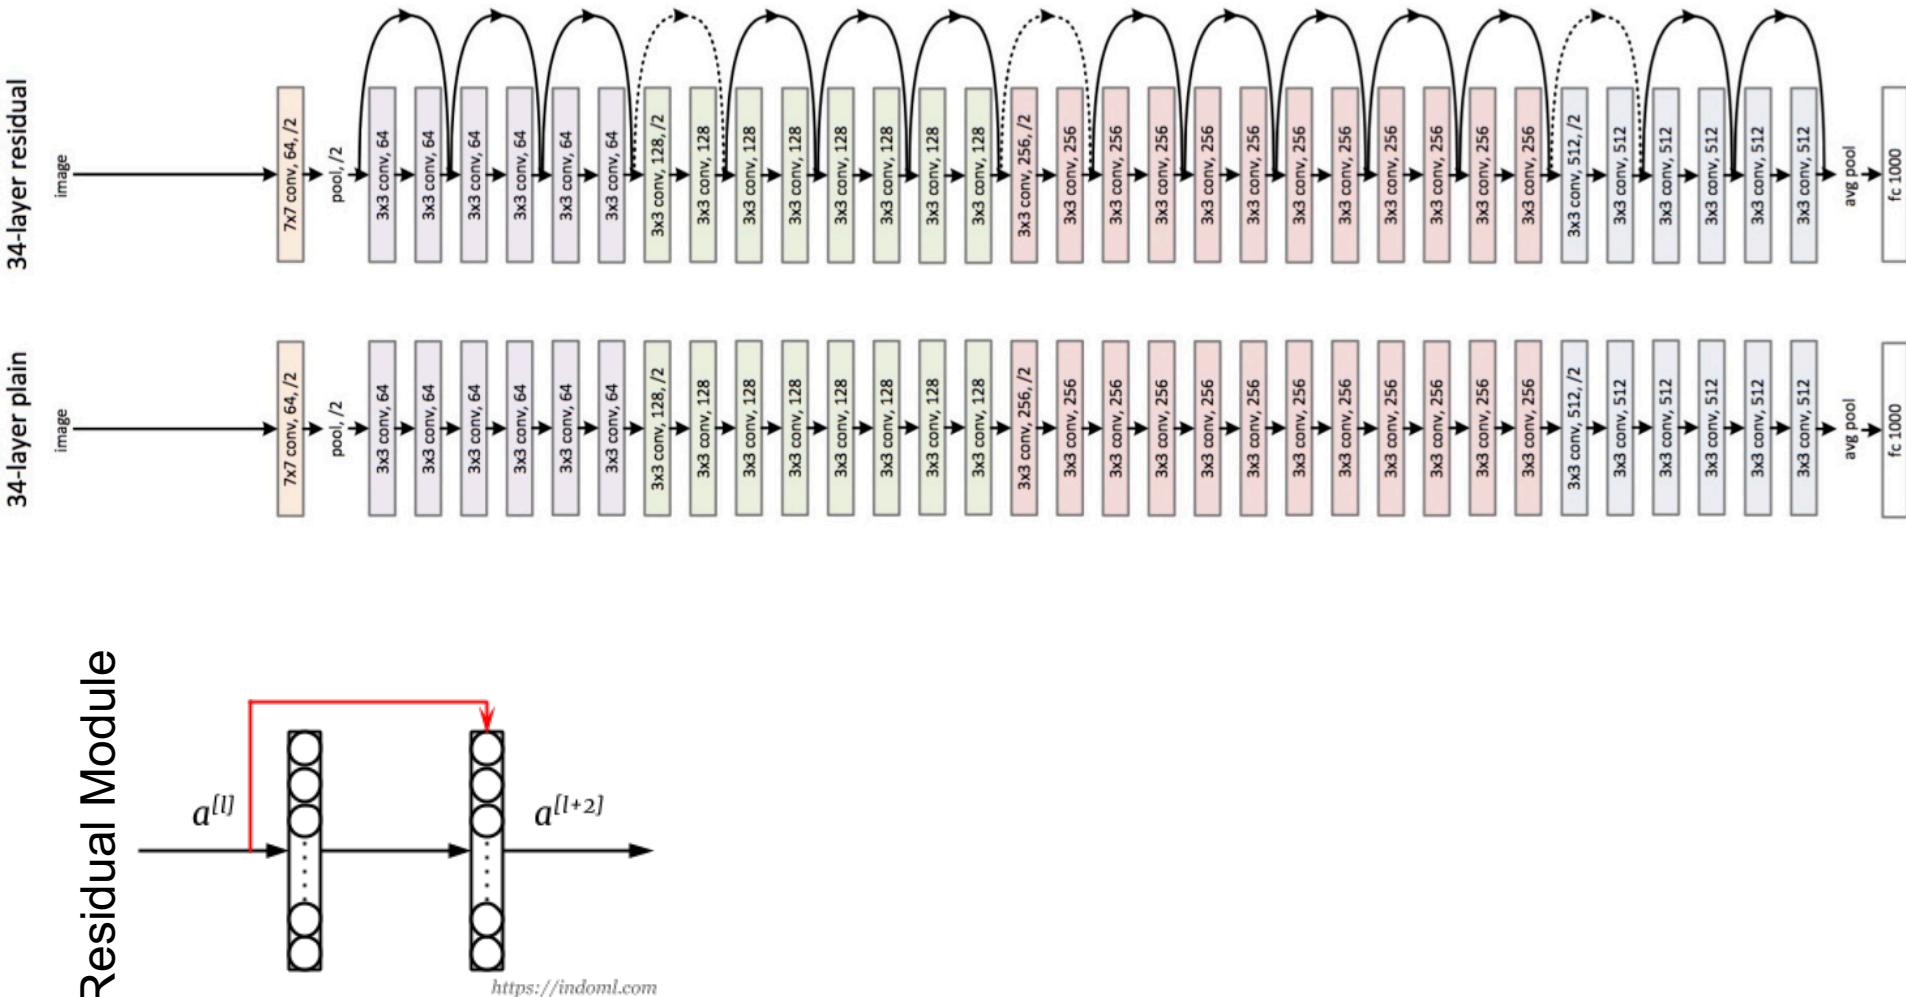


<https://indoml.com>

VGG-16

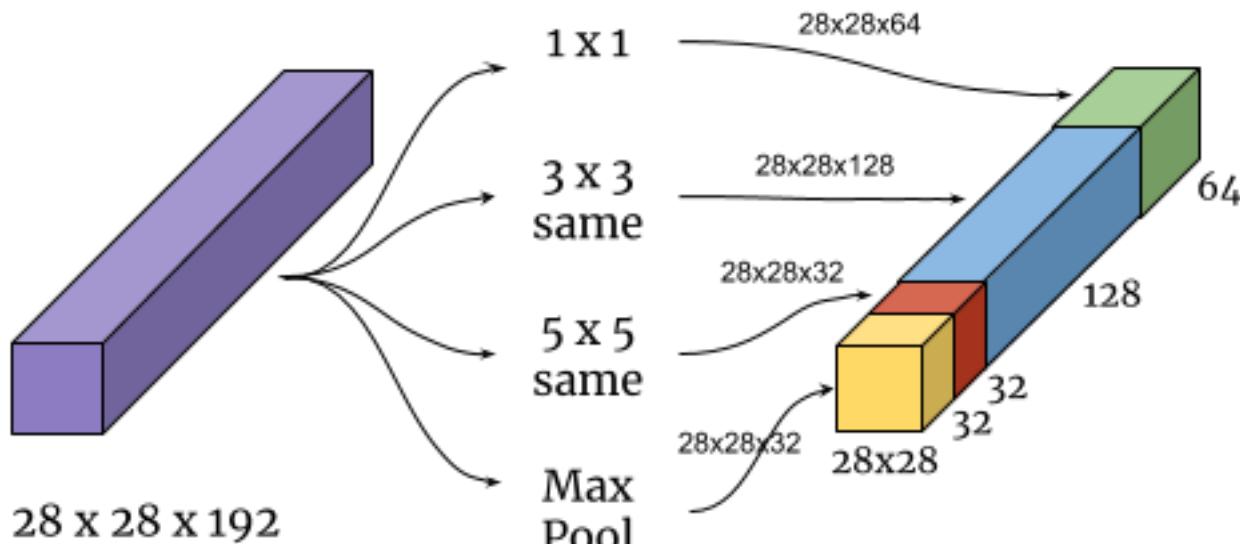


ResNet

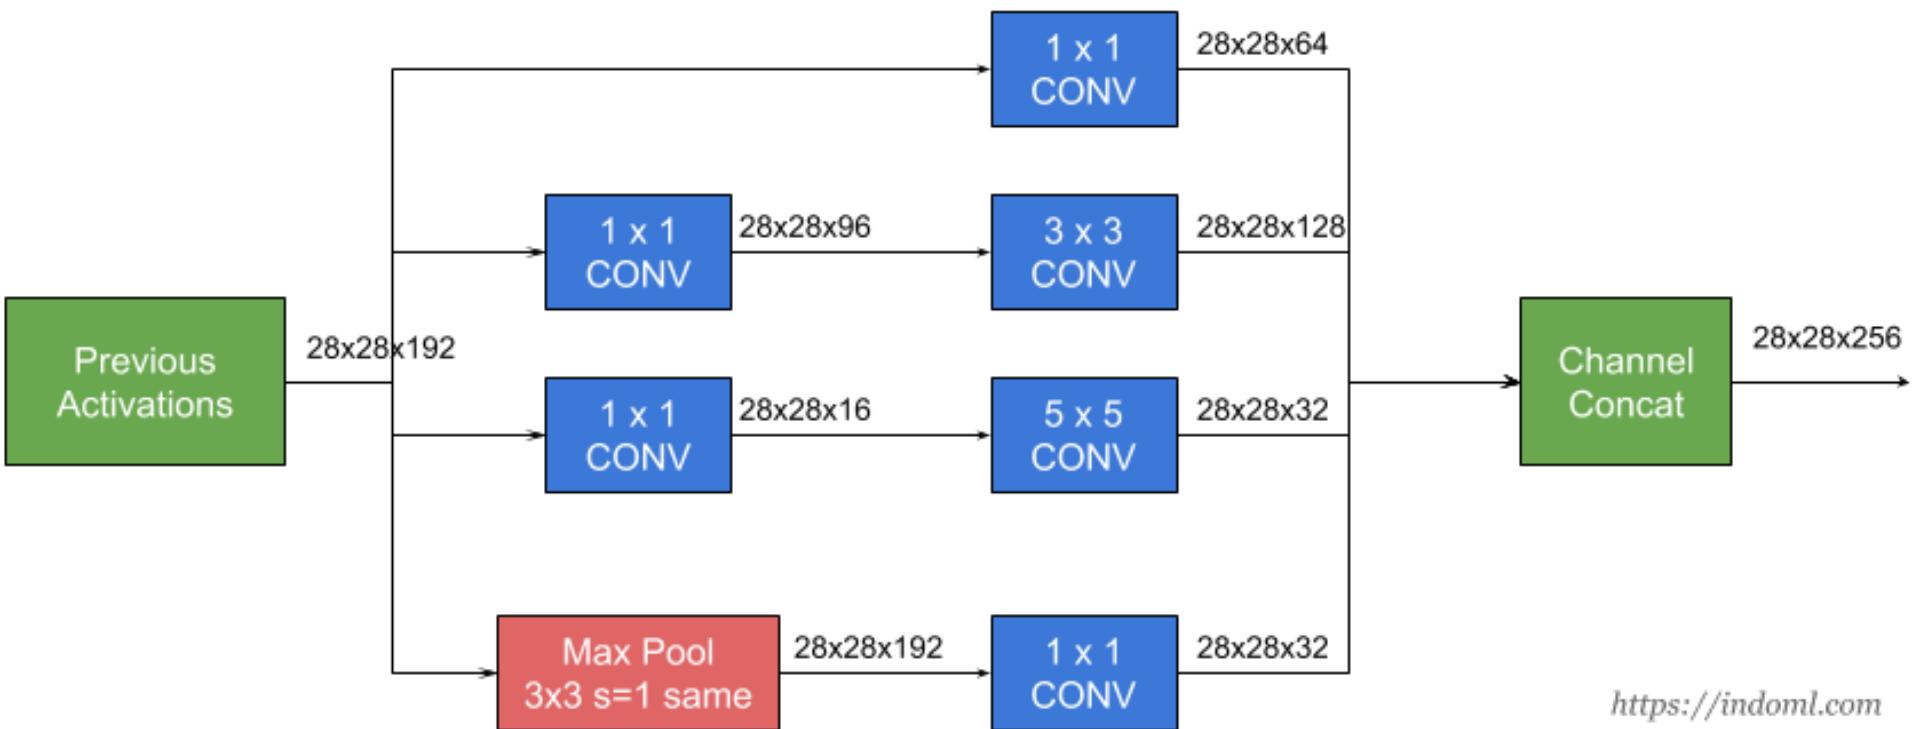


GoogLeNet

Inception Module

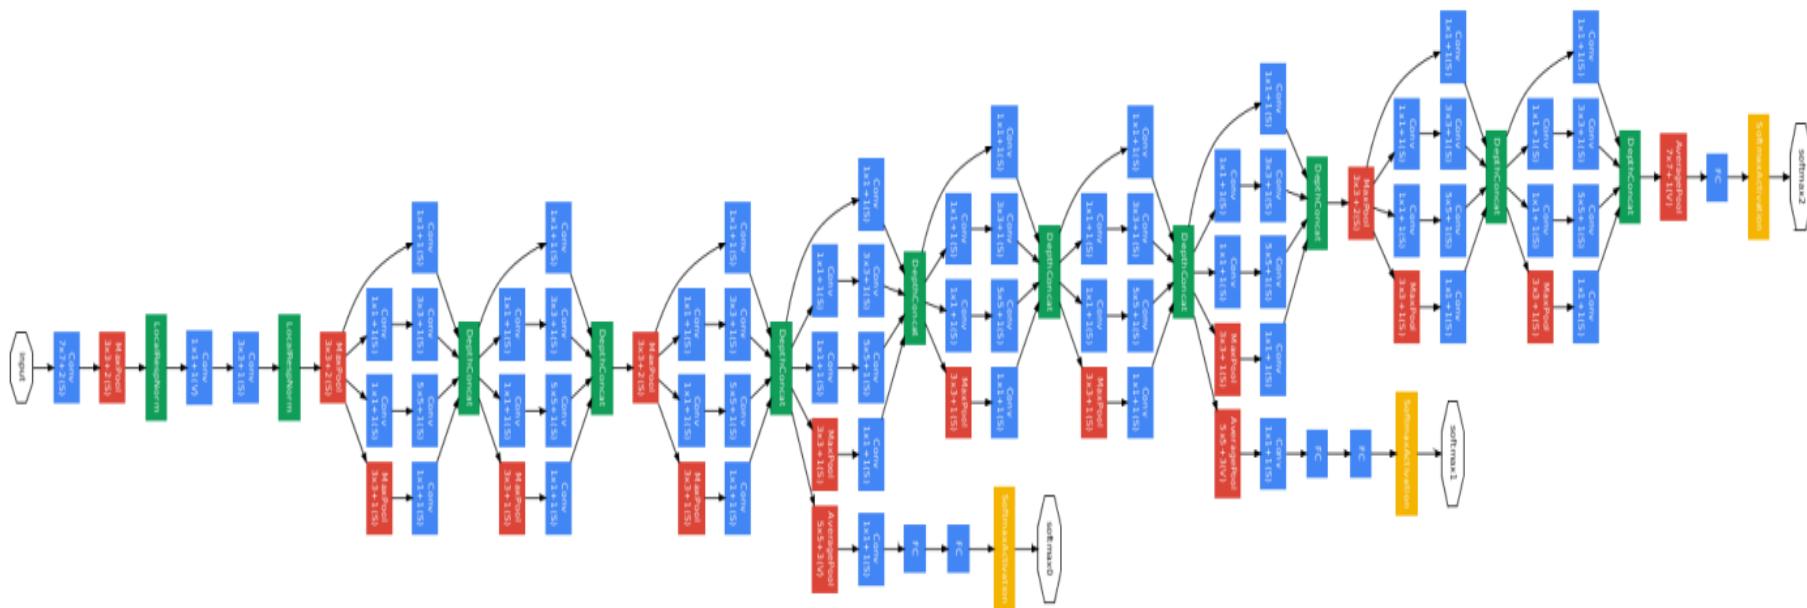


<https://indoml.com>



<https://indoml.com>

GoogLeNet:



Summary of Networks

Year	CNN	Developed by	Place	Top-5 error rate	No. of parameters
1998	LeNet(8)	Yann LeCun et al			60 thousand
2012	AlexNet(7)	Alex Krizhevsky, Geoffrey Hinton, Ilya Sutskever	1st	15.3%	60 million
2014	GoogLeNet(19)	Google	1st	6.67%	4 million
2014	VGG Net(16)	Simonyan, Zisserman	2nd	7.3%	138 million
2015	ResNet(152)	Kaiming He	1st	3.6%	

References

- CNN: ML Lecture 10: Convolutional Neural Network by Hongyi Li (youtube)
- Chapter 20 from Berkeley Artificial Intelligence: A Modern Approach (link: <http://aima.eecs.berkeley.edu/>)
- Module 2: Convolutional Neural Networks from course notes of Stanford CS231n (link: <http://cs231n.github.io/convolutional-networks/>)
- Full connection blog (link:
<https://www.superdatascience.com/blogs/convolutional-neural-networks-cnn-step-4-full-connection>)
- <http://aima.eecs.berkeley.edu/>
- <http://cs231n.github.io/convolutional-networks/>