

Machine Learning, Spring 2019

Generative Adversarial Networks (GAN)

Instructor: Prof. Yi Fang

yfang@nyu.edu

Python tutorial: <http://learnpython.org/>

TensorFlow tutorial: <https://www.tensorflow.org/tutorials/>

PyTorch tutorial: <https://pytorch.org/tutorials/>

Acknowledge: The slides are partially referred to the online materials by Taegyun Joen, <https://www.slideshare.net/TaegyunJeon1/pr12-you-only-look-once-yolo-unified-realtime-object-detection> and online YOLO paper and other materials (from ECS289g by Prof. Lee)

Magic of GANs

- Which one is Computer generated?



Src: Ledig, Christian, et al. "Photo-realistic single image super-resolution using a generative adversarial network." arXiv preprint arXiv:1609.04802 (2016).

- ARTISTIC STYLE TRANSFER



Src: <https://neurohive.io/en/state-of-the-art/twin-gan-cross-domain-translation-of-human-portraits/>



← → 🔍 https://www.nyu.edu ☆ Profile Logout

Information For: Students Faculty Alumni Employees Community [Login to NYU Home](#) [All NYU](#)

 **NYU** About NYU Admissions Academics University Life Research Search Search icon



Academic Programs



"Games, Roadblocks, and Glitches": A Guide to Our New Tax Code
1040 Label



Courant's LeCun Wins Turing Award for Breakthroughs in AI

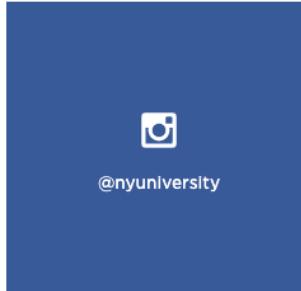
A photograph of Washington Square Park in New York City, featuring the Washington Square Arch and a large fountain in the foreground.



Graduate Admissions



Delivey Science
Feeling Self(ie)-Conscious



 @nyuniversity



Two students walking on a campus path.

May 5, 2019

LeCun, NIPS 2016

- Reinforcement learning (**cherry**)
- Supervised learning (**Chocolate**)
- Unsupervised/Predictive learning (**Cake**)
 - Generative adversarial nets (GAN)



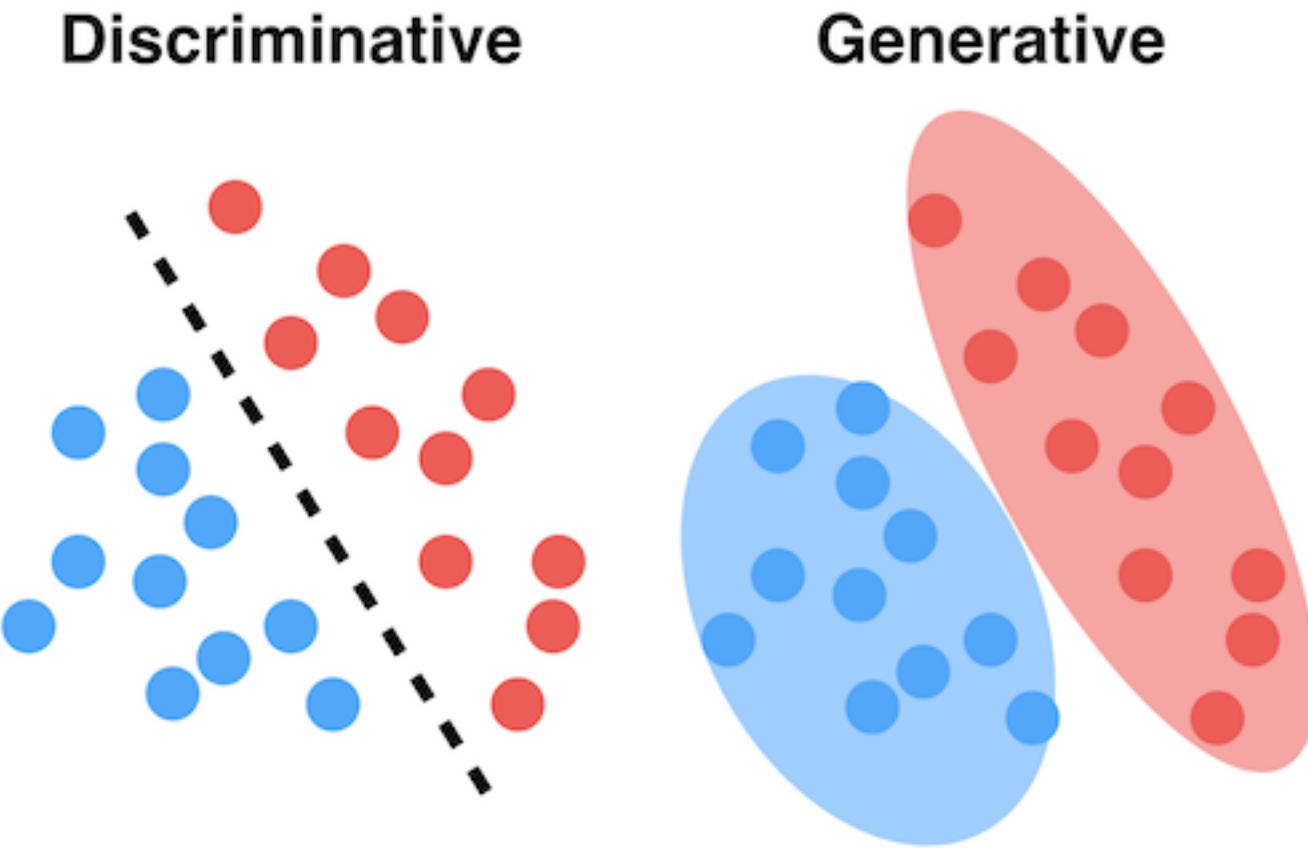
"GANs are the most interesting idea in the last 10 years in ML"

Yann LeCun.

- Introduction of GAN
- GAN applications

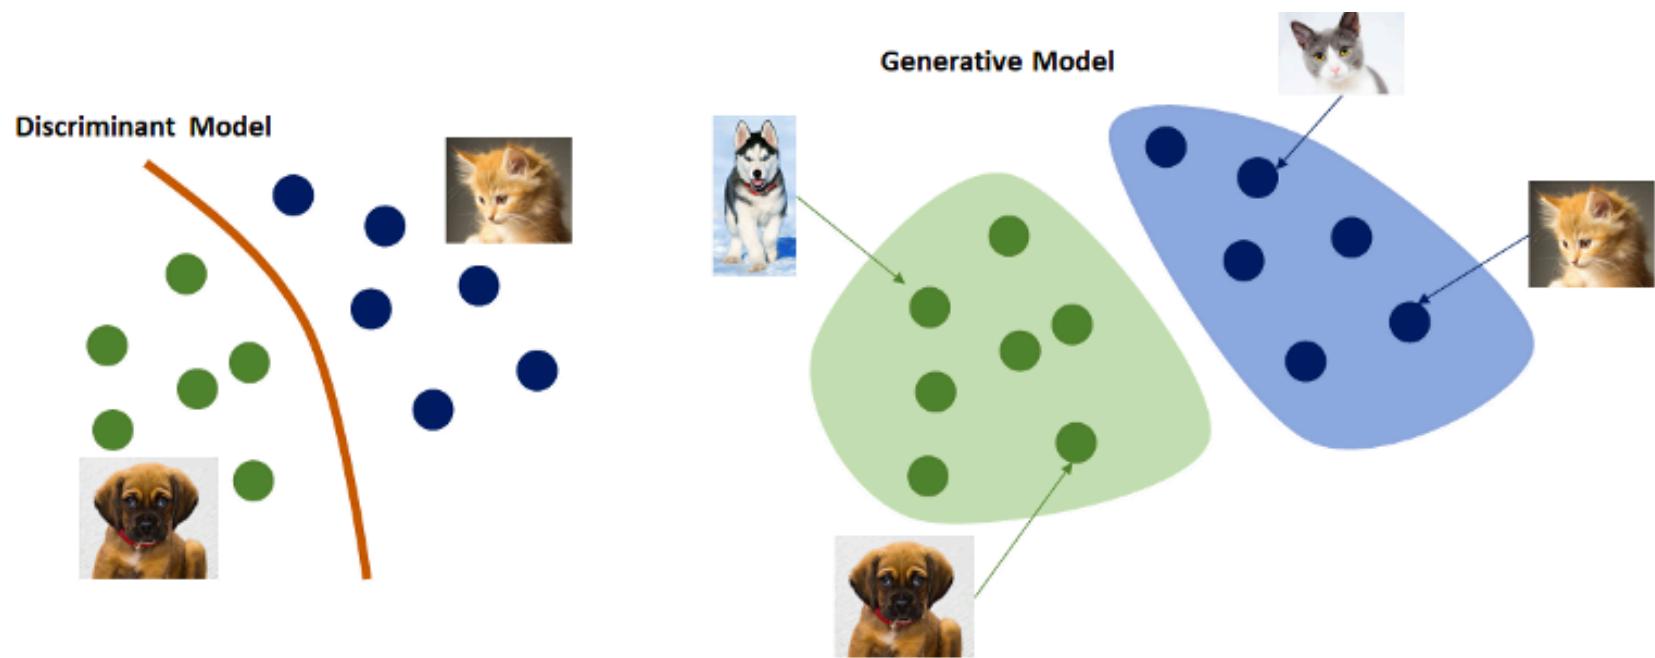
- Introduction of GAN
- GAN applications

Generative and Discriminative models

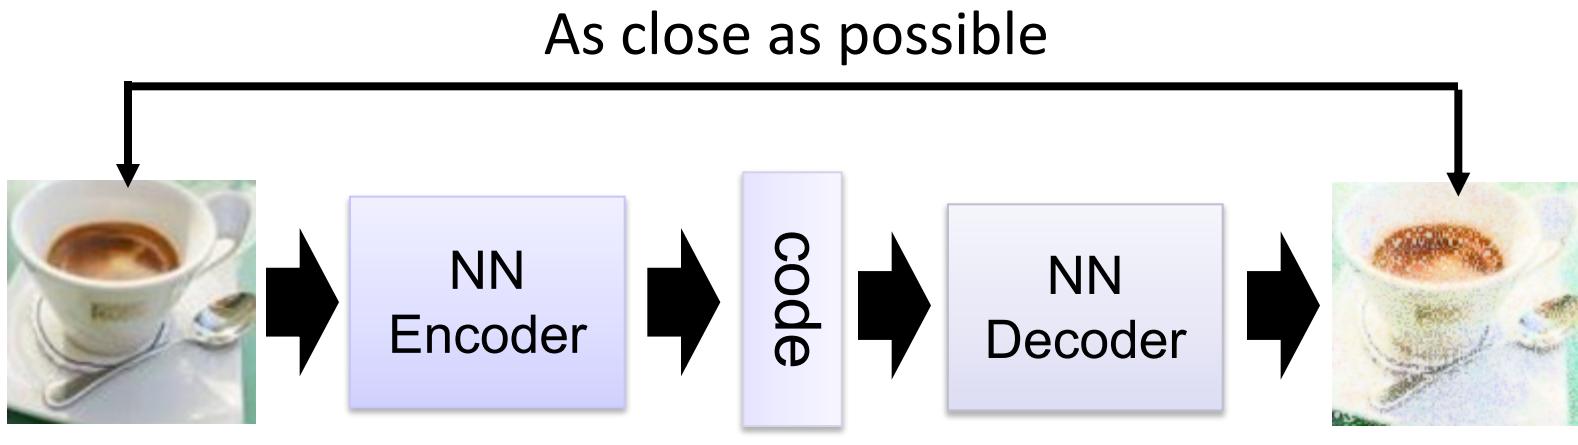


Src: <https://medium.com/@jordi299/about-generative-and-discriminative-models-d8958b67ad32>

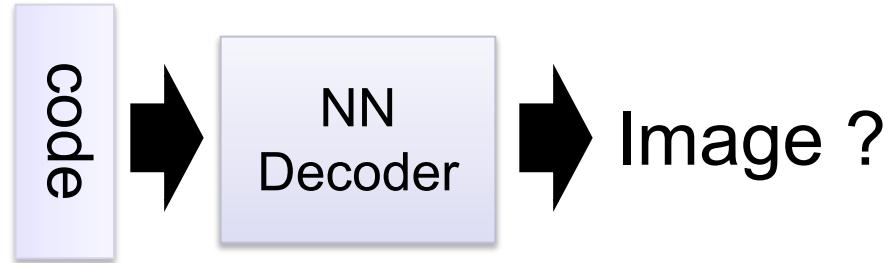
- In supervised learning, we have data x and response (label) y and the goal is to learn a function to map x to y e.g. regression, classification, object detection;
- In unsupervised learning, there are no labels and the goal is to find some underlying hidden structure of the data e.g. clustering, dimensionality reduction, feature learning. The goal of generative models is to generate new samples of data from a distribution. These models are used in problems such as density estimation, a problem of unsupervised learning.



Autoencoder

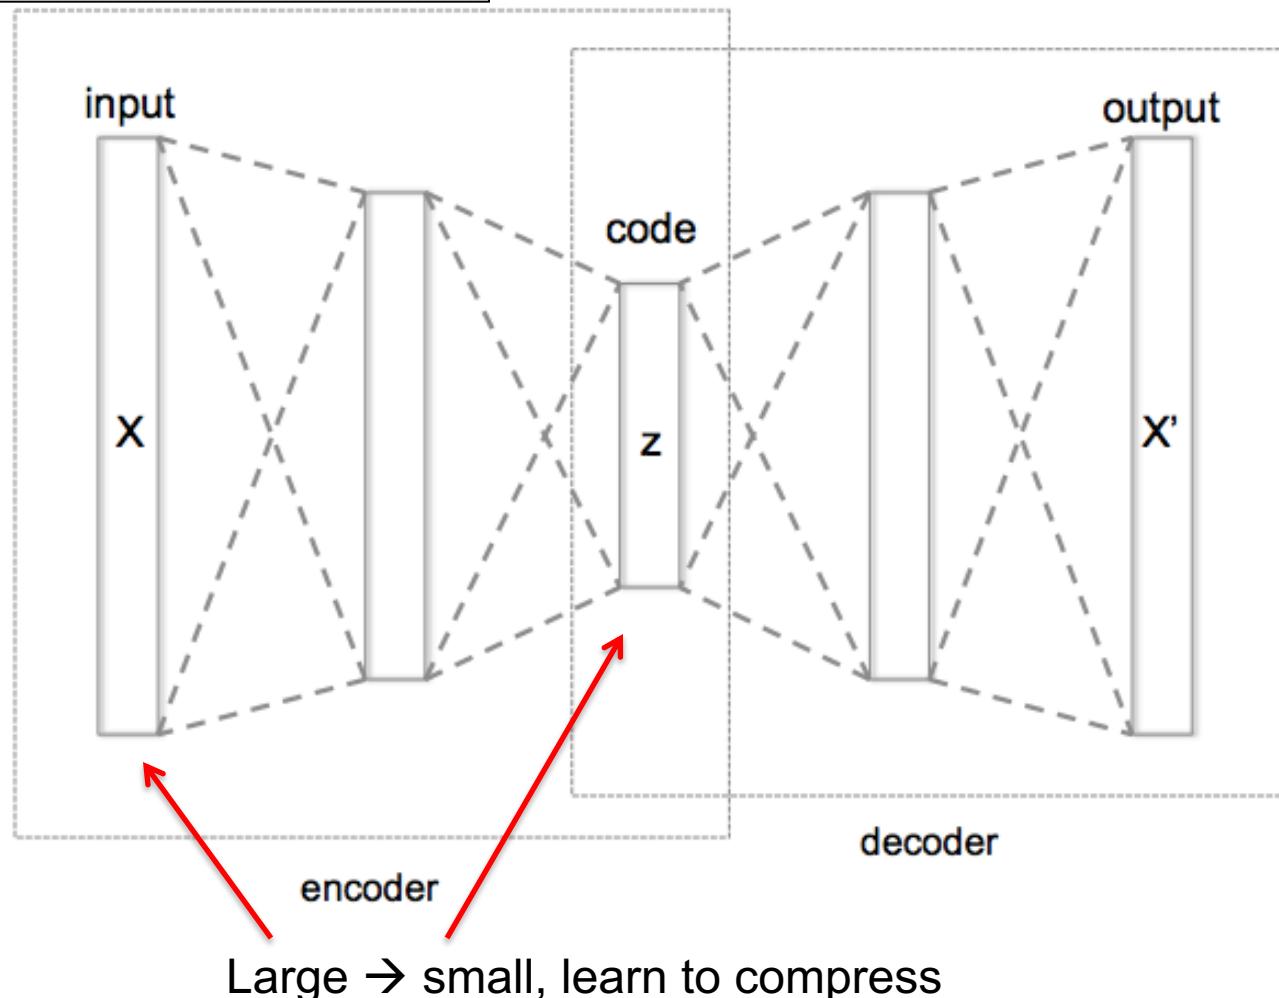


Randomly
generate a vector
as code

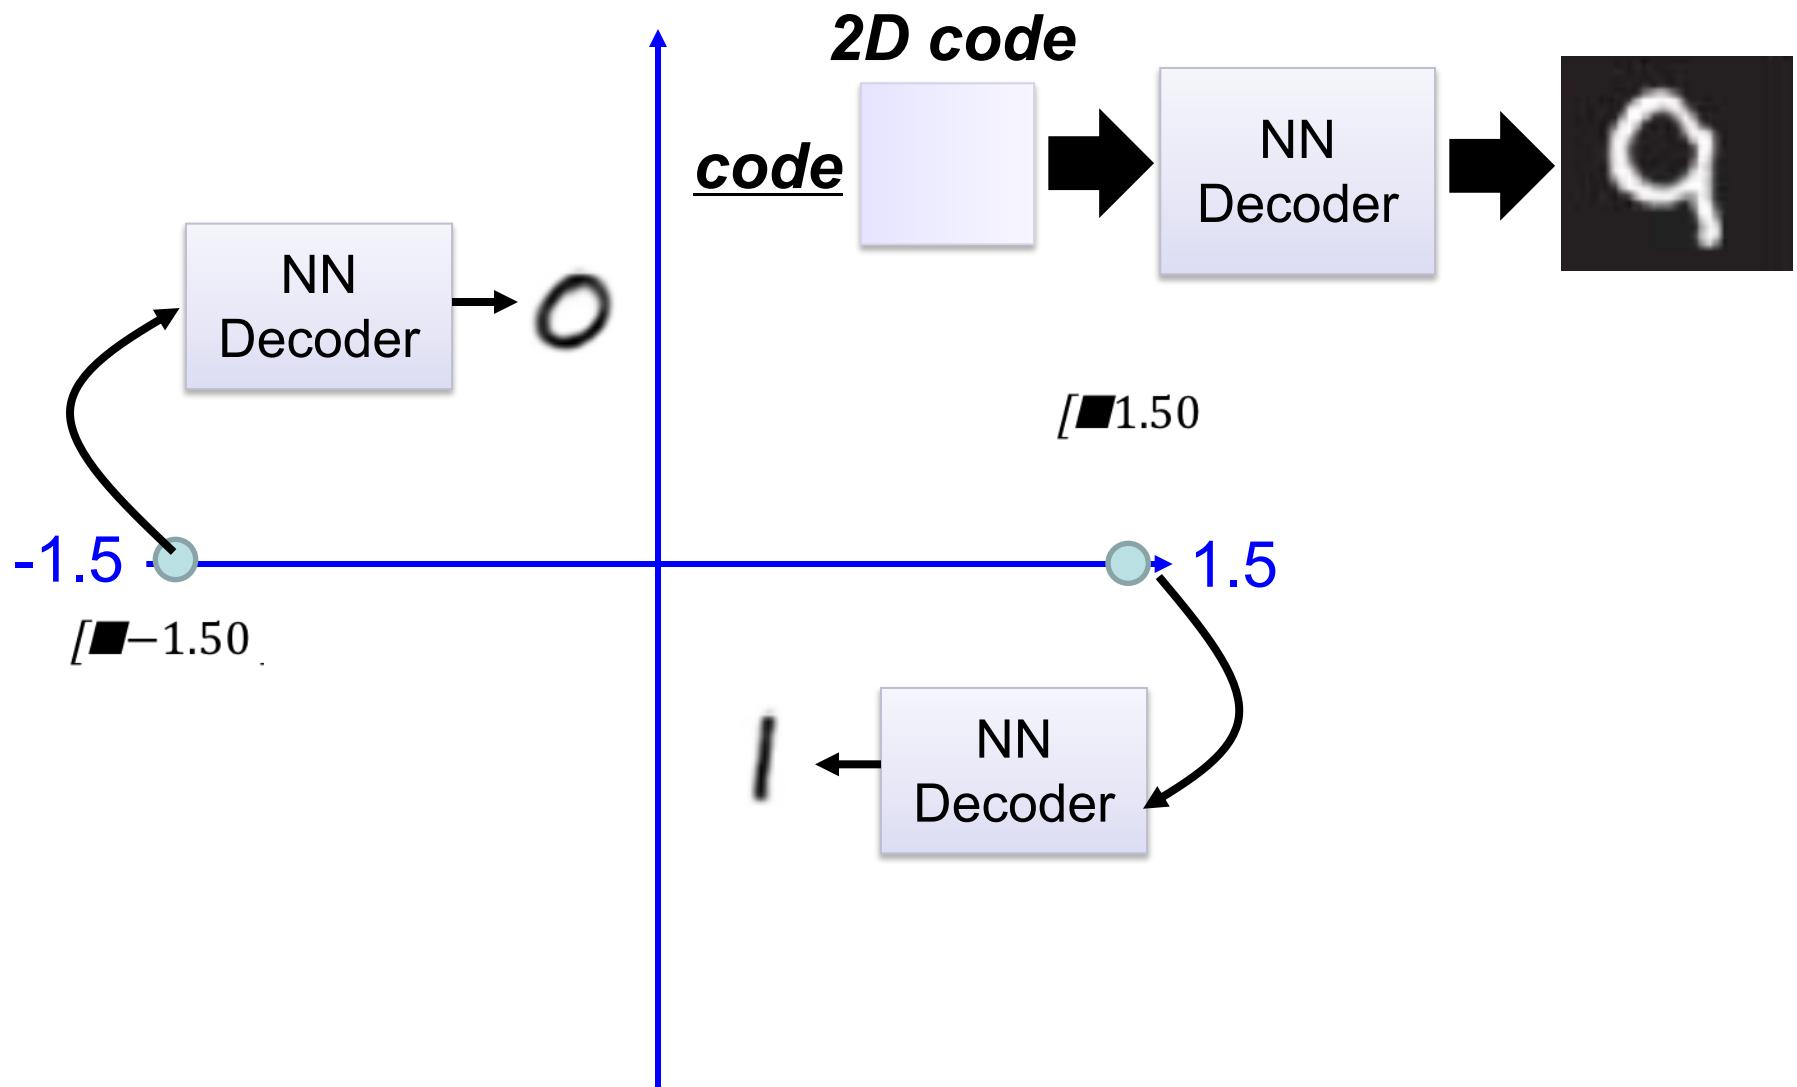


Autoencoder with 3 fully connected layers

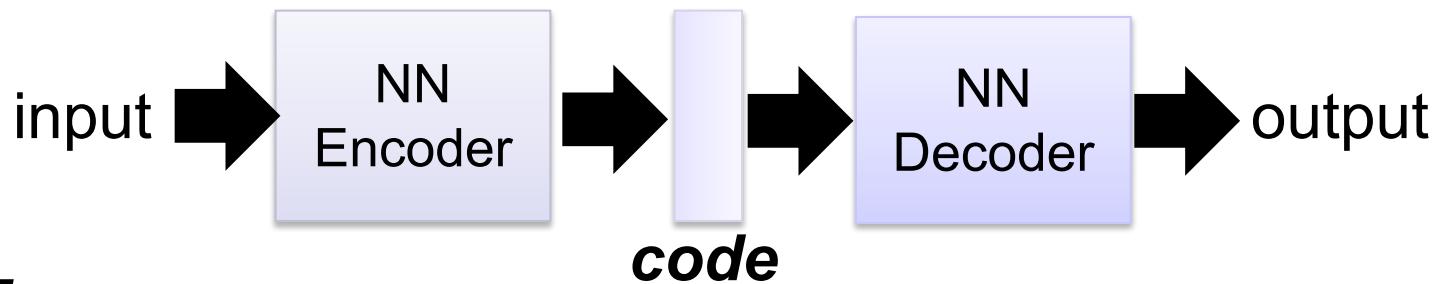
Training: `model.fit(X,X)`
Cost function: $\sum_{k=1..N} (x_k - x'_k)^2$



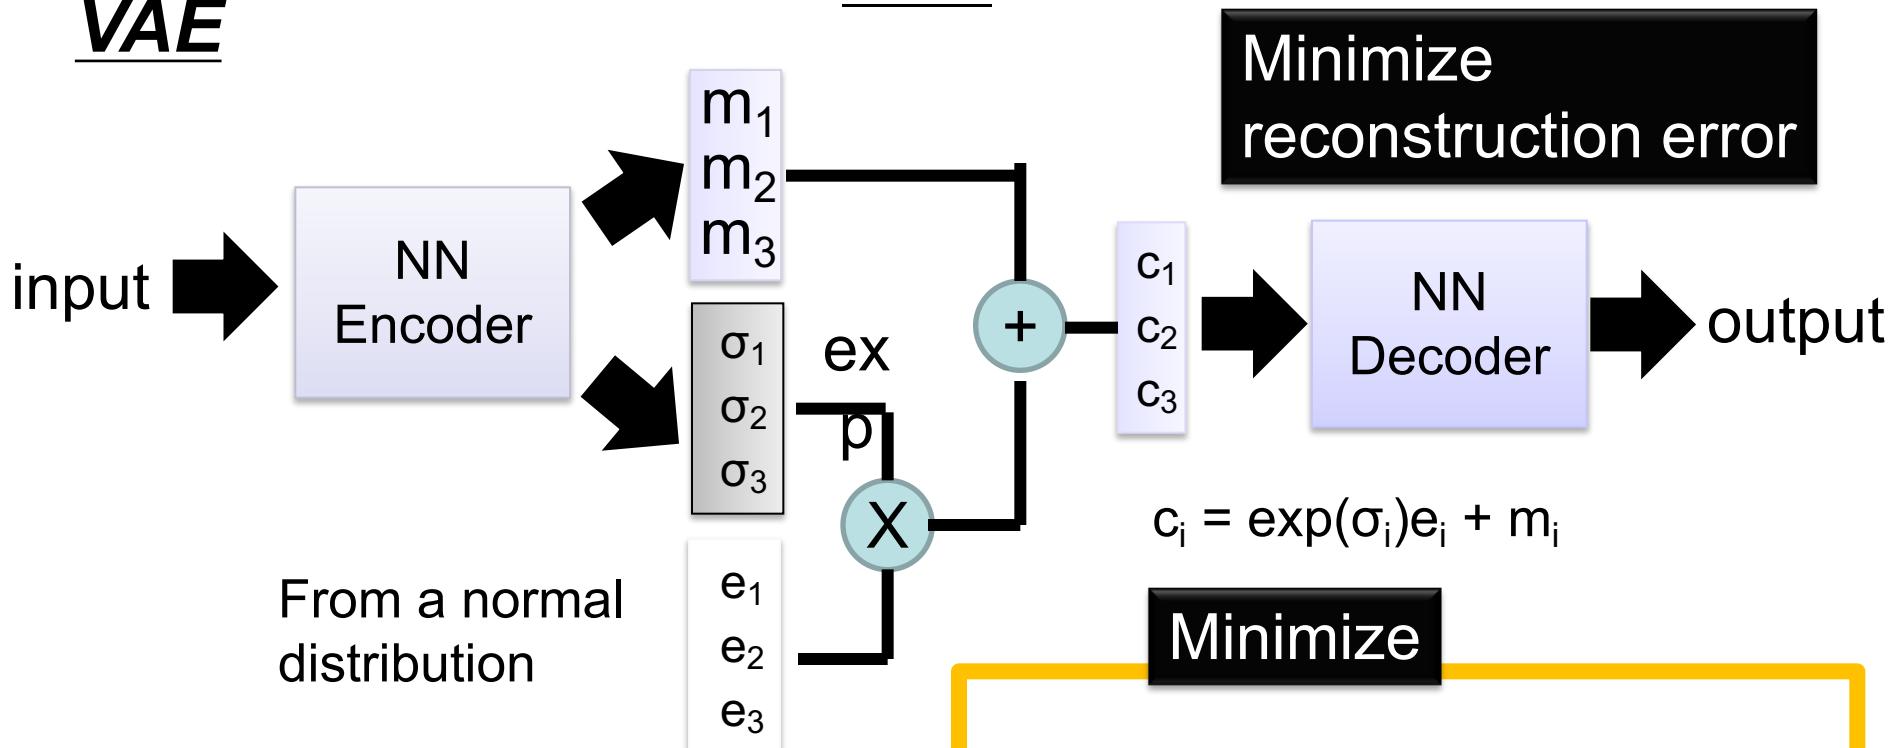
Auto-encoder



Auto-encoder



VAE



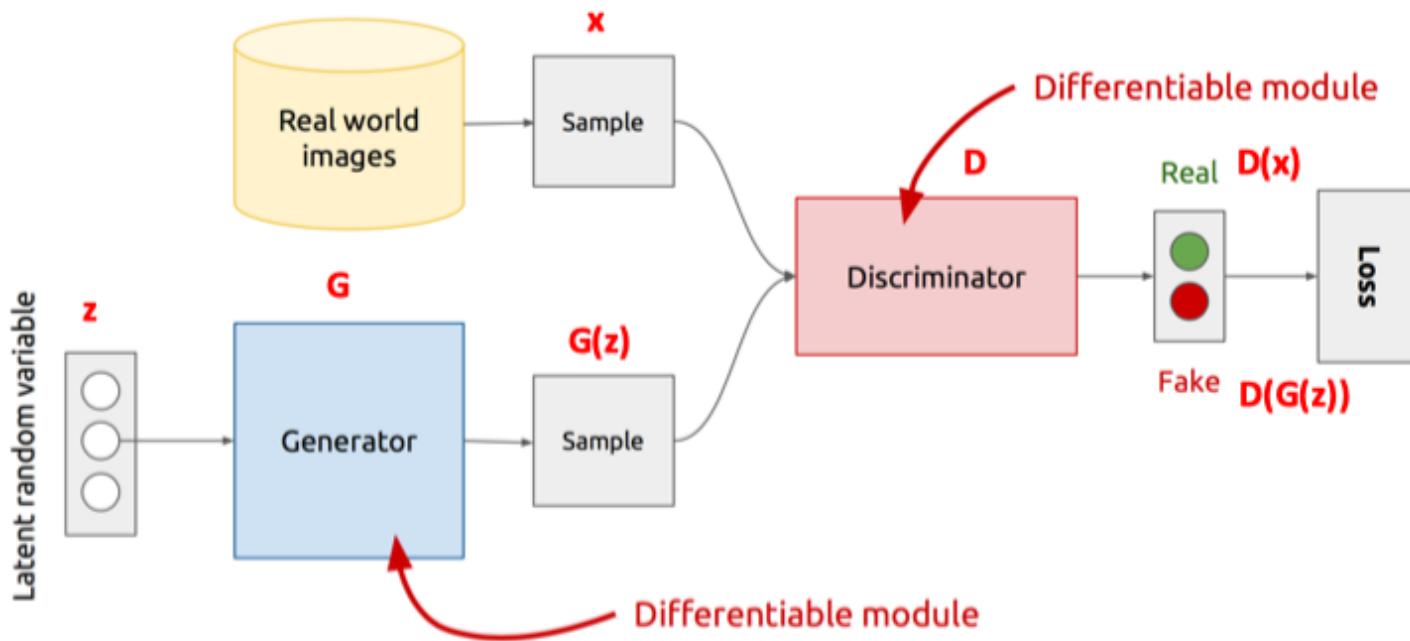
Auto-Encoding Variational Bayes,
<https://arxiv.org/abs/1312.6114>

$$\sum_{i=1..3} [\exp(\sigma_i) - (1 + \sigma_i) + (m_i)^2]$$

This constrains σ_i approaching 0 is good

Generative Adversarial Network

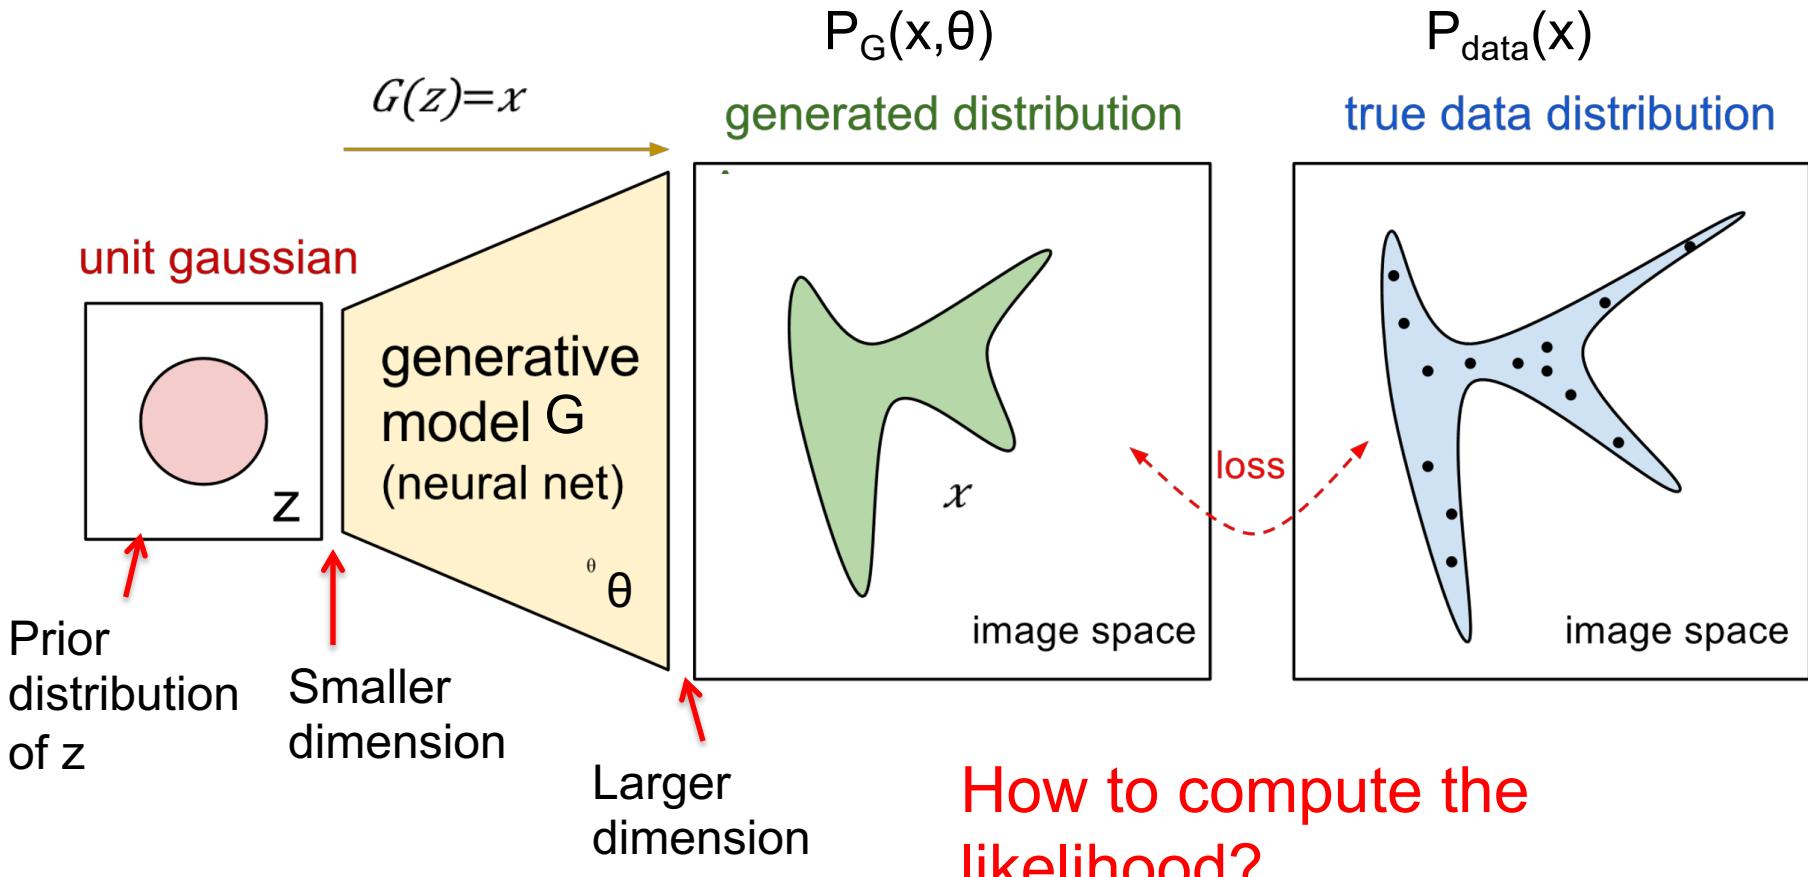
GAN's Architecture



- Z is some random noise (Gaussian/Uniform).
- Z can be thought as the latent representation of the image.

<https://www.slideshare.net/xavigiro/deep-learning-for-computer-vision-generative-models-and-adversarial-training-upc-2016>

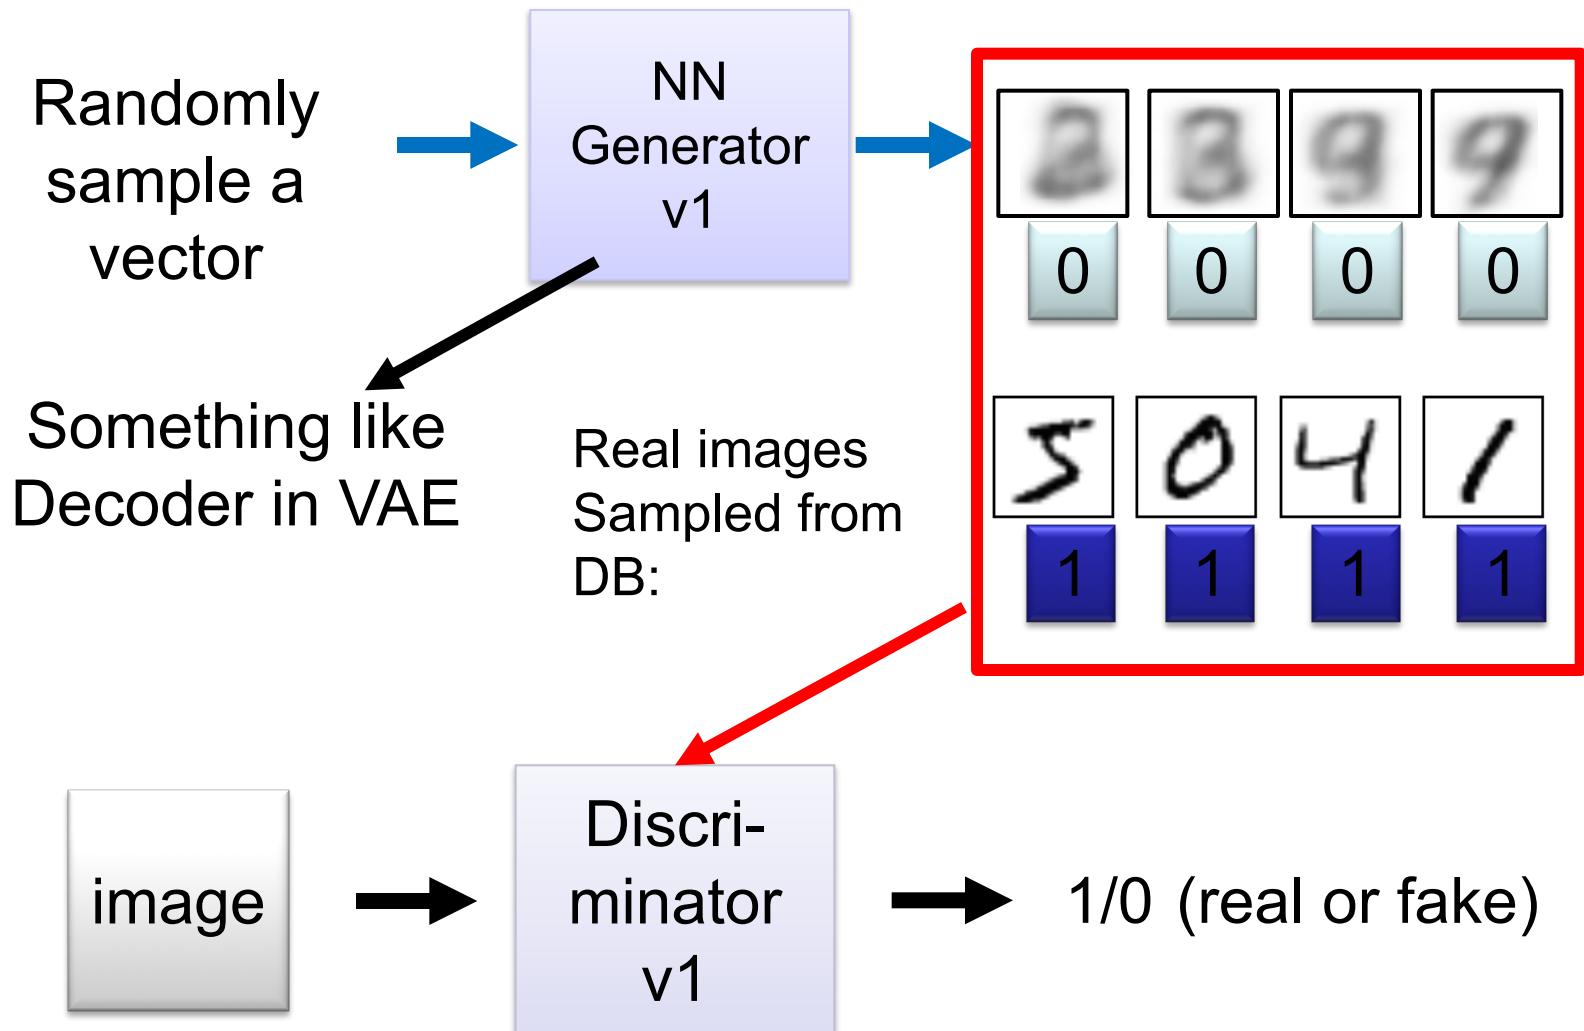
Neural Network as $P_G(x; \theta)$



$$P_G(x) = \text{Integration}_z P_{\text{prior}}(z) I_{[G(z)=x]} dz$$

<https://blog.openai.com/generative-models/>

GAN – Learn a discriminator



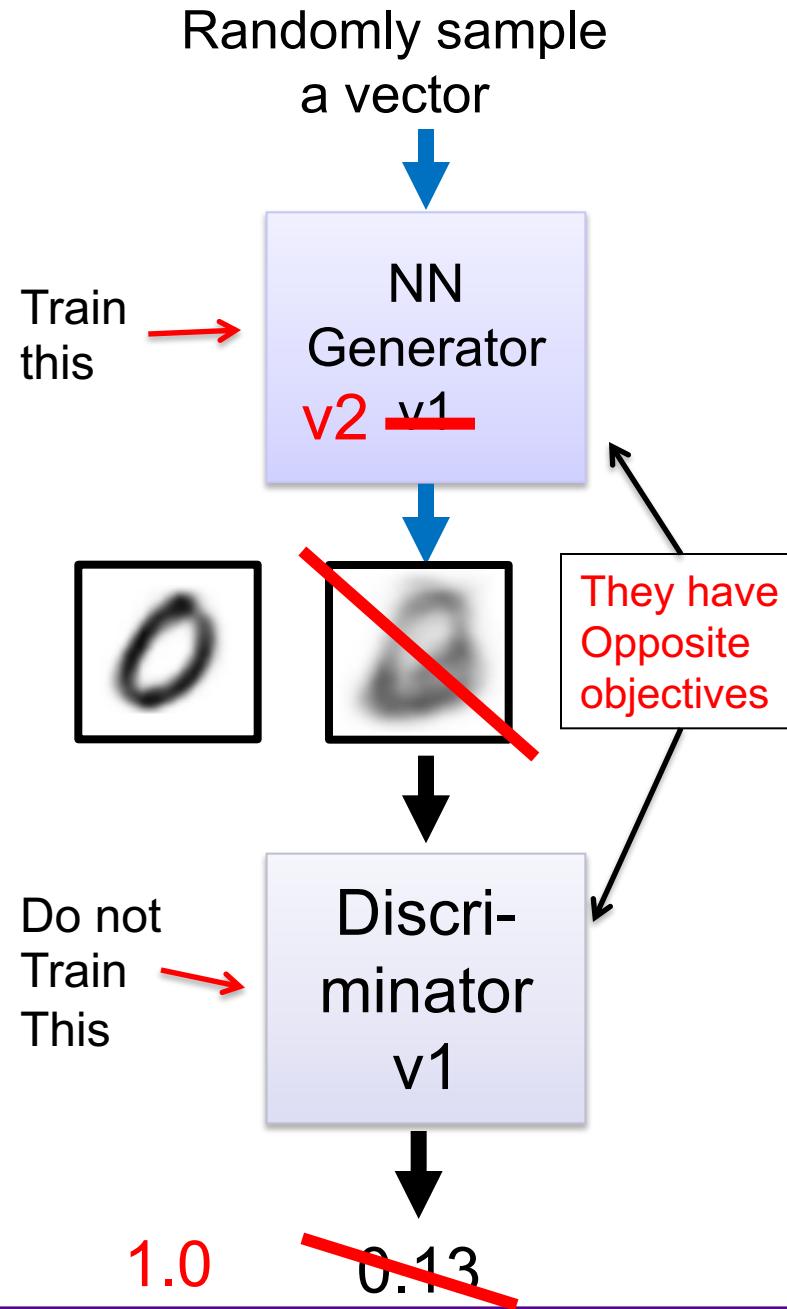
GAN – Learn a generator

Updating the parameters of generator

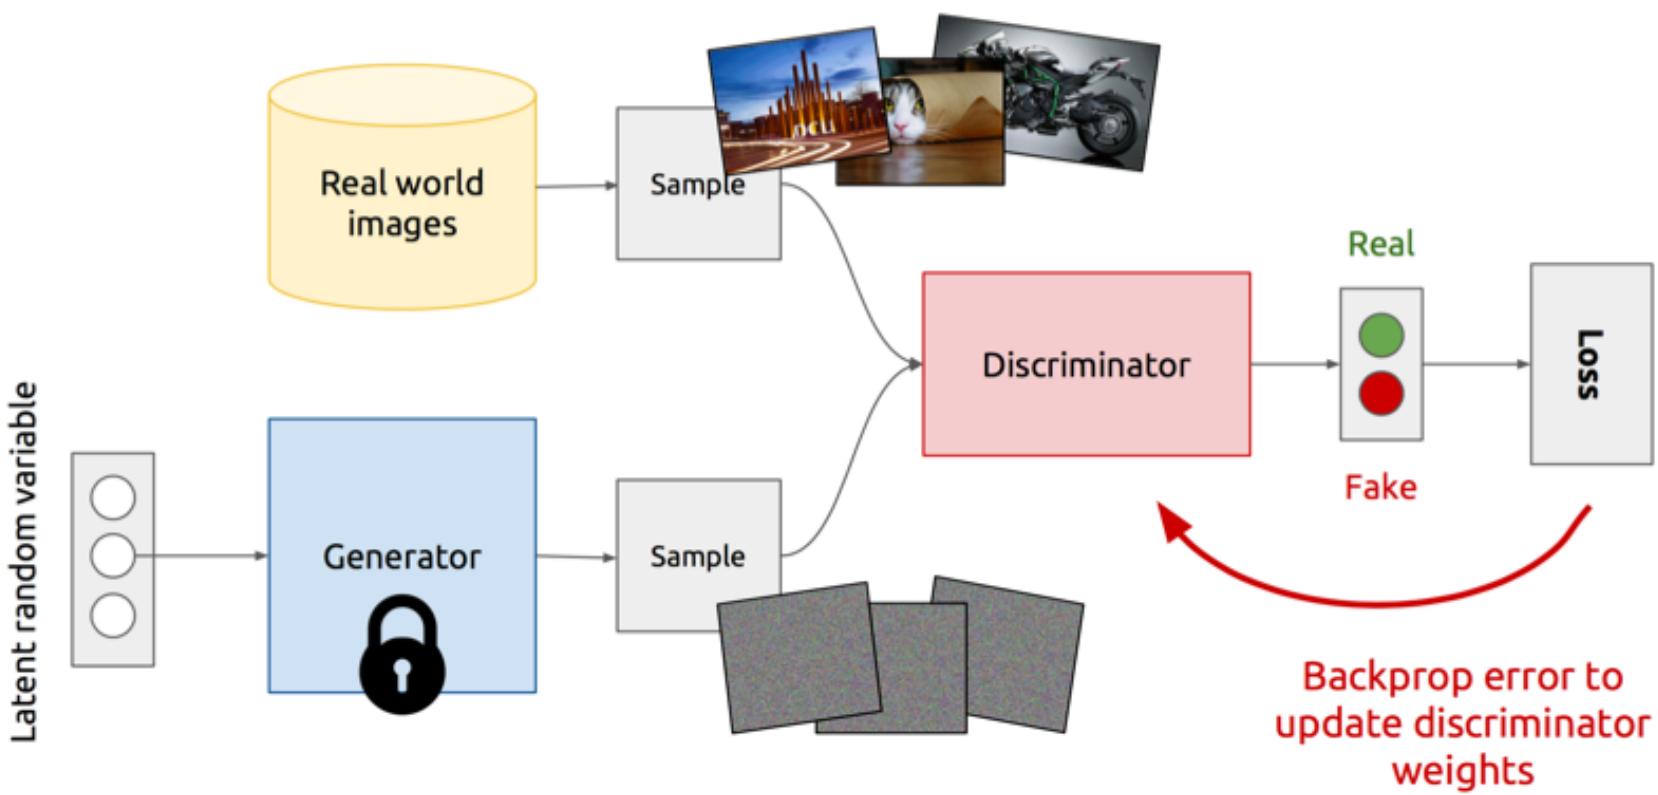
→ The output be classified as “real” (as close to 1 as possible)

Generator + Discriminator =
a network

Using gradient descent to update the parameters in the generator, but fix the discriminator

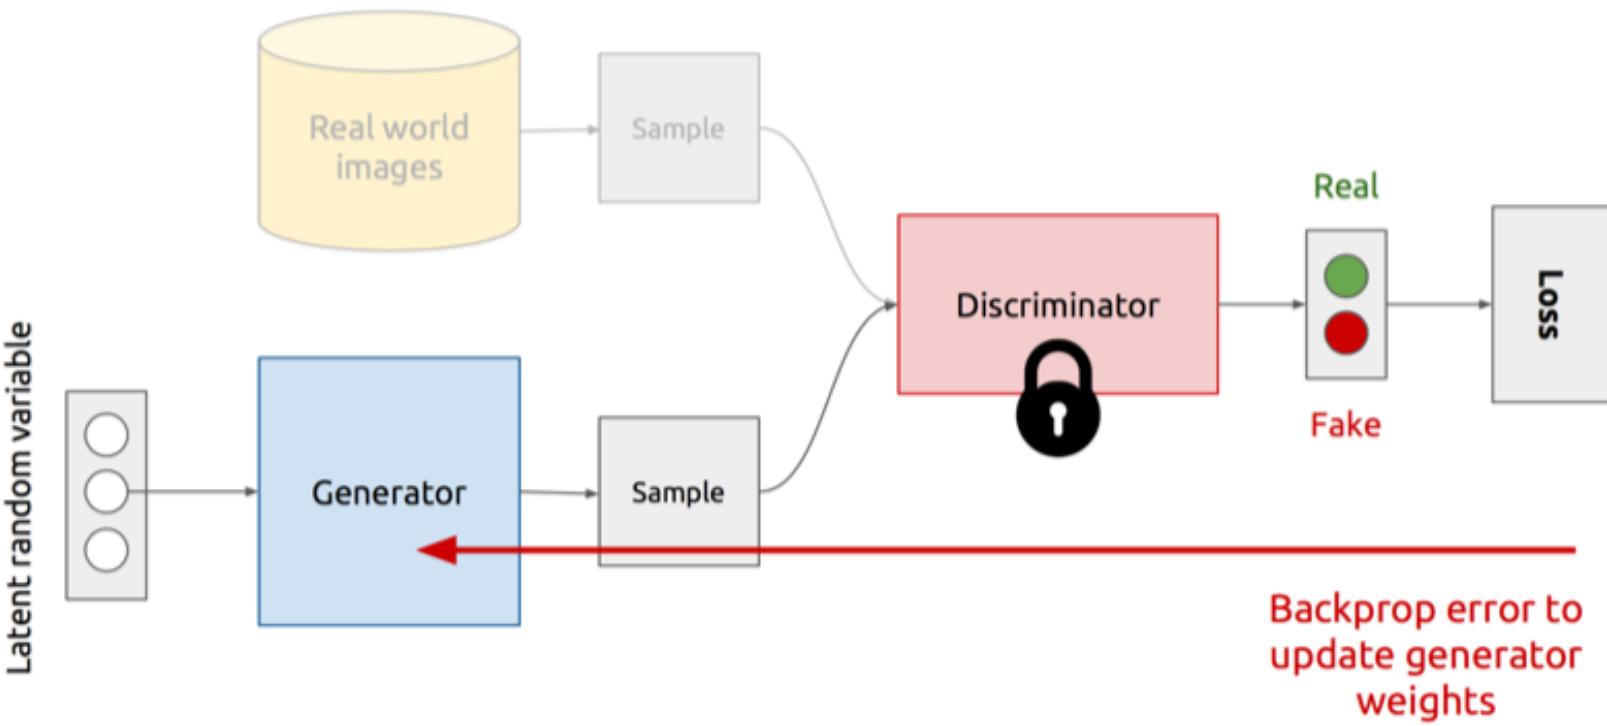


Training Discriminator



<https://www.slideshare.net/xavigiro/deep-learning-for-computer-vision-generative-models-and-adversarial-training-upc-2016>

Training Generator



<https://www.slideshare.net/xavigiro/deep-learning-for-computer-vision-generative-models-and-adversarial-training-upc-2016>

GAN's formulation

$$\min_G \max_D V(D, G)$$

- It is formulated as a **minimax game**, where:
 - The Discriminator is trying to maximize its reward $V(D, G)$
 - The Generator is trying to minimize Discriminator's reward (or maximize its loss)
- The Nash equilibrium of this particular game is achieved at:
 - $P_{data}(x) = P_{gen}(x) \quad \forall x$
 - $D(x) = \frac{1}{2} \quad \forall x$

$$V(D, G) = \boxed{\mathbb{E}_{x \sim p(x)} [\log D(x)]} + \boxed{\mathbb{E}_{z \sim q(z)} [\log(1 - D(G(z)))]}$$

Algorithm 1 Minibatch stochastic gradient descent training of generative adversarial nets. The number of steps to apply to the discriminator, k , is a hyperparameter. We used $k = 1$, the least expensive option, in our experiments.

for number of training iterations **do**

for k steps **do**

- Sample minibatch of m noise samples $\{\mathbf{z}^{(1)}, \dots, \mathbf{z}^{(m)}\}$ from noise prior $p_g(\mathbf{z})$.
- Sample minibatch of m examples $\{\mathbf{x}^{(1)}, \dots, \mathbf{x}^{(m)}\}$ from data generating distribution $p_{\text{data}}(\mathbf{x})$.
- Update the discriminator by ascending its stochastic gradient:

$$\nabla_{\theta_d} \frac{1}{m} \sum_{i=1}^m \left[\log D(\mathbf{x}^{(i)}) + \log (1 - D(G(\mathbf{z}^{(i)}))) \right].$$

end for

- Sample minibatch of m noise samples $\{\mathbf{z}^{(1)}, \dots, \mathbf{z}^{(m)}\}$ from noise prior $p_g(\mathbf{z})$.
- Update the generator by descending its stochastic gradient:

$$\nabla_{\theta_g} \frac{1}{m} \sum_{i=1}^m \log (1 - D(G(\mathbf{z}^{(i)}))).$$

end for

The gradient-based updates can use any standard gradient-based learning rule. We used momentum in our experiments.

**Discriminator
updates**

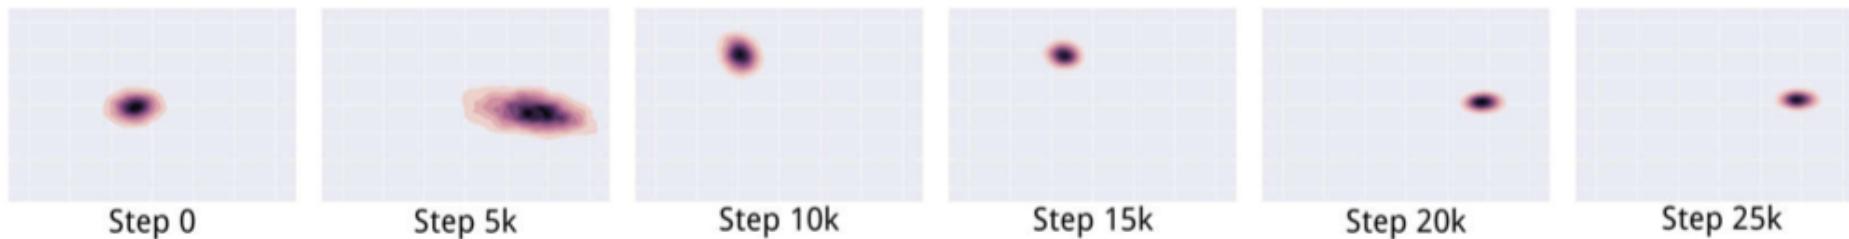
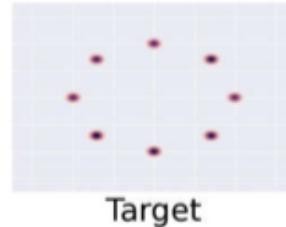
**Generator
updates**

GAN Challenges

Mode Collapse

$$\min_G \max_D V(G, D) \neq \max_D \min_G V(G, D)$$

- D in inner loop: convergence to correct distribution
- G in inner loop: place all mass on most likely point



Vanishing gradient strikes back again...

$$V(D, G) = \min_G \max_D V(D, G)$$
$$V(D, G) = \mathbb{E}_{x \sim p(x)} [\log D(x)] + \boxed{\mathbb{E}_{z \sim q(z)} [\log(1 - D(G(z)))]}$$

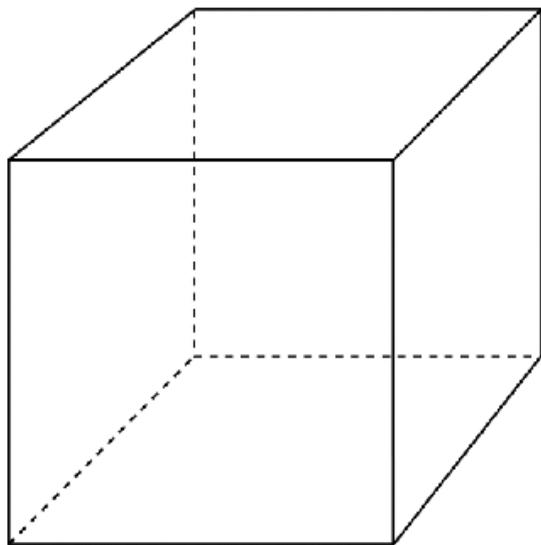
$$\nabla_{\theta_G} V(D, G) = \nabla_{\theta_G} \mathbb{E}_{z \sim q(z)} [\log(1 - D(G(z)))]$$

- $\nabla_a \log(1 - \sigma(a)) = \frac{-\nabla_a \sigma(a)}{1 - \sigma(a)} = \frac{-\sigma(a)(1 - \sigma(a))}{1 - \sigma(a)} = -\sigma(a) = -D(G(z))$
- Gradient goes to 0 if D is confident, i.e. $D(G(z)) \rightarrow 0$
- Minimize $-\mathbb{E}_{z \sim q(z)} [\log D(G(z))]$ for **Generator** instead (keep Discriminator as it is)

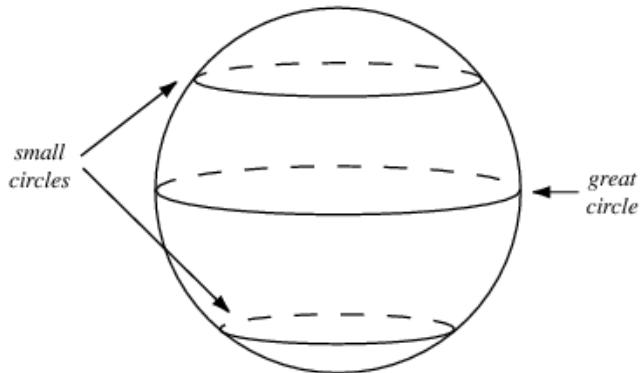
How to Train GAN (Tricks)

- Spherical z for sampling

- Dont sample from a Uniform distribution



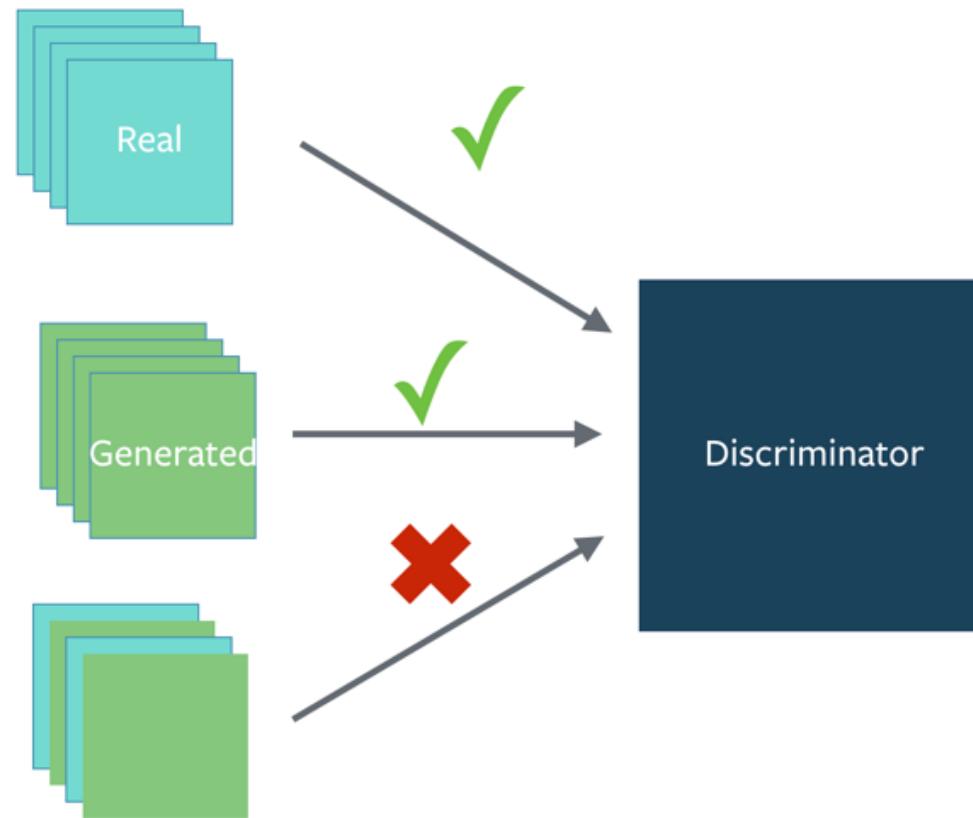
- Sample from a gaussian distribution



Src: <https://lanpartis.github.io/deep%20learning/2018/03/12/tricks-of-gans.html>

• BatchNorm

- Construct different mini-batches for real and fake, i.e. each mini-batch needs to contain only all real images or all generated images.
- when batchnorm is not an option use instance normalization (for each sample, subtract mean and divide by standard deviation).



- Conditional GANs

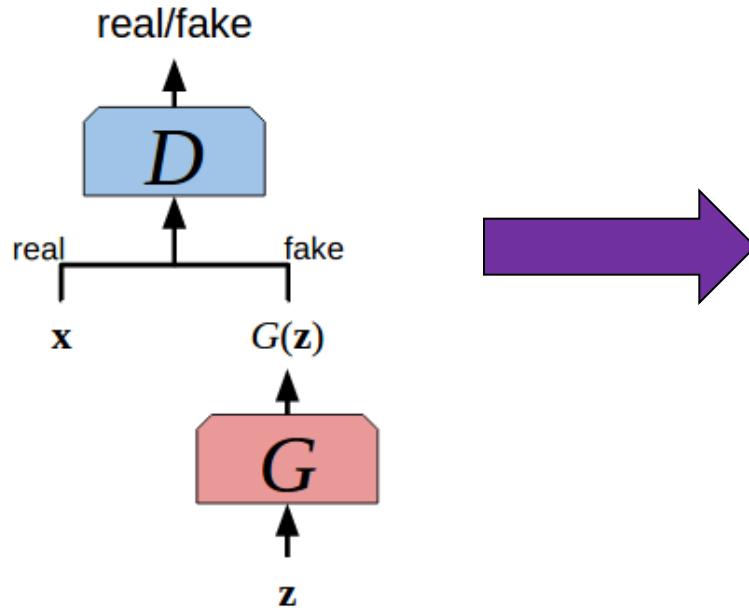


Figure 1: Generative Adversarial Networks

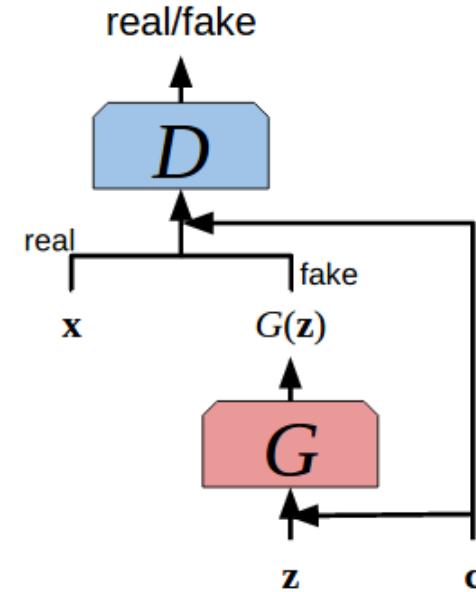
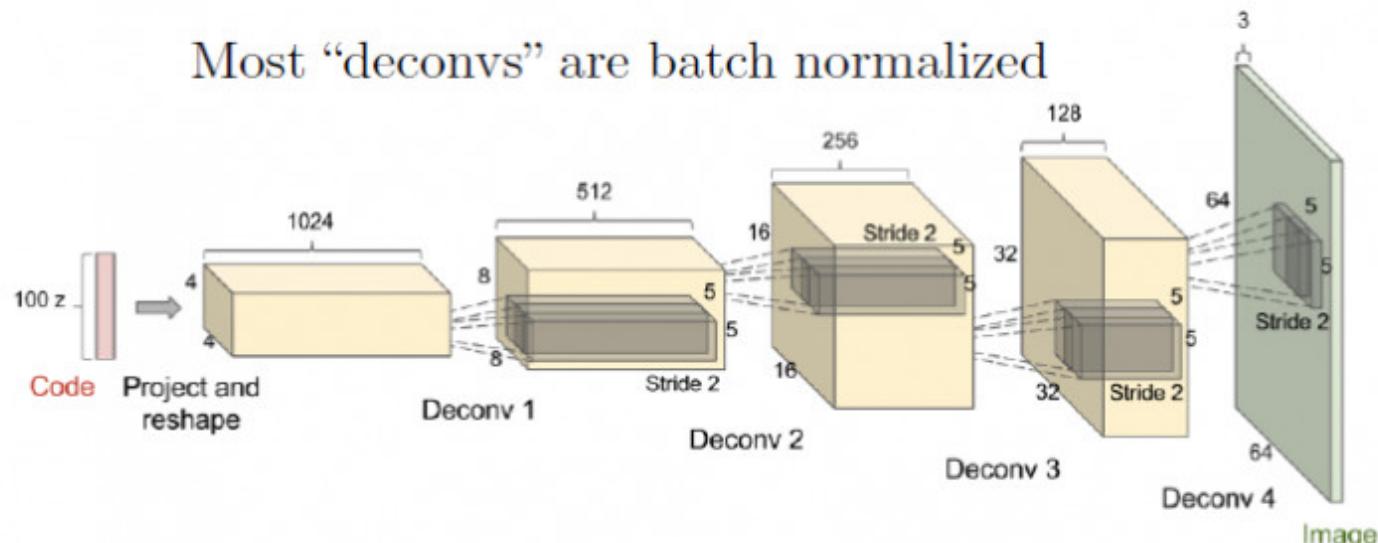


Figure 3: Conditional GAN

Src: <https://medium.com/jungle-book/towards-data-set-augmentation-with-gans-9dd64e9628e6>

Unsupervised Representation Learning

DCGAN Architecture:



(Radford et al 2015)

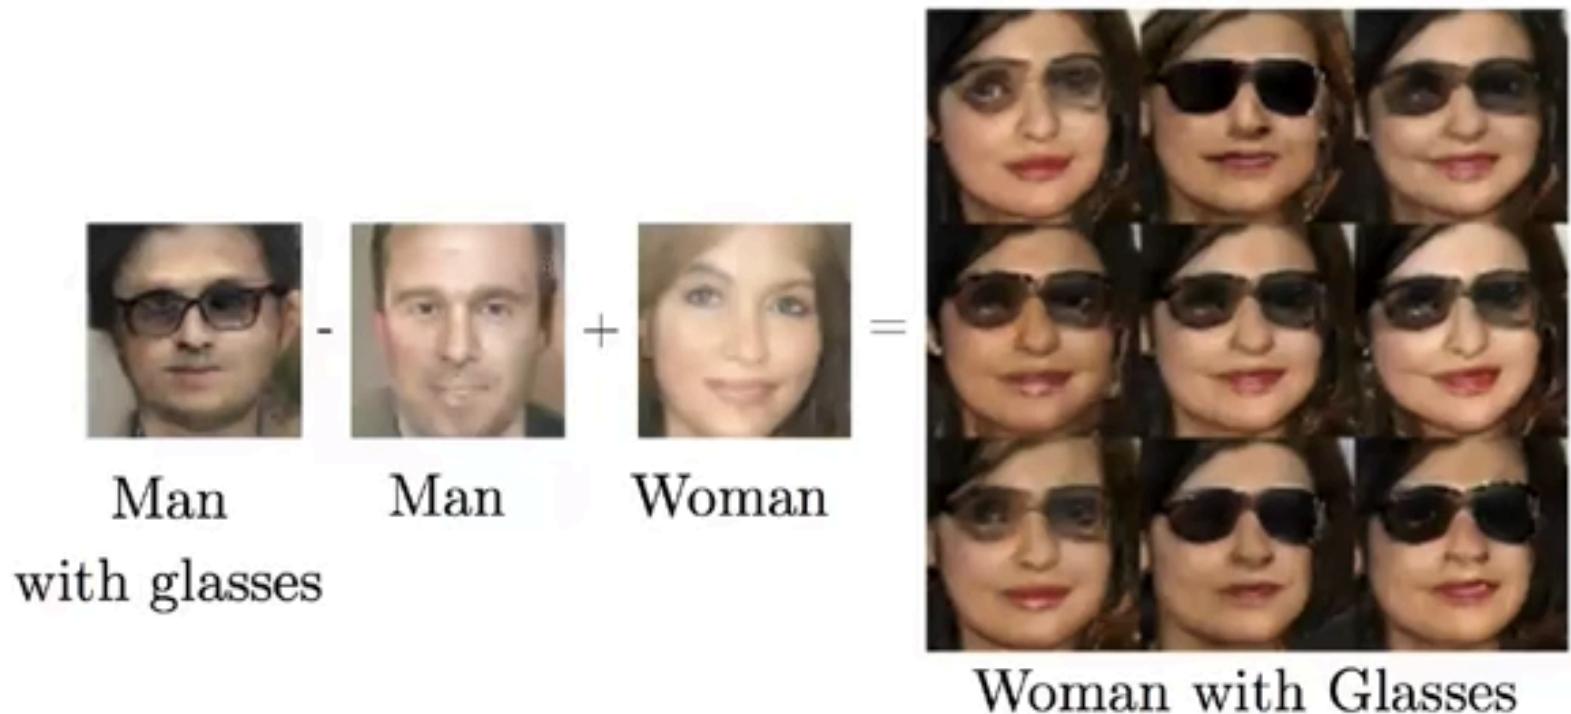
(Goodfellow 2016)

DCGAN for LSUN Image Generation:



(Radford et al 2015)

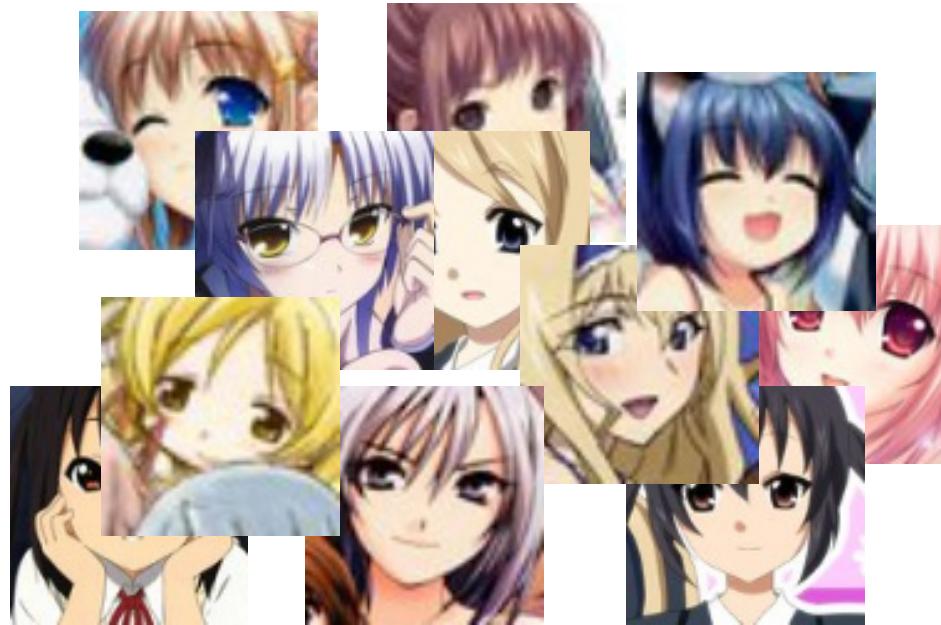
DCGAN for Vector Space Arithmetic:



(Radford et al, 2015)

Visualization of Image Generation Process

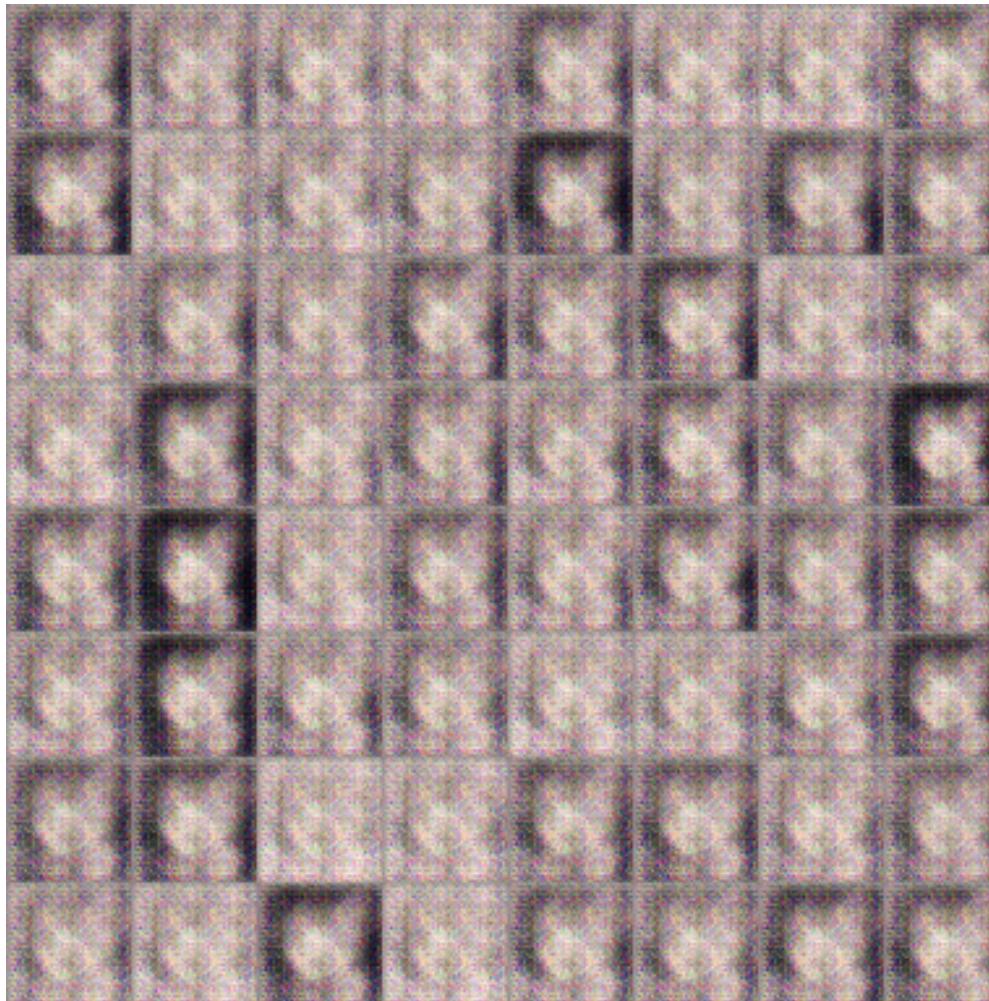
Generating 2nd element figures



Source of images: <https://zhuanlan.zhihu.com/p/24767059>
From Dr. HY Lee's notes.

DCGAN: <https://github.com/carpedm20/DCGAN-tensorflow>

GAN – generating 2nd element figures



100 rounds

This is fast, I think you can use your CPU

GAN – generating 2nd element figures



1000 rounds

GAN – generating 2nd element figures

2000 rounds



GAN – generating 2nd element figures

5000 rounds



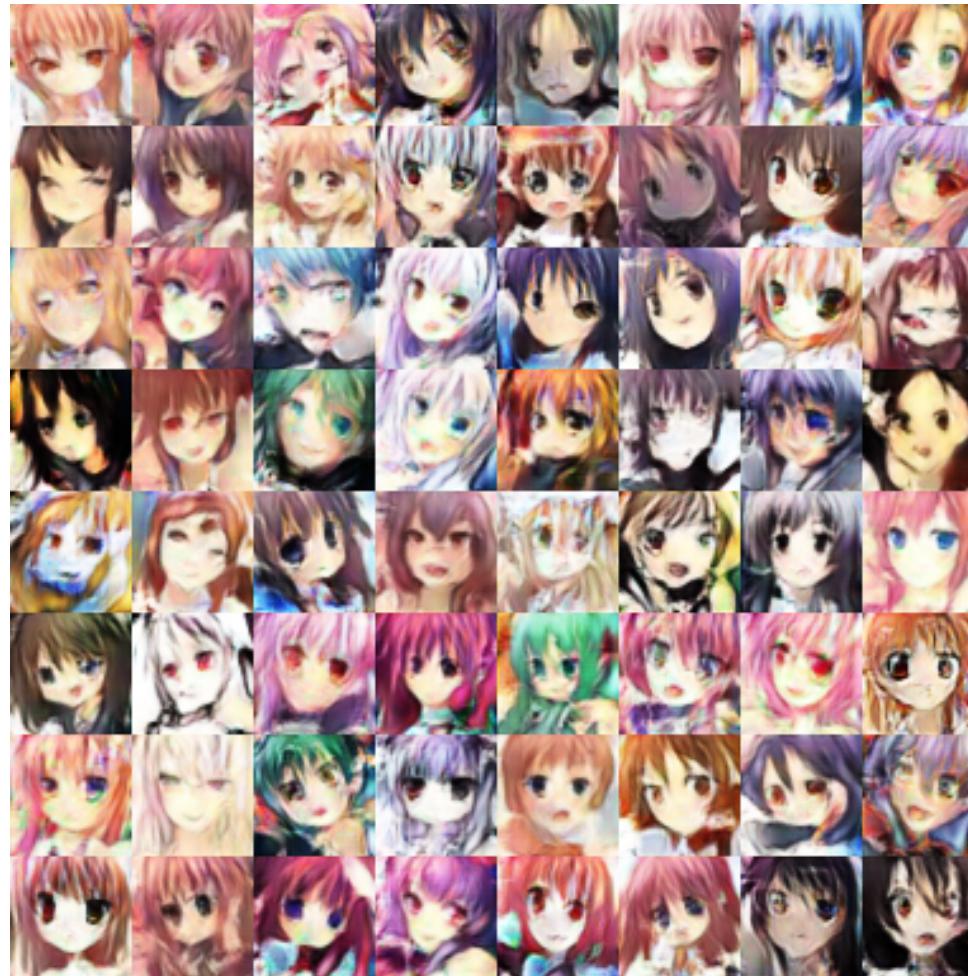
GAN – generating 2nd element figures

10,000 rounds

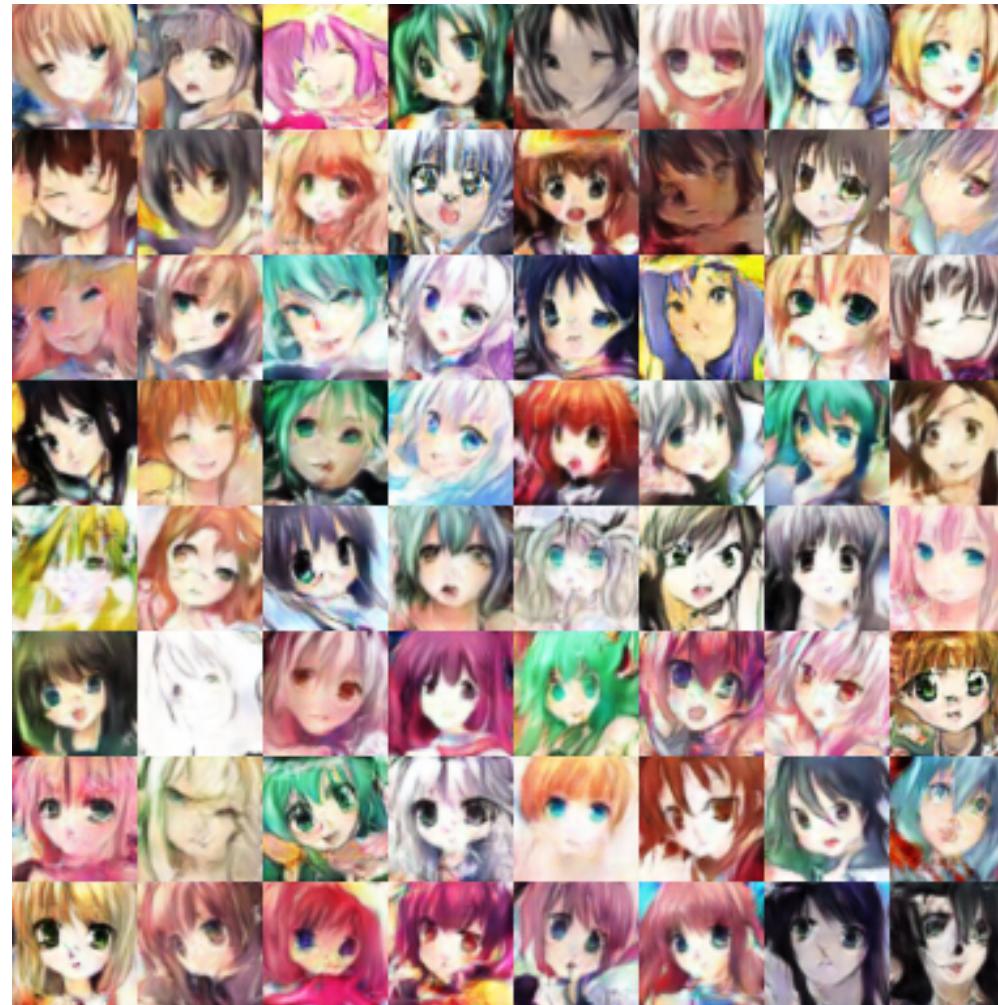


GAN – generating 2nd element figures

20,000 rounds



GAN – generating 2nd element figures



50,000 rounds

- Introduction of GAN
- GAN applications

Applications of GAN

- Adversarial Learning
- Low Level Vision
- Domain Adaptation
- Image translation

References

- <http://slazebni.cs.illinois.edu/spring17/>
- <https://cs.uwaterloo.ca/~mli/Deep-Learning-2017-Lecture7GAN.ppt>